

Introduction to Processes

01/19/2023

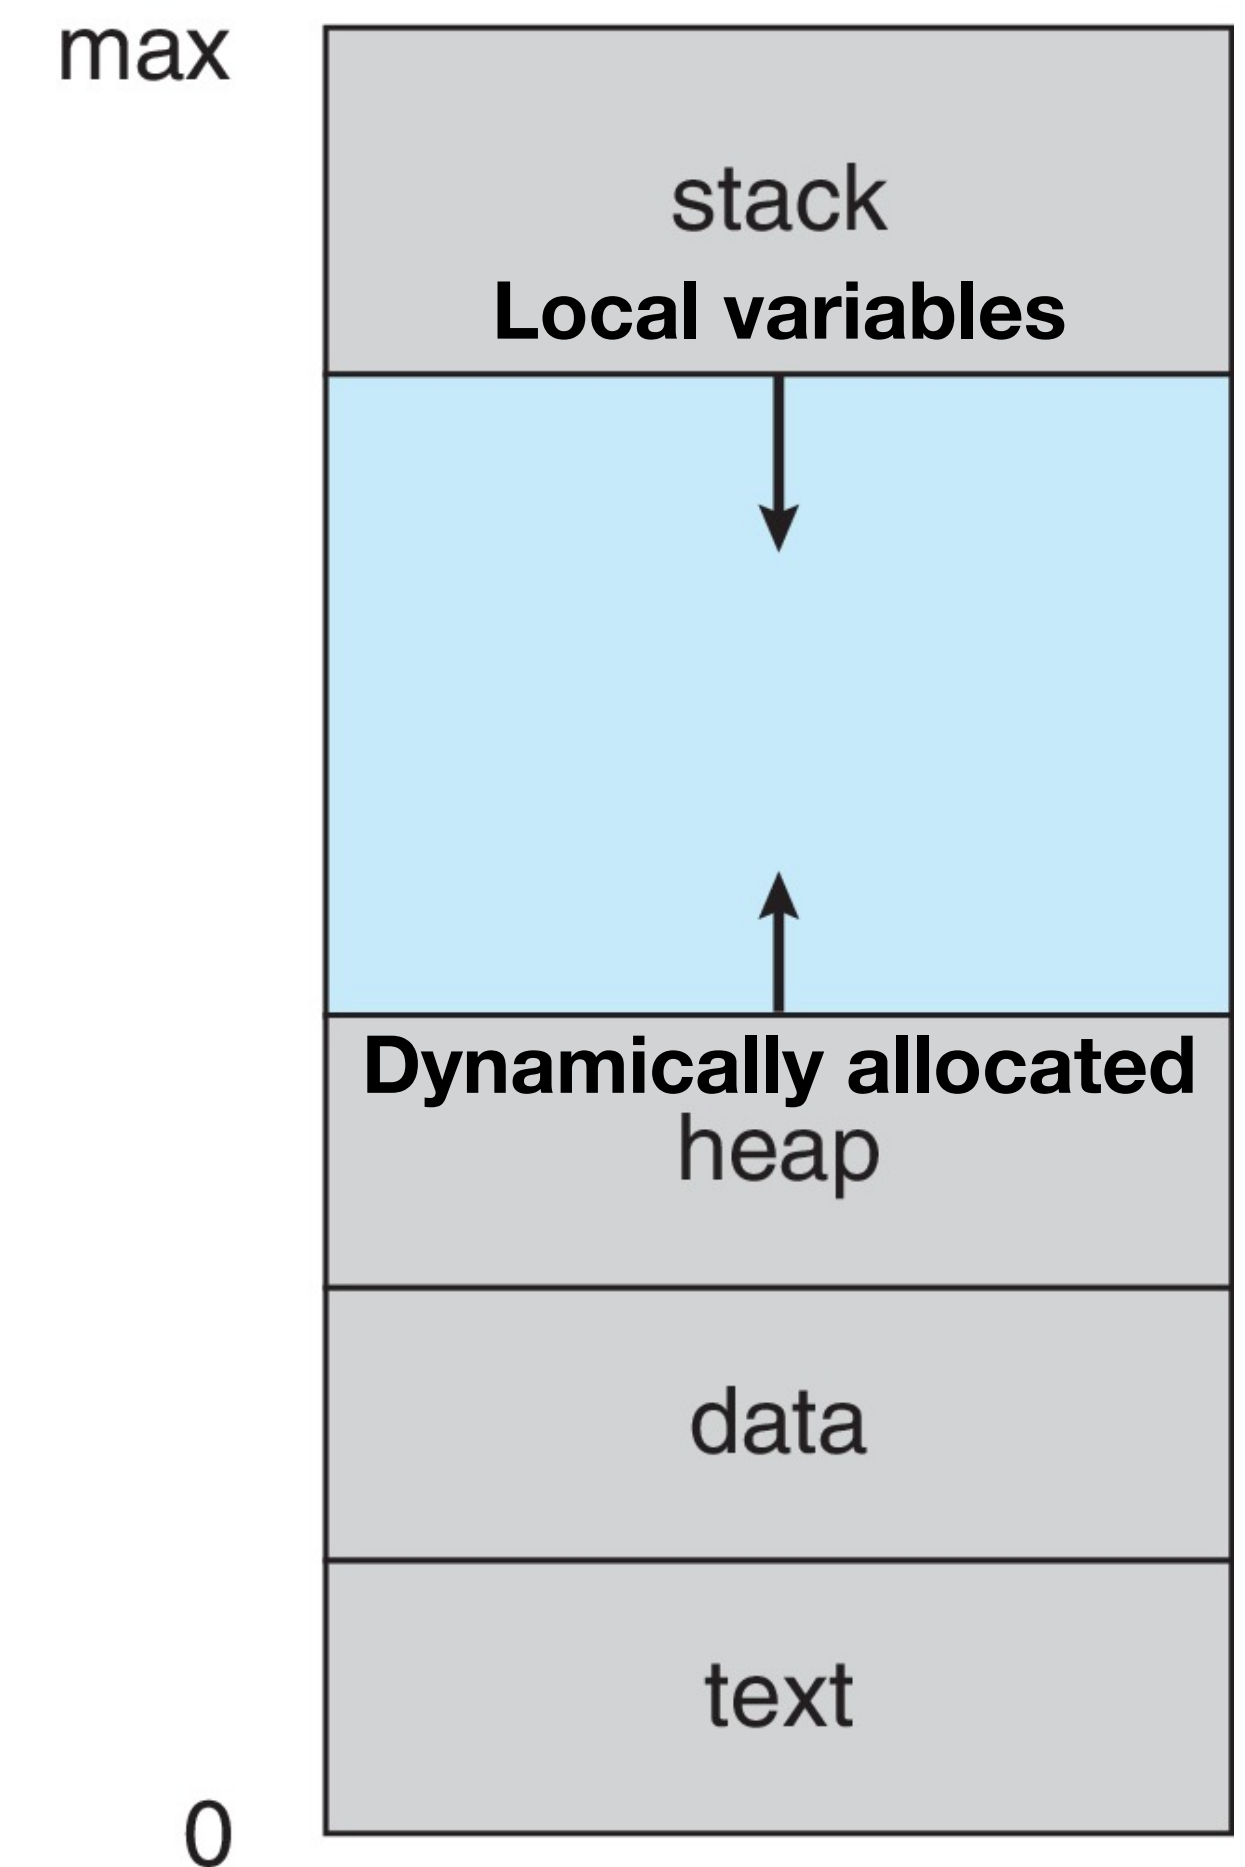
Professor Amanda Bienz

CPU Virtualization

- Provides illusion of many CPUs
- **Time sharing:** Running one process, then stopping it and running another
- Potential cost is *performance*

A Process

- A process is a **running program**
- A process is comprised of:
 - Virtual memory (address space)
 - Instructions
 - Data section
 - Registers
 - Program counter : register that holds the address of the instruction being executed
 - Stack pointer : register that stores address of last program request in stack



Process API

- **Create:** create a new process to run a program
- **Destroy:** halt a runaway process
- **Wait:** wait for a process to stop running
- **Miscellaneous control:** some method to suspend a process and then resume it
- **Status:** get some status information about a process

Process Creation

- **Load a program code into memory** (into address space of the process)
 - Programs initially reside on disk in an executable format
 - OS performs the loading process **lazily**
 - Loading pieces of code or data as they are needed during the program
- **The program's run-time stack is allocated**
 - Use the stack for *local variables, function parameters, and return address*
 - Initialize the stack with arguments (argc, argv)

Process Creation (cont'd)

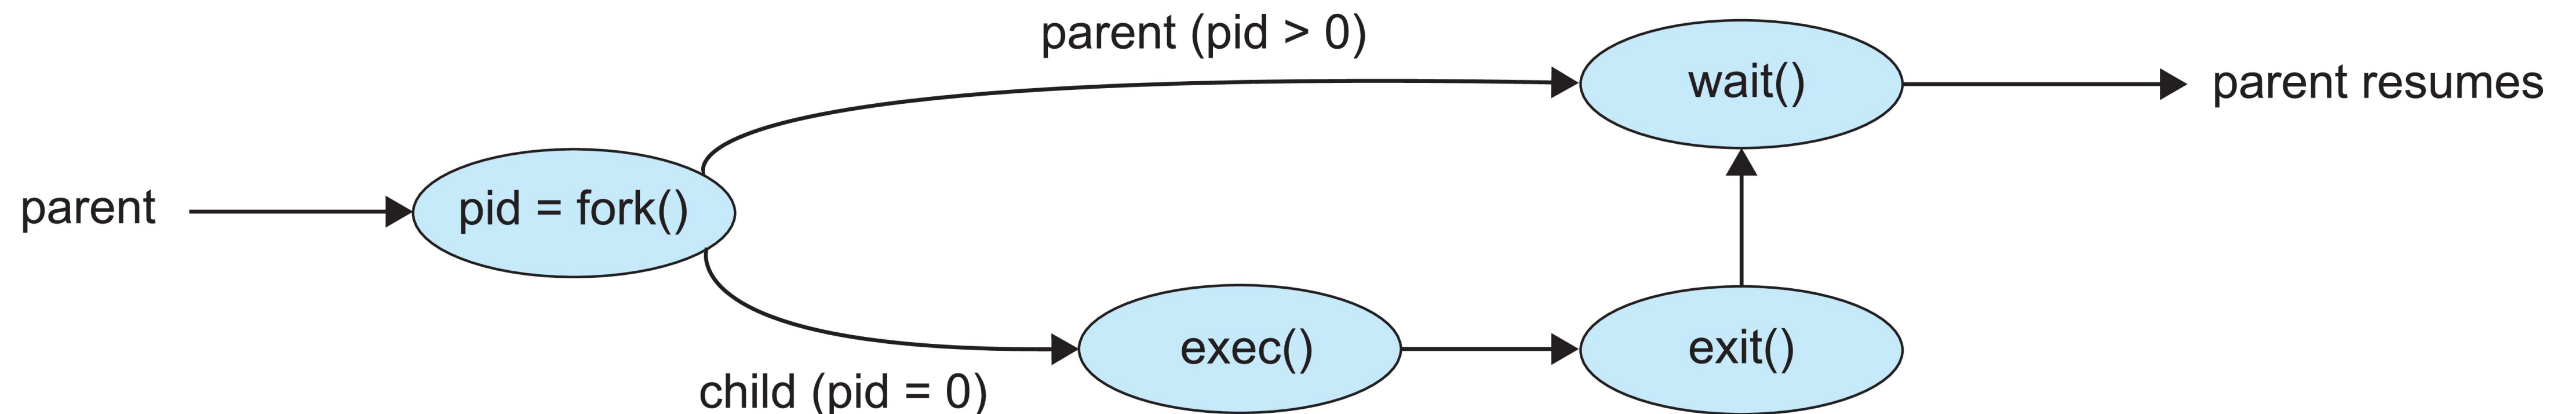
- **The program's heap is created**
 - Used for explicitly requested dynamically allocated data
 - Program request space with `malloc()` and free space by calling `free()`
- **The OS does some other initialization tasks**
 - Input/Output setup
 - Each process, by default, has three open file descriptors
 - Standard input, output, error
- **Start the program running at the entry point, namely `main()`**
 - The OS transfers control of the CPU to the newly-created process

Parent/Child Processes

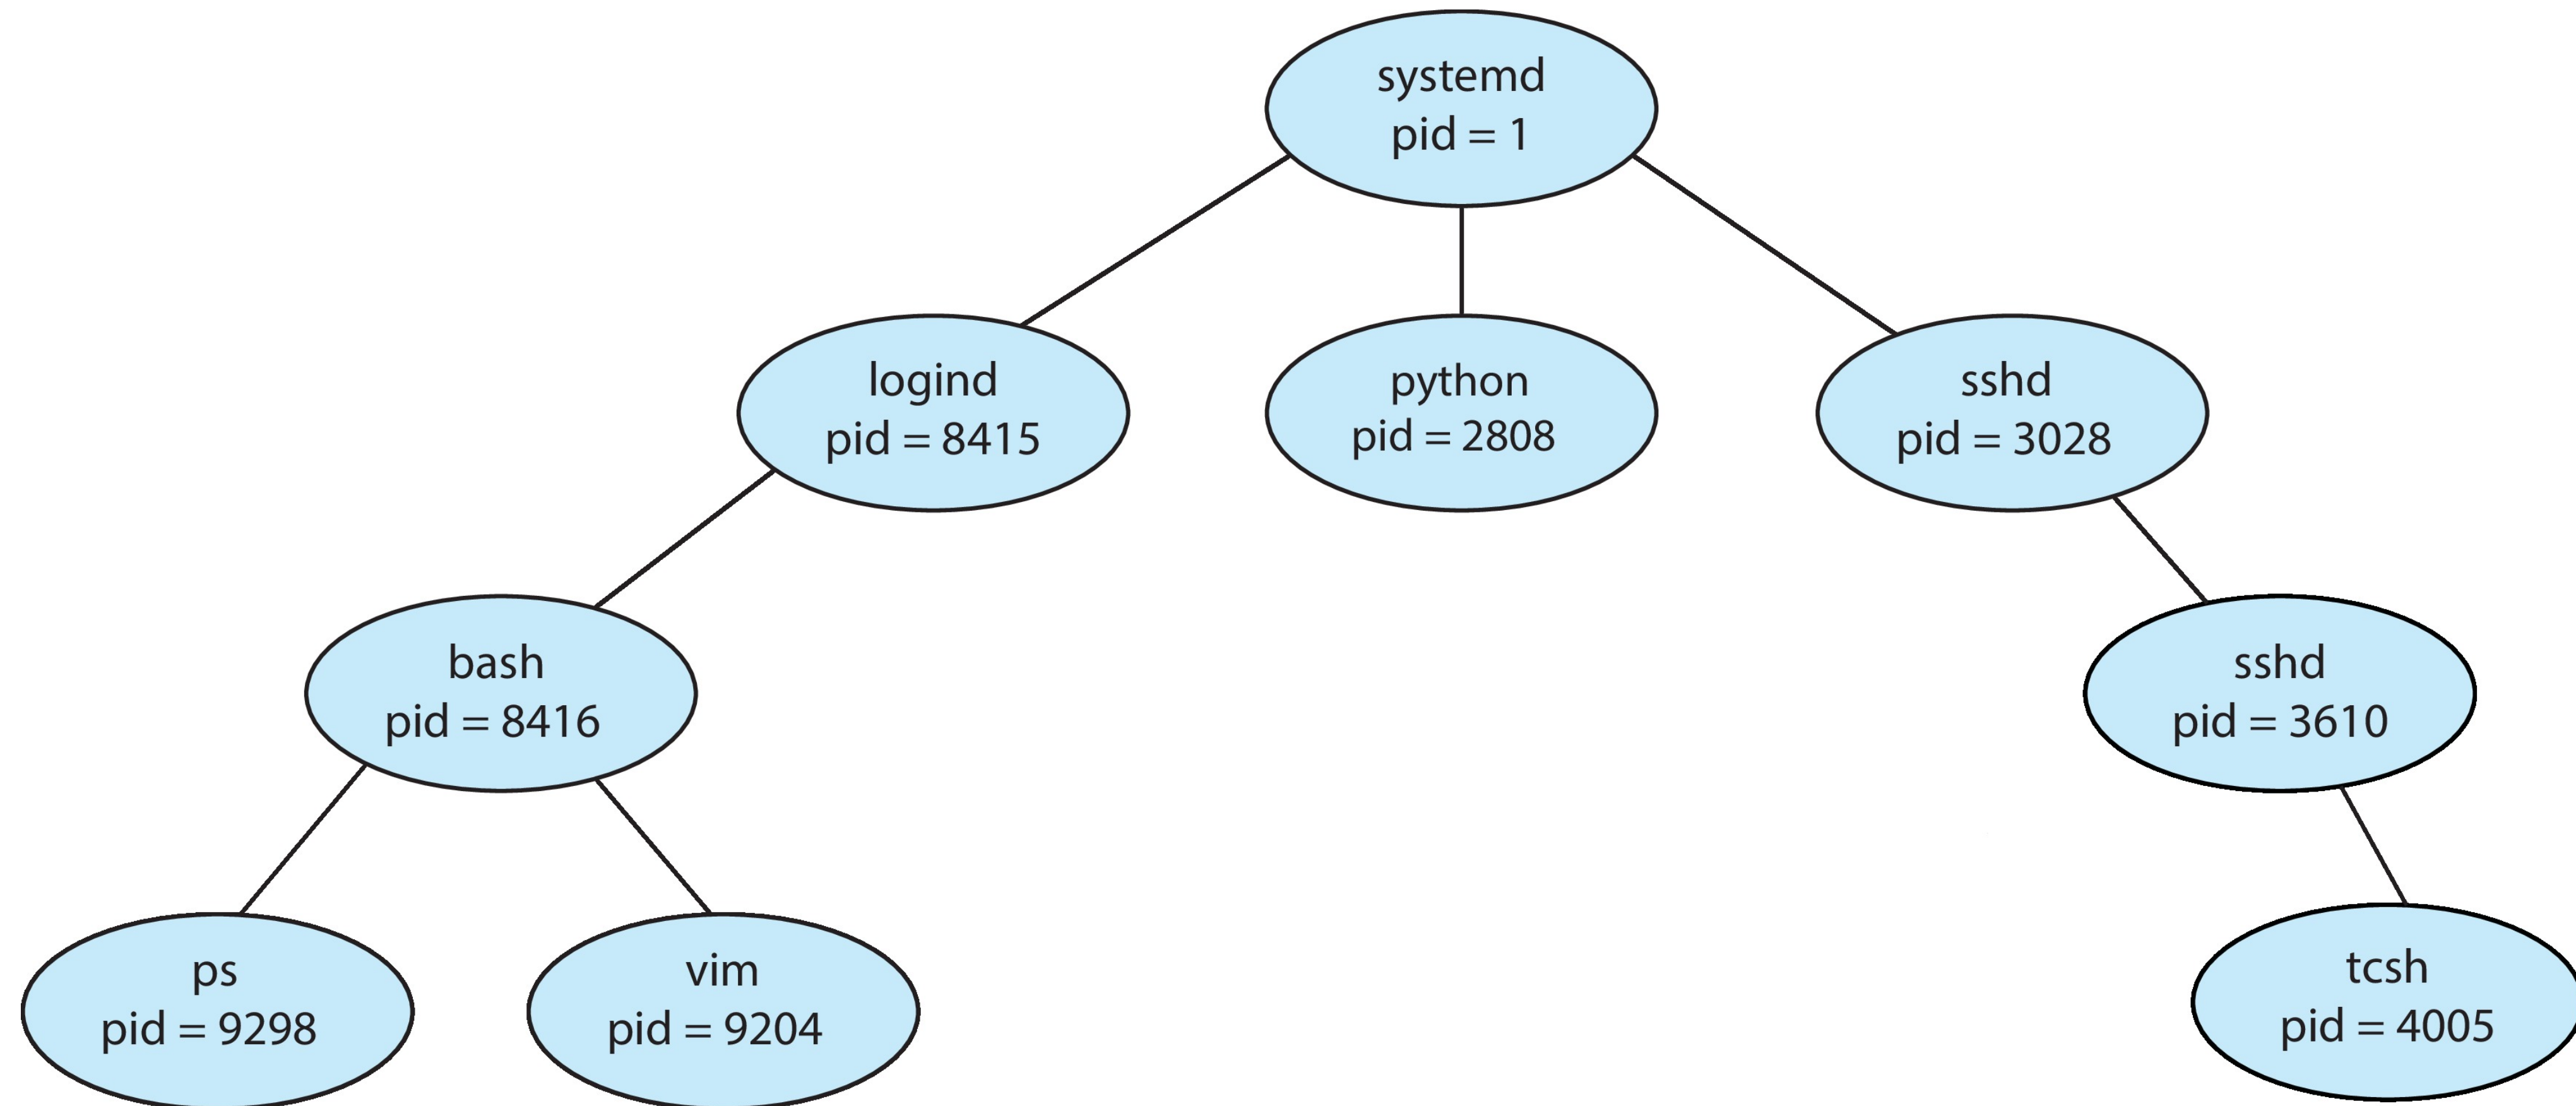
- Parent process can create a child process
 - Child process can create other children, etc
- Each process has a process identifier (pid)
- Resource sharing:
 - Parent and children can share all resources, children share a subset, or they share no resources
- Execution options:
 - Parent and children execute concurrently, or parent waits until children terminate

Parent/Children Processes

- **fork()** system call creates new process
- **exec()** system call used after a **fork()** to replace the process' memory space with a new program
- Parent process calls **wait()** waiting for the child to terminate



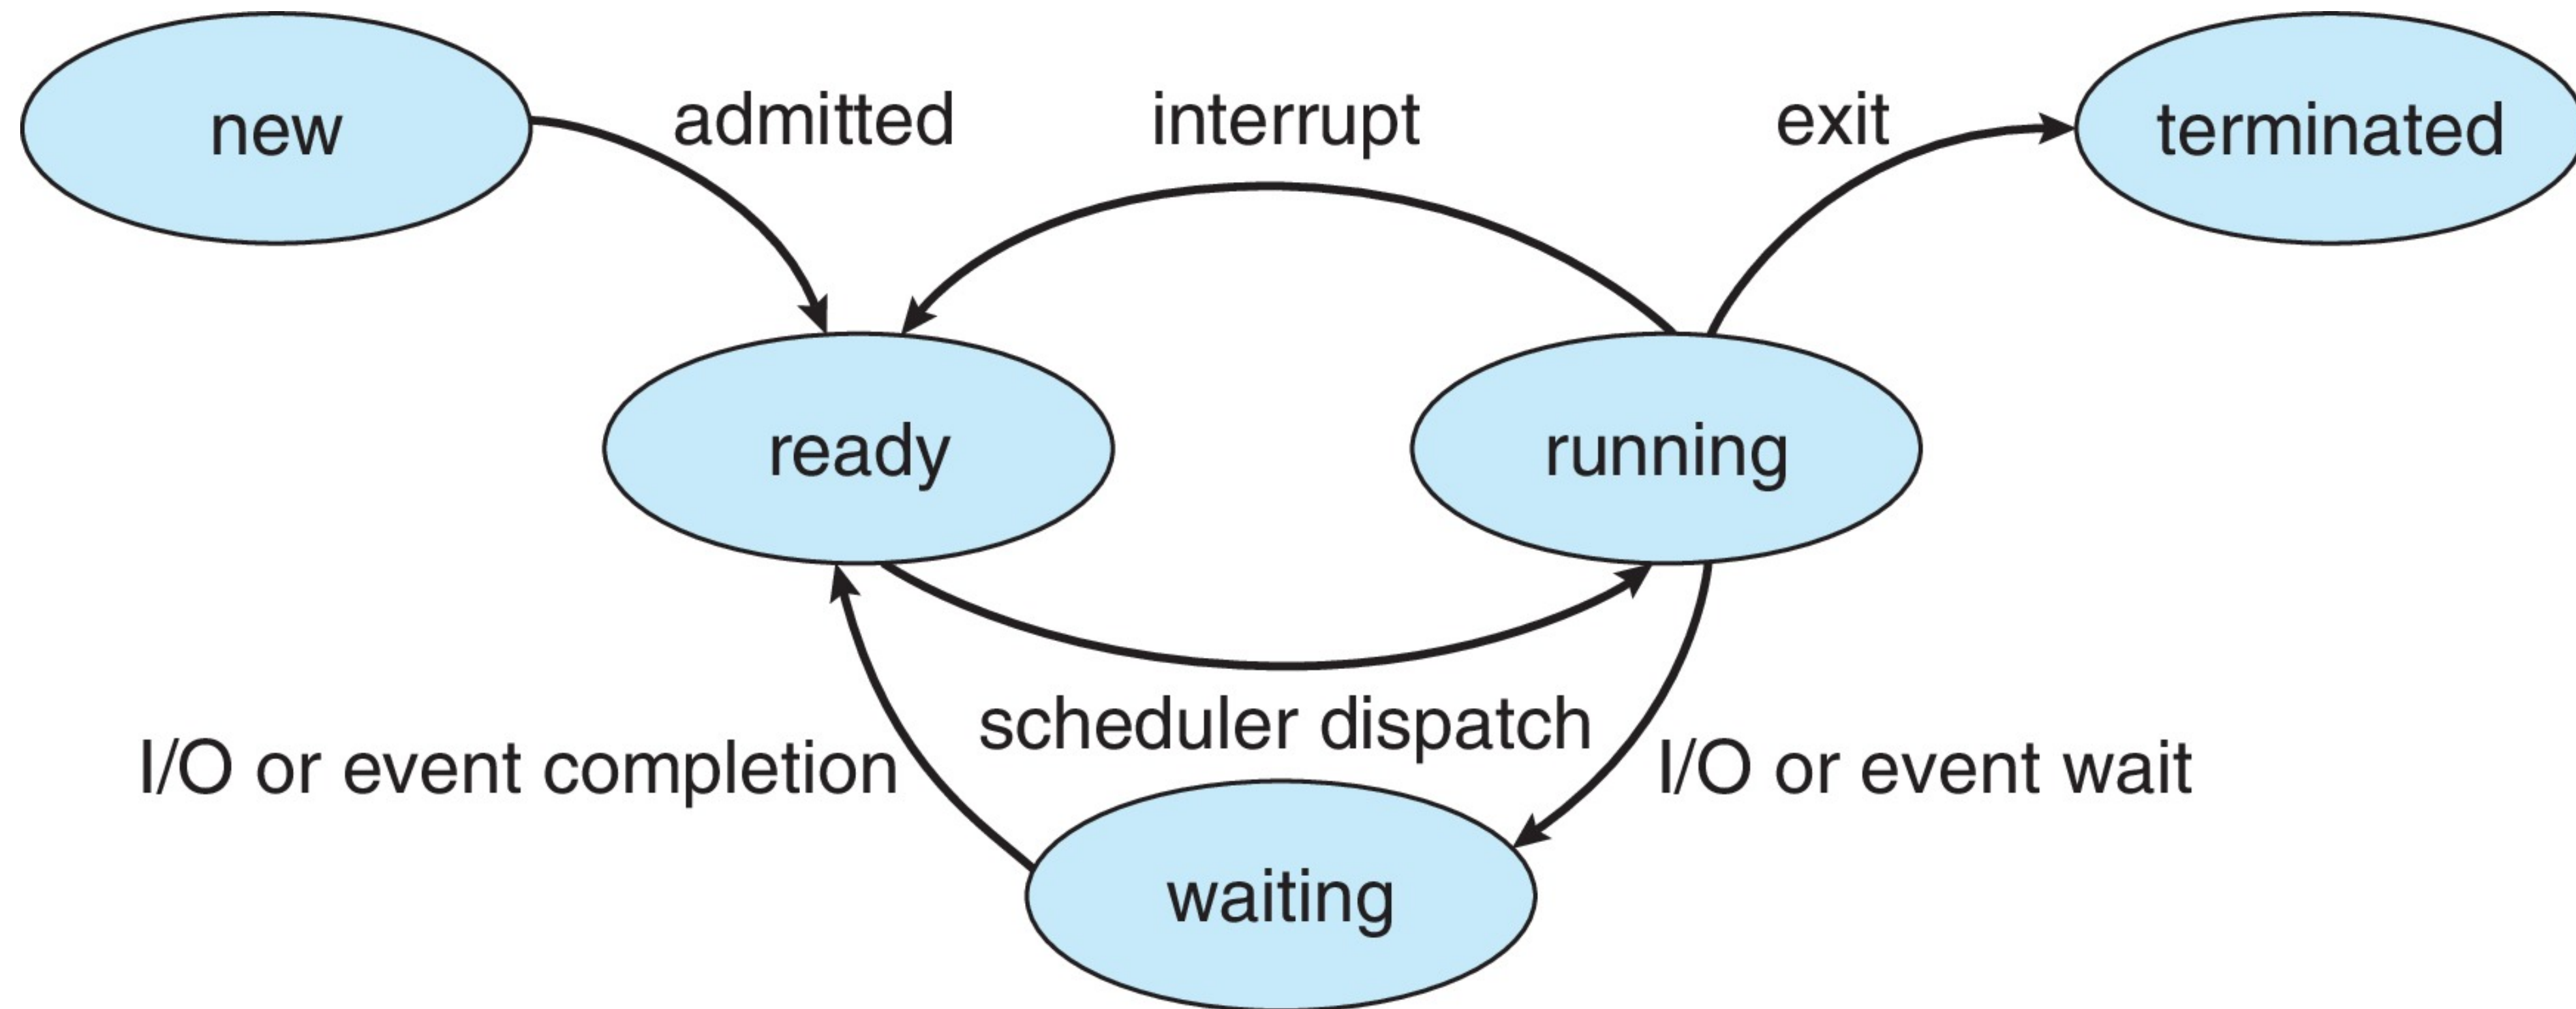
Linux Process Tree



Process State

- The state of a process is defined (in part) by the current activity of that process
 - **New:** process is being created
 - **Running:** instructions are being executed
 - **Waiting (or blocked):** process is waiting for some event to occur (i.e. I/O, completion signal)
 - **Ready:** process waiting to be assigned a processor
 - **Terminated:** process has finished execution

Process State



Process Control Block

Information associated with each process

- Process state : running, waiting, etc
- Program counter : location of instruction to next execute
- CPU registers : contents of all process-centric registers
- CPU scheduling information : priorities, scheduling queue pointers
- Memory-management information : memory allocated to the process
- Accounting information : CPU used, clock time elapsed since start, time limits
- I/O status information : I/O devices allocated to process, list of open files

