

Concurrency : Common Errors

04/06/2021

Professor Amanda Bienz

Lots of ways to mess up with concurrency

- Bugs are often unpredictable and hard to reproduce
- Subtle interactions between multiple actors that are timing dependent
- Adding debugging code or using the debugger can change the timing, making the bug temporarily hide
- Need to understand common concurrency problems
 1. Critical Section Problems (data races) — we've seen this
 2. Deadlock
 3. Starvation

What Types of Bugs Exist

- Focus on four major open-source applications
 - MySQL, Apache, Mozilla, OpenOffice

Application	What it does	Non-Deadlock	Deadlock
MySQL	Database Server	14	9
Apache	Web Server	13	4
Mozilla	Web Browser	41	16
Open Office	Office Suite	6	2
Total		74	31

Bugs In Modern Applications

Non-Deadlock Bugs

- Make up a majority of concurrency bugs
- Two major types of non deadlock bugs:
 - Atomicity violation
 - Order violation

Atomicity-Violation Bugs

- The desired serializability among multiple memory accesses is violated
 - Simple example found in MySQL:
 - Two different threads access the field `proc_info`

```
1      Thread1 ::  
2      if (thd->proc_info) {  
3          ...  
4          fputs (thd->proc_info , ...);  
5          ...  
6      }  
7  
8      Thread2 ::  
9      thd->proc_info = NULL;
```

Atomicity-Violation Bugs (Cont.)

- Solution : Simply add locks around the shared-variable references

```
1  pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
2
3  Thread1 :
4  pthread_mutex_lock(&lock);
5  if(thd->proc_info) {
6      ...
7      fputs(thd->proc_info , ...);
8      ...
9  }
10 pthread_mutex_unlock(&lock);
11
12 Thread2 :
13 pthread_mutex_lock(&lock);
14 thd->proc_info = NULL;
15 pthread_mutex_unlock(&lock);
```

Order-Violation Bugs

- The desired order between two memory accesses is flipped
 - i.e. A should always be executed before B, but the order is not enforced during execution
- Example: The code in Thread2 seems to assume that the variable mThread has already been initialized (and is not NULL).

```
1  Thread1 ::
2  void init() {
3      mThread = PR_CreateThread(mMain, ...);
4  }
5
6  Thread2 ::
7  void mMain(...) {
8      mState = mThread->State
9  }
```

Order-Violation Bugs (Cont.)

- Solution : enforce ordering using condition variables


```
1  pthread_mutex_t mtLock = PTHREAD_MUTEX_INITIALIZER;
2  pthread_cond_t mtCond = PTHREAD_COND_INITIALIZER;
3  int mtInit = 0;
4
5  Thread 1::
6  void init() {
7      ...
8      mThread = PR_CreateThread(mMain,...);
9
10     // signal that the thread has been created.
11     pthread_mutex_lock(&mtLock);
12     mtInit = 1;
13     pthread_cond_signal(&mtCond);
14     pthread_mutex_unlock(&mtLock);
15     ...
16 }
17
18 Thread2: :
19 void mMain(...) {
20     ...
```


Order-Violation Bugs (Cont.)

```
21      // wait for the thread to be initialized ...
22      pthread_mutex_lock(&mtLock);
23      while(mtInit == 0)
24          pthread_cond_wait(&mtCond, &mtLock);
25      pthread_mutex_unlock(&mtLock);
26
27      mState = mThread->State;
28      ...
29  }
```

The Dining Philosophers

- Assume there are five “philosophers” sitting around a table
 - Between each pair of philosophers is a **single fork** (five total)
 - The philosophers each have times where they **think**, and don’t need any forks, and times where they **eat**
 - In order to eat, a philosopher needs **two forks**, both the one on their left and the one on their right
 - **There is contention for these forks**



The Dining Philosophers (Cont.)

- Key Challenge:
 - There is no deadlock
 - No philosopher starves and never gets to eat
 - Concurrency is high

```
while (1) {  
    think();  
    getforks();  
    eat();  
    putforks();  
}
```

Basic loop of each philosopher

```
// helper functions  
int left(int p) { return p; }  
  
int right(int p) {  
    return (p + 1) % 5;  
}
```

Helper functions (Downey's solutions)

- Philosopher p wishes to refer to the fork on their left : call left(p)
- Philosopher p wishes to refer to the fork on their right : call right(p)

The Dining Philosophers (Cont.)

- We need some semaphore, one for each fork: `set_t forks[5]`

```
1  void getforks() {  
2    sem_wait(forks[left(p)]);  
3    sem_wait(forks[right(p)]);  
4  }  
5  
6  void putforks() {  
7    sem_post(forks[left(p)]);  
8    sem_post(forks[right(p)]);  
9  }
```

The `getforks()` and `putforks()` Routines (Broken Solution)

- What is wrong with this solution?

The Dining Philosophers (Cont.)

- We need some semaphore, one for each fork: `set_t forks[5]`

```
1  void getforks() {
2      sem_wait(forks[left(p)]);
3      sem_wait(forks[right(p)]);
4  }
5
6  void putforks() {
7      sem_post(forks[left(p)]);
8      sem_post(forks[right(p)]);
9  }
```

The `getforks()` and `putforks()` Routines (Broken Solution)

- Deadlock occur!
 - If each philosopher happens to grab the fork on their left before any grabs a fork on their right
 - Each will be stuck holding one fork and waiting for the other, forever

A Solution : Breaking the Dependency

- Change how forks are acquired
 - Let's assume that philosopher 4 acquire the forks in a different order

```
1  void getforks() {  
2    if (p == 4) {  
3        sem_wait(forks[right(p)]);  
4        sem_wait(forks[left(p)]);  
5    } else {  
6        sem_wait(forks[left(p)]);  
7        sem_wait(forks[right(p)]);  
8    }  
9 }
```

- There is no situation where each philosopher grabs one fork and is stuck waiting for the other. The cycle of waiting is broken.

Deadlock Bugs

Thread 1:

`lock(L1);`


`lock(L2);`

Thread 2:

`lock(L2);`

`lock(L1);`

- The presence of **a cycle**
 - Thread1 is holding a lock L1 and waiting for another one, L2.
 - Thread2 that holds lock L2 is waiting for L1 to be release.



Why Do Deadlocks Occur?

- Reason 1:
 - In large code bases, complex dependencies arise between components
- Reason 2:
 - Due to the nature of encapsulation
 - Hide details of implementations and make software easier to build in a modular way
 - Such modularity does not mesh well with locking

Why Do Deadlocks Occur? (Cont.)

- Example : Java Vector class and the method AddAll()

```
1    Vector v1,v2;  
2    v1.AddAll(v2);
```

- Locks for both the vector being added to (v1) and the parameter (v2) need to be acquired
 - The routine acquires said locks in some arbitrary order (v1 then v2)
 - If some other thread calls v2.AddAll(v1) at nearly the same time : we have the potential for deadlock

Conditional for Deadlock

- Four conditions need to hold for a deadlock to occur

Condition	Description
Mutual Exclusion	Threads claim exclusive control of resources that they require.
Hold-and-wait	Threads hold resources allocated to them while waiting for additional resources
No preemption	Resources cannot be forcibly removed from threads that are holding them.
Circular wait	There exists a circular chain of threads such that each thread holds one more resources that are being requested by the next thread in the chain

- If any of these four conditions are not met, deadlock cannot occur

Prevention : Circular Wait

- Provide a total ordering on lock acquisition
 - This approach requires careful design of global locking strategies
- Example :
 - There are two locks in the system (L1 and L2)
 - We can prevent deadlock by always acquiring L1 before L2

Prevention : Hold-and-wait

- Acquire all locks at once, atomically

```
1    lock (prevention) ;  
2    lock (L1) ;  
3    lock (L2) ;  
4    ...  
5    unlock (prevention) ;
```

- This code guarantees that no untimely thread switch can occur in the midst of lock acquisition
- Problem:
 - Require us to know when calling a routine exactly which locks must be held and to acquire them ahead of time
 - Decrease concurrency

Prevention : No Preemption

- Multiple lock acquisition often gets us into trouble because when waiting for one lock we are holding another
- trylock()
 - Used to build a deadlock-free, ordering-robust lock acquisition protocol
 - Grab the lock (if it is available)
 - Or, return -1 : you should try again later

```
1    top:
2        lock(L1);
3        if( tryLock(L2) == -1 ) {
4            unlock(L1);
5            goto top;
6        }
```

Prevention : No Preemption (Cont.)

- Livelock:
 - Both systems are running through the code sequence over and over again
 - Progress is not being made
 - Solution : add a random delay before looping back and trying the entire thing over again

Prevention : Mutual Exclusion

- Wait-Free
- Using powerful hardware instruction
- You can build data structures in a manner that does not require explicit locking

```
1  int CompareAndSwap(int *address, int expected, int new) {  
2      if(*address == expected) {  
3          *address = new;  
4          return 1; // success  
5      }  
6      return 0;  
7  }
```

Prevention : Mutual Exclusion (Cont.)

- We now wanted to atomically increment a value by a certain amount:

```
1  void AtomicIncrement(int *value, int amount) {  
2      do{  
3          int old = *value;  
4      }while( CompareAndSwap(value, old, old+amount) == 0 );  
5  }
```

- Repeatedly tried to update the value to the new amount and uses the compare-and-swap to do so
- No lock is acquired
- No deadlock can arise
- Livelock is still a possibility

Prevention : Mutual Exclusion (Cont.)

- More complex example : list insertion

```
1  void insert(int value){  
2      node_t * n = malloc(sizeof(node_t));  
3      assert( n != NULL );  
4      n->value = value ;  
5      n->next  = head;  
6      head    = n;  
7  }
```

- If called by multiple threads at the same time, this code has a race condition

Prevention : Mutual Exclusion (Cont.)

- Solution :
- Surrounding this code with a lock acquire and release

```
1  void insert(int value){
2      node_t * n = malloc(sizeof(node_t));
3      assert( n != NULL );
4      n->value = value ;
5      lock(listlock); // begin critical section
6      n->next  = head;
7      head     = n;
8      unlock(listlock) ; //end critical section
9  }
```

- Wait-free manner using the compare-and-swap instruction

```
1  void insert(int value) {
2      node_t *n = malloc(sizeof(node_t));
3      assert(n != NULL);
4      n->value = value;
5      do {
6          n->next = head;
7      } while (CompareAndSwap(&head, n->next, n));
8  }
```

Automatic Anti-Deadlock Techniques

- Deadlock avoidance : OS or runtime should be smart enough to keep deadlock from being able to happen
- Deadlock prevention : detect that deadlock is about to or is occurring and then break it (by rebooting or killing threads?)

Deadlock Avoidance via Scheduling


- In some scenarios deadlock avoidance is preferable
 - Global knowledge is required
 - Which locks various threads might grab during their execution
 - Subsequently schedules said threads in a way as to guarantee no deadlock can occur

Example of Deadlock Avoidance via Scheduling(1)

- We have two processors and four threads:
- Lock acquisition demand of the threads:

	T1	T2	T3	T4
L1	yes	yes	no	no
L2	yes	yes	yes	no

- A smart scheduler could compute that as long as T1 and T2 are not run at the same time, no deadlock could ever arise



Example of Deadlock Avoidance via Scheduling(2)

- More contention for the same resources

	T1	T2	T3	T4
L1	yes	yes	yes	no
L2	yes	yes	yes	no

- A possible schedule that guarantees that no deadlock could ever occur



- The total time to complete the jobs is lengthened considerably

Detect and Recover

- Allow deadlock to occasionally occur and then take some action
 - Example : If OS froze, you would reboot it
- Many database systems employ deadlock detection and recovery technique
 - A deadlock detector runs periodically
 - Building a resource graph and checking it for cycles
 - In deadlock, the system needs to be restarted