

Introduction to Concurrency

03/09/2021

Professor Amanda Bienz

Textbook pages 159-171

Thread

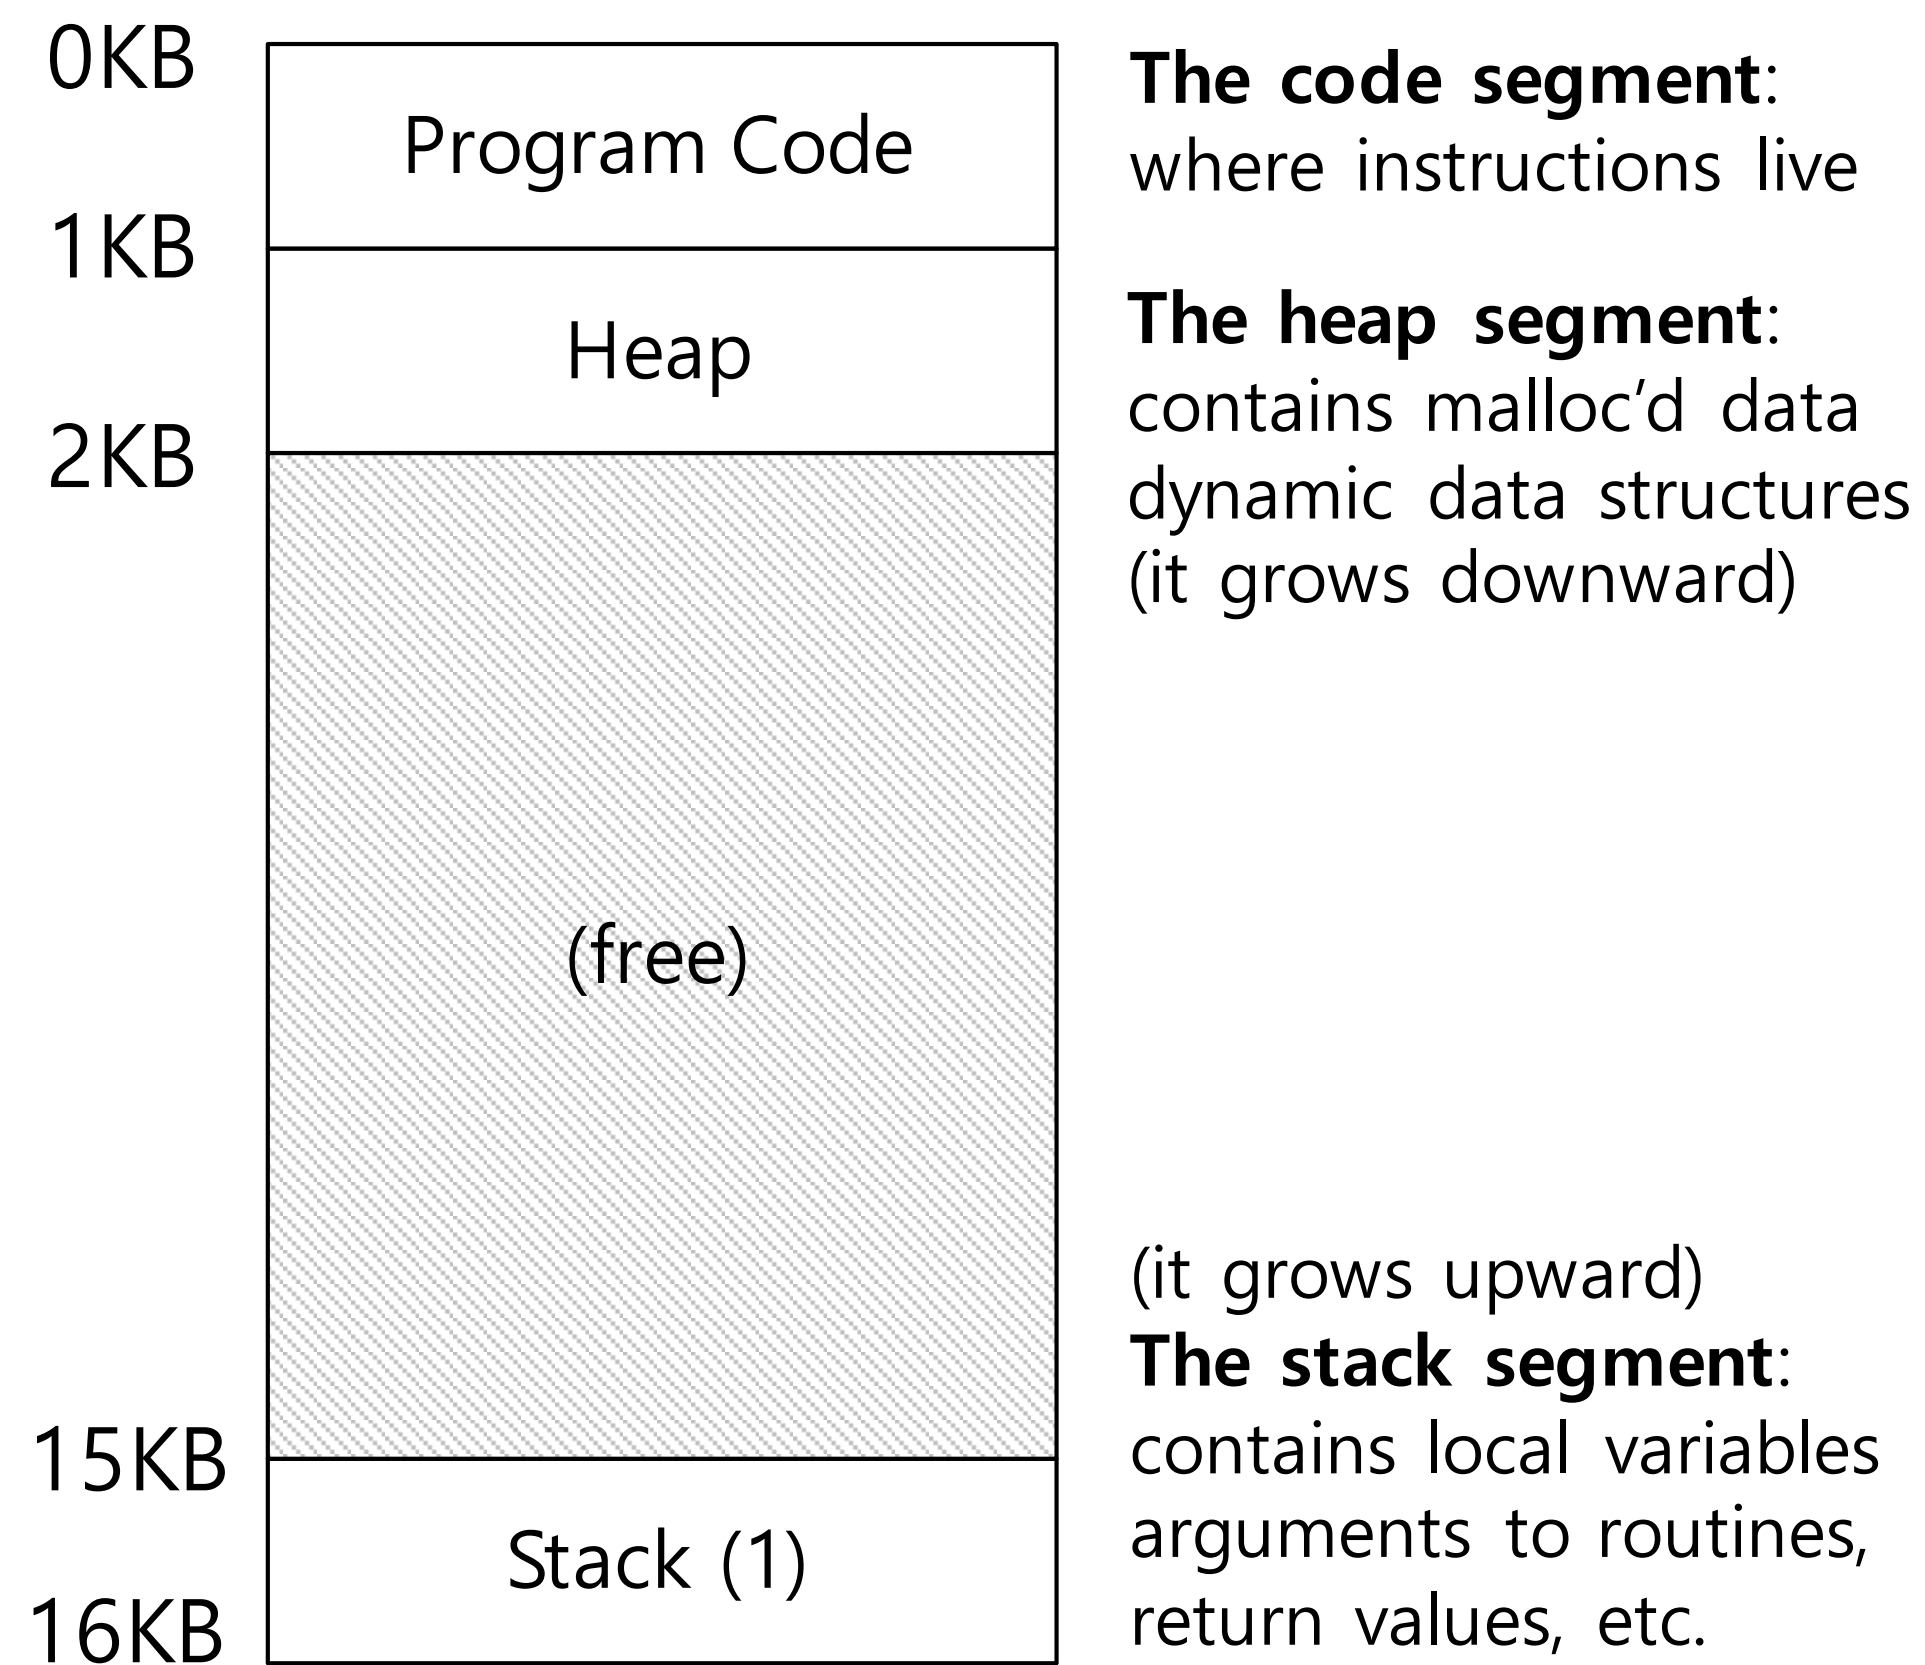
- New abstraction for **single running process**
- Multi-threaded program:
 - Multi-threaded program has more than one point of execution
 - Multiple program counters
 - They **share** the same **address space**

Context switch between threads


- Each thread has its own **program counter** and **set of registers**
 - One or more **thread control blocks (TCBs)** are needed to store the state of each thread
- When switching from running one (T1) to running the other (T2):
 - Save the register state of T1
 - Restore the register state of T2
 - Address space remains the same

The stack of the relevant thread

- There will be one stack per thread



**A Single-Threaded
Address Space**



**Two threaded
Address Space**

Race Condition

- Example with two threads
 - Assume we have a counter variable that holds the number 50
 - Each thread executes the line : `counter = counter + 1`
 - What do you think the result will be? Why?

Critical Section Problem

- A piece of code that accesses a shared variable and must not be concurrently executed by more than one thread
 - Multiple threads executing critical section can result in a race condition
 - Need to support atomicity for critical sections (mutual exclusion)
- **Goal 1 : carefully specify what a good solution to the critical section problem looks like**
- **Goal 2 : provide programming abstractions that let the user solve this problem**

Example Solution : Locks

- Ensure that any such critical section executes as if it were a single atomic instruction
 - Execute a series of instructions atomically
- lock_t mutex
lock(&mutex)
counter += 1
unlock(&mutex)

Thread Creation

- How to create an control threads?

```
#include <pthread.h>

int
pthread_create(      pthread_t*      thread,
                    const pthread_attr_t* attr,
                    void*             (*start_routine) (void*),
                    void*             arg);
```

- `thread`: Used to interact with this thread.
- `attr`: Used to specify any attributes this thread might have.
 - Stack size, Scheduling priority, ...
- `start_routine`: the function this thread start running in.
- `arg`: the argument to be passed to the function (`start_routine`)
 - *a void pointer* allows us to pass in *any type of* argument.

Thread Creation (Cont.)

- If start_routine instead required another type of argument, the declaration would look like this:

- An integer argument:

```
int
pthread_create(..., // first two args are the same
                  void*   (*start_routine) (int),
                  int     arg);
```

- Return an integer:

```
int
pthread_create(..., // first two args are the same
                  int    (*start_routine) (void*),
                  void*   arg);
```

Example : Creating a Thread

```
#include <pthread.h>

typedef struct __myarg_t {
    int a;
    int b;
} myarg_t;

void *mythread(void *arg) {
    myarg_t *m = (myarg_t *) arg;
    printf("%d %d\n", m->a, m->b);
    return NULL;
}

int main(int argc, char *argv[]) {
    pthread_t p;
    int rc;

    myarg_t args;
    args.a = 10;
    args.b = 20;
    rc = pthread_create(&p, NULL, mythread, &args);
    ...
}
```

Wait for a thread to complete

```
int pthread_join(pthread_t thread, void **value_ptr);
```

- `thread`: Specify which thread *to wait for*
- `value_ptr`: A pointer to the return value
 - Because `pthread_join()` routine changes the value, you need to **pass in a pointer** to that value.

Example : Waiting for Thread Completion

```
1  #include <stdio.h>
2  #include <pthread.h>
3  #include <assert.h>
4  #include <stdlib.h>
5
6  typedef struct __myarg_t {
7      int a;
8      int b;
9  } myarg_t;
10
11 typedef struct __myret_t {
12     int x;
13     int y;
14 } myret_t;
15
16 void *mythread(void *arg) {
17     myarg_t *m = (myarg_t *) arg;
18     printf("%d %d\n", m->a, m->b);
19     myret_t *r = malloc(sizeof(myret_t));
20     r->x = 1;
21     r->y = 2;
22     return (void *) r;
23 }
```

Example : Waiting for Thread Completion (Cont.)

```
25  int main(int argc, char *argv[]) {
26      int rc;
27      pthread_t p;
28      myret_t *m;
29
30      myarg_t args;
31      args.a = 10;
32      args.b = 20;
33      pthread_create(&p, NULL, mythread, &args);
34      pthread_join(p, (void **) &m); // this thread has been
                                     // waiting inside of the
// pthread_join() routine.
35      printf("returned %d %d\n", m->x, m->y);
36      return 0;
37 }
```

Example : Dangerous Code

- Be careful with **how values are returned** from a thread

```
1  void *mythread(void *arg) {
2      myarg_t *m = (myarg_t *) arg;
3      printf("%d %d\n", m->a, m->b);
4      myret_t r; // ALLOCATED ON STACK: BAD!
5      r.x = 1;
6      r.y = 2;
7      return (void *) &r;
8  }
```

Example : Simpler Argument Passing to a Thread

- Just passing a single value

```
1  void *mythread(void *arg) {
2      int m = (int) arg;
3      printf("%d\n", m);
4      return (void *) (arg + 1);
5  }
6
7  int main(int argc, char *argv[]) {
8      pthread_t p;
9      int rc, m;
10     pthread_create(&p, NULL, mythread, (void *) 100);
11     pthread_join(p, (void **) &m);
12     printf("returned %d\n", m);
13     return 0;
14 }
```

Locks

- Provide mutual exclusion to a critical section

- Interface

```
int pthread_mutex_lock(pthread_mutex_t *mutex);  
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

- Usage (w/o *lock initialization* and *error check*)

```
pthread_mutex_t lock;  
pthread_mutex_lock(&lock);  
x = x + 1; // or whatever your critical section is  
pthread_mutex_unlock(&lock);
```

- No other thread holds the lock → the thread will acquire the lock and **enter the critical section**.
- If another thread hold the lock → the thread will **not return from the call** until it has acquired the lock.

Locks (Cont.)

- All locks must be properly initialized

- One way: using `PTHREAD_MUTEX_INITIALIZER`

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
```

- The dynamic way: using `pthread_mutex_init()`

```
int rc = pthread_mutex_init(&lock, NULL);  
assert(rc == 0); // always check success!
```

Locks (Cont.)

- Check errors code when calling lock and unlock
 - An example wrapper

```
// Use this to keep your code clean but check for failures
// Only use if exiting program is OK upon failure
void Pthread_mutex_lock(pthread_mutex_t *mutex) {
    int rc = pthread_mutex_lock(mutex);
    assert(rc == 0);
}
```

- These two calls are used in lock acquisition

```
int pthread_mutex_trylock(pthread_mutex_t *mutex);
int pthread_mutex_timelock(pthread_mutex_t *mutex,
                           struct timespec *abs_timeout);
```

- trylock: return failure if the lock is already held
- timelock: return after a timeout

Compiling and Running

- To compile them, you must include the header `pthread.h`

```
prompt> gcc -o main main.c -Wall -pthread
```

- Explicitly link with the [pthreads library](#), by adding the `-pthread` flag.

```
man -k pthread
```

- There are a lot more pieces to pthreads. We'll see more abstractions and useful examples later

Reading

- For Thursday, read pages 257-275
 - Synchronization (critical sections, atomic operations, locks)