

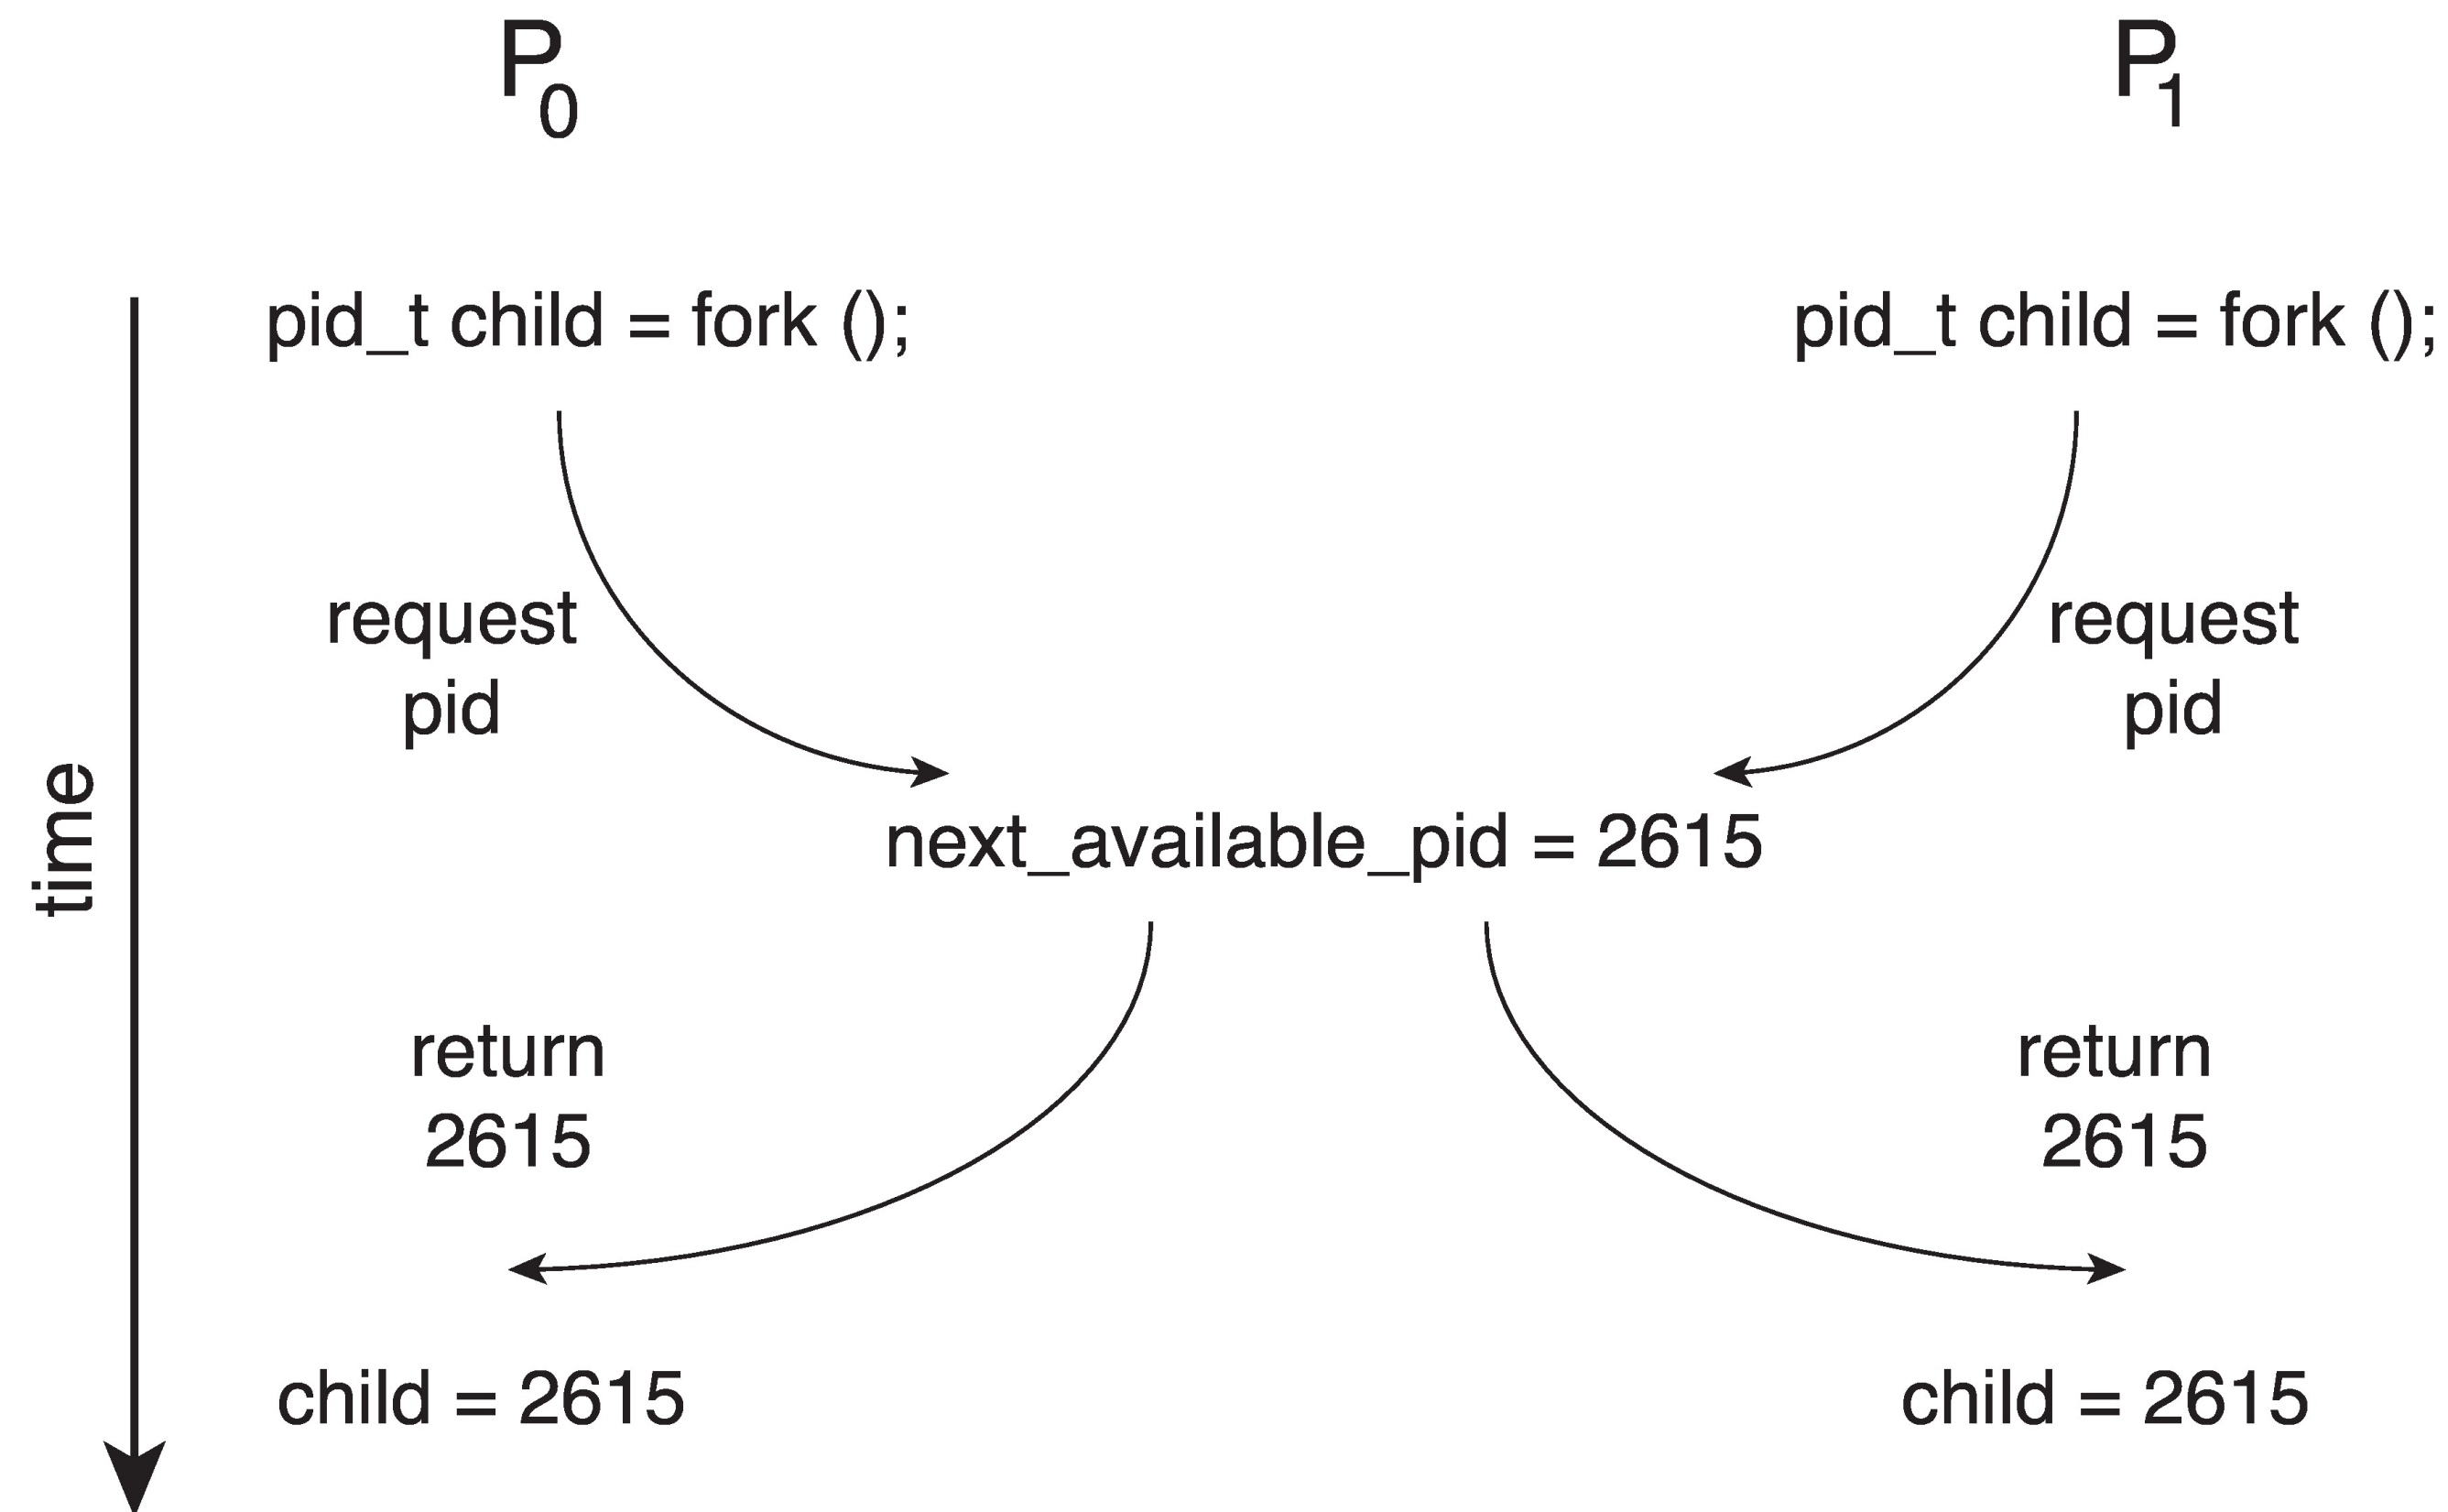
Concurrency : Critical Sections

03/09/2021

Professor Amanda Bienz

Textbook pages 257-275

Race Condition



- Unless there is a mechanism to prevent P_0 and P_1 from accessing the variable `next_available_pid`, the same pid could be assigned to two processes!

Critical Section Problem

- Consider system of n processes $\{P_0, P_1, \dots, P_{N-1}\}$
- Each process has **critical section** segment of code
 - Process may be changing common variables, updating table, writing file, etc
 - When one process in critical section, no other may be in its critical section
- **Critical section problem** : design protocol to solve this
- Each process must ask permission to enter critical section

Locks : The Basic Idea

- Ensure that any critical section executes as if it were a single **atomic instruction**

- An example: the canonical update of a shared variable

```
balance = balance + 1;
```

- Add some code around the critical section

```
1    lock_t mutex; // some globally-allocated lock 'mutex'  
2    ...  
3    lock(&mutex);  
4    balance = balance + 1;  
5    unlock(&mutex);
```

Locks : The Basic Idea

- Lock variable holds **the state of** the lock
- **Available** (unlocked, free) : no thread holds the lock
- **Aquired** (locked, held) : exactly one thread holds the lock and presumably is in a critical section

The semantics of the lock()

- lock()
 - Try to acquire the lock
 - If **no other thread holds** the lock, the thread will acquire the lock
 - Enter the critical section
 - This thread is said to be **the owner of the lock**
 - Other threads are prevented from entering the critical section while the first thread that holds the lock is in there

Pthread Locks - mutex

- The name that the POSIX library uses for a lock
- Used to provide mutual exclusion between threads

```
1  pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;  
2  
3  Pthread_mutex_lock(&lock); // wrapper for pthread_mutex_lock()  
4  balance = balance + 1;  
5  Pthread_mutex_unlock(&lock);
```

- We may be using different locks to protect different variables
 - Increase concurrency (a more fine-grained approach)

Building A Lock

- **Efficient locks** provide mutual exclusion at low cost
- Building a lock needs some help from hardware and the OS

Evaluating C.S. Solutions - Basic Criteria

- Correctness : Mutual exclusion and progress
 - **Mutual exclusion** : does the lock work, preventing multiple threads from entering a critical section?
 - **Progress** : if no one else is in the critical section, do we get in?
- Fairness : **Bounded Waiting**
 - If multiple threads are trying to get in, does each get in in a reasonable number of attempts? (Can a thread starve?)
- Also care about performance if we're going to do this often (and generally we are!)

Controlling Interrupts

- Disable interrupts for critical sections
 - One of the earliest solutions used to provide mutual exclusion
 - Invented for **single processor** systems

```
1  void lock() {  
2      DisableInterrupts();  
3  }  
4  void unlock() {  
5      EnableInterrupts();  
6  }
```

Controlling Interrupts : Problems

- Requires too much **trust** in applications
 - Greedy (or malicious) program could monopolize the processor
- Does not work on **multiprocessors**
- Code that masks or unmask interrupts is executed **slowing** by modern CPUs

Software Solution 1

- Two process solution
- Assume **load** and **store** instructions are ***atomic*** : cannot be interrupted
- The processes share the variable **int turn**
- The variable **turn** indicates whose turn it is to enter the critical section
- Initially, the value of **turn** is set to *i*

Algorithm for Process P_i

```
while (true) {  
  
    while (turn == j);  
  
    /* critical section */  
  
    turn = j;  
  
    /* remainder section */  
  
}
```

Does this work well?

- Mutual exclusion?
- Progress requirement?
- Bounded-waiting requirement?

Peterson's Solution

- Two process solution, share two variables:
 - **int turn**
 - **boolean flag[2]**
- Variable **turn** indicates whose turn it is to enter critical section
- The **flag** array is used to indicate if a process is ready to enter the critical section
 - $\text{flag}[i] = \text{true}$ implies P_i is ready

Algorithm for Process P_i

```
while (true) {  
  
    flag[i] = true;  
    turn = j;  
    while (flag[j] && turn == j)  
        ;  
  
    /* critical section */  
  
    flag[i] = false;  
  
    /* remainder section */  
  
}
```


Does this work well?

- Mutual exclusion?
- Progress requirement?
- Bounded-waiting requirement?

Why is hardware support needed?

- First attempt : using a flag denoting whether the lock is held or not
 - This code has problems

```
1  typedef struct __lock_t { int flag; } lock_t;
2
3  void init(lock_t *mutex) {
4  // 0 → lock is available, 1 → held
5  mutex->flag = 0;
6  }
7
8  void lock(lock_t *mutex) {
9  while (mutex->flag == 1) // TEST the flag
10     ; // spin-wait (do nothing)
11  mutex->flag = 1; // now SET it !
12  }
13
14 void unlock(lock_t *mutex) {
15  mutex->flag = 0;
16  }
```

Why is hardware support needed? (cont.)

- Problem 1 : No mutual exclusion (assume flag = 0 to begin)

Thread1	Thread2
<pre>call lock() while (flag == 1) interrupt: switch to Thread 2</pre>	<pre>call lock() while (flag == 1) flag = 1; interrupt: switch to Thread 1</pre>
<pre>flag = 1; // set flag to 1 (too!)</pre>	

- Problem 2 : Spin-waiting wastes time waiting for another thread
- So, we need an atomic instruction supported by Hardware!
 - Test-and-set instruction, also known as an atomic exchange

Test and Set (Atomic Exchange)

- An instruction to support the creation of simple locks

```
1  int TestAndSet(int *ptr, int new) {  
2  int old = *ptr;    // fetch old value at ptr  
3  *ptr = new;        // store 'new' into ptr  
4  return old;        // return the old value  
5  }
```

- Return(*testing*) old value pointed to by the ptr
- Simultaneously update(*setting*) said value to new
- This sequence of operations is **performed atomically**

A Simple Spin Lock using test-and-set

```
1  typedef struct __lock_t {
2  int flag;
3  } lock_t;
4
5  void init(lock_t *lock) {
6  // 0 indicates that lock is available,
7  // 1 that it is held
8  lock->flag = 0;
9  }
10
11 void lock(lock_t *lock) {
12 while (TestAndSet(&lock->flag, 1) == 1)
13     ;    // spin-wait
14 }
15
16 void unlock(lock_t *lock) {
17 lock->flag = 0;
18 }
```

- Note : to work correctly on a **single processor**, it requires a preemptive scheduler

Evaluating Spin Locks

- Correctness : yes
 - The spin lock only allows a single thread to enter the critical section
 - If no one is waiting, a thread will get in
- Fairness : no
 - Spin locks don't provide any fairness guarantees
 - Indeed, a thread spinning may spin forever
- Performance :
 - Single CPU : performance overheads can be quite painful
 - If number of threads roughly equals the number of CPUs, spin locks work reasonably well

Compare-And-Swap

- Test whether the value at the address(ptr) is equal to expected
 - If so, update the memory location pointed to by ptr with the new value
- In either case, return the actual value at that memory location

```
1  int CompareAndSwap(int *ptr, int expected, int new) {
2  int actual = *ptr;
3  if (actual == expected)
4      *ptr = new;
5  return actual;
6  }
```

Compare-and-Swap hardware atomic instruction (C-style)

```
1  void lock(lock_t *lock) {
2  while (CompareAndSwap(&lock->flag, 0, 1) == 1)
3      ; // spin
4  }
```

Spin lock with compare-and-swap

Load-Linked and Store-Conditional

```
1  int LoadLinked(int *ptr) {
2  return *ptr;
3  }
4
5  int StoreConditional(int *ptr, int value) {
6  if (no one has updated *ptr since the LoadLinked to this address) {
7      *ptr = value;
8      return 1; // success!
9  } else {
10     return 0; // failed to update
11 }
12 }
```

Load-linked And Store-conditional

- The store-conditional only succeeds if no intermittent store to the address has taken place
 - Success : return 1 and update the value at ptr to value
 - Fail : the value at ptr is not updated and 0 is returned

Load-Linked and Store-Conditional (Cont.)

```

1  void lock(lock_t *lock) {
2  while (1) {
3      while (LoadLinked(&lock->flag) == 1)
4          ; // spin until it's zero
5      if (StoreConditional(&lock->flag, 1) == 1)
6          return; // if set-it-to-1 was a success: all done
7                  otherwise: try it all over again
8  }
9  }
10
11 void unlock(lock_t *lock) {
12     lock->flag = 0;
13 }

```

Using LL/SC To Build A Lock

```

1  void lock(lock_t *lock) {
2  while (LoadLinked(&lock->flag) || !StoreConditional(&lock->flag, 1))
3      ; // spin
4  }

```

A more concise form of the lock () using LL/SC

Fetch-and-Add

- Atomically increment a value while returning the old value at a particular address

```
1  int FetchAndAdd(int *ptr) {  
2      int old = *ptr;  
3      *ptr = old + 1;  
4      return old;  
5  }
```

Fetch-And-Add Hardware atomic instruction (C-style)

Ticket Lock

- Ticket lock can be built with fetch-and-add
- Ensure progress for all threads : **fairness**

```
1  typedef struct __lock_t {
2      int ticket;
3      int turn;
4  } lock_t;
5
6  void lock_init(lock_t *lock) {
7      lock->ticket = 0;
8      lock->turn = 0;
9  }
10
11 void lock(lock_t *lock) {
12     int myturn = FetchAndAdd(&lock->ticket);
13     while (lock->turn != myturn)
14         ; // spin
15 }
16 void unlock(lock_t *lock) {
17     FetchAndAdd(&lock->turn);
18 }
```

So Much Spinning

- Hardware-based spin locks are simple and they work
- In some cases, these solution can be quite inefficient
- Any time a thread gets caught spinning, it wastes and entire time splice doing nothing but checking a value

How To Avoid *Spinning*?
We'll need OS Support too!

A Simple Approach : Just Yield

- When you are going to spin, give up the CPU to another thread
- OS system call moves the caller from the **running state** to the **ready state**
- Cost of a **context switch** can be substantial and the **starvation** problem exists

```
1  void init() {  
2      flag = 0;  
3  }  
4  
5  void lock() {  
6      while (TestAndSet(&flag, 1) == 1)  
7          yield(); // give up the CPU  
8  }  
9  
10 void unlock() {  
11     flag = 0;  
12 }
```

Lock with Test-and-set and Yield

Using Queues : Sleeping Instead of Spinning

- Queue to keep track of which threads are waiting to enter the lock
- `park()` : put a calling thread to sleep
- `unpark(threadID)` : wake a particular thread as designated the threadID

Using Queues : Sleeping Instead of Spinning

```
1  typedef struct __lock_t { int flag; int guard; queue_t *q; } lock_t;
2
3  void lock_init(lock_t *m) {
4      m->flag = 0;
5      m->guard = 0;
6      queue_init(m->q);
7  }
8
9  void lock(lock_t *m) {
10     while (TestAndSet(&m->guard, 1) == 1)
11         ; // acquire guard lock by spinning
12     if (m->flag == 0) {
13         m->flag = 1; // lock is acquired
14         m->guard = 0;
15     } else {
16         queue_add(m->q, gettid());
17         m->guard = 0;
18         park();
19     }
20 }
21 ...
```

Lock With Queues, Test-and-set, Yield, And Wakeup

Using Queues : Sleeping Instead of Spinning

```
• void unlock(lock_t *m) {  
•     while (TestAndSet(&m->guard, 1) == 1)  
•         ; // acquire guard lock by spinning  
•     if (queue_empty(m->q))  
•         m->flag = 0; // let go of lock; no one wants it  
•     else  
•         unpark(queue_remove(m->q)); // hold lock (for next thread!)  
•     m->guard = 0;  
• }
```

Lock With Queues, Test-and-set, Yield, And Wakeup (Cont.)

Wakeup / Waiting Race

- In case of releasing the lock (thread A) just before the call to park() (thread B) : thread B would potentially sleep forever
- Solaris solves this problem by adding a third system call : setpark()
 - By calling this routine, a thread can indicate it is about to park
 - If it happens to be interrupted and another thread calls unpack before park is actually called, the subsequent park returns immediately instead of sleeping

```
1      queue_add(m->q, gettid());  
2      setpark() ; // new code  
3      m->guard = 0;  
4      park();
```

Code modification inside of lock()

Futex

- Linux provides a futex (similar to Solaris's park and unpack)
 - `futex_wait(address, expected)`
 - Put calling thread to sleep
 - If value at address is not equal to expected, call returns immediately
 - `futex_wake(address)`
 - Wake one thread that is waiting on the queue

Futex (Cont.)

- Snippet from lowlevellock.h in the nptl library
 - The high bit of the integer v : track whether the lock is held or not
 - All other bits : number of waiters

```
1  void mutex_lock(int *mutex) {
2      int v;
3      /* Bit 31 was clear, we got the mutex (this is the fastpath) */
4      if (atomic_bit_test_set(mutex, 31) == 0)
5          return;
6      atomic_increment(mutex);
7      while (1) {
8          if (atomic_bit_test_set(mutex, 31) == 0) {
9              atomic_decrement(mutex);
10             return;
11         }
12         /* We have to wait now. First make sure the futex value
13            we are monitoring is truly negative (i.e. locked). */
14         v = *mutex;
15         ...
```

Futex (Cont.)

```
16         if (v >= 0)
17             continue;
18         futex_wait(mutex, v);
19     }
20 }
21
22 void mutex_unlock(int *mutex) {
23     /* Adding 0x80000000 to the counter results in 0 if and only if
24        there are not other interested threads */
25     if (atomic_add_zero(mutex, 0x80000000))
26         return;
27     /* There are other threads waiting for this mutex,
28        wake one of them up */
29     futex_wake(mutex);
30 }
```

Linux-based Futex Locks (Cont.)

Two-Phase Locks

- A two-phase lock realizes that spinning can be useful if the lock is about to be released
 - First phase :
 - Lock spins for awhile, hoping that it can acquire the lock
 - If the lock is not acquired during the first spin phase, a second spin phase is entered
 - Second phase :
 - The caller is put to sleep
 - The caller is only woken up when the lock becomes free later

Semaphores

- More sophisticated way to synchronize processes than mutex locks
- Semaphore **S** : integer variable
- Can only be accessed via two atomic operations
 - **wait()**
 - **signal()**

Wait and Signal Operations

- Definition of wait() operation:

wait(S)

{

while (S <= 0)

; // busy wait

S—;

}

- Definition of signal() operation:

signal(S)

{

S++;

}

Semaphore (Cont.)

- **Counting semaphore** : integer value can range over an unrestricted domain
- **Binary semaphore** : integer value can range only between 0 and 1
 - Same as a **mutex lock**
- Can implement a counting semaphore S as a binary semaphore
- With semaphores, we can solve various synchronization problems

Critical Sections and Semaphores

- Create a semaphore “**mutex**” initialized to 1

```
wait(mutex);  
Critical Section  
signal(mutex);
```

Semaphores and Synchronization

- Two processes : P1 and P2
- Two statements : S1 (by P1) and S2 (by P2)
- S1 must happen before S2

Semaphores and Synchronization

- Two processes : P1 and P2
- Two statements : S1 (by P1) and S2 (by P2)
- S1 must happen before S2

P1 :

S1;

signal(synch);

P2 :

wait(synch);

S2;

Semaphore Implementation

- Must guarantee that no two processes can execute the ***wait()*** and ***signal()*** on the same semaphore at the same time
- Thus, implementation becomes the critical section problem where **wait** and **signal** code are place in the critical section
- Could now have **busy waiting** in critical section implementation
 - But implementation code is short
 - Little busy waiting if critical section rarely occupied
- Note that applications may spend lots of time in critical sections and therefore this is not a good solution

Semaphore Implementation with no Busy Waiting

- With each semaphore, there is an associated waiting queue
- Each entry in waiting queue has two data items:
 - Value (integer)
 - Pointer to next record in list
- Two operations:
 - **Block** : place process invoking the operating on the appropriate waiting queue
 - **Wakeup** : remove one of the processes in the waiting queue and place it in the ready queue

Implementation with no Busy Waiting

- Waiting queue
typedef struct
{
 int value;
 struct process * list;
} semaphore;

Implementation with no Busy Waiting

- *Wait (semaphore * S)*
{
 S->value--;
 if (S->value < 0)
 {
 add this process to S->list;
 block();
 }
}

Implementation with no Busy Waiting

- *Signal (semaphore* S)*
{
 S->value++;
 if (S->value <= 0)
 {
 remove a process P from S->list;
 wakeup(P);
 }
}

Problems with Semaphores

- Incorrect use of semaphore operations:
 - `signal(mutex) ... wait(mutex)`
 - `wait(mutex) ... wait(mutex)`
 - Omitting `wait(mutex)` and or `single(mutex)`

Reading

- Deadlocks : pg 283
- Synchronization problems : pg 289-294
- POSIX : pg 299-303
- Alternative approaches : pg 311-313