

File System Interface

04/22/2021

Professor Amanda Bienz

Persistent Storage

- Keep a data intact, even if there is **power loss**
 - Hard disk drive
 - Solid-state storage device
- **Two key abstractions in the virtualization of storage**
 - File
 - Directory


File

- A linear array of bytes
- Each file has low-level name as inode number
 - User is not aware of this name
- File system has a responsibility to store data persistently on disk

Directory

- Directory is like a file, also has a low-level name
 - It contains a list of (user-readable name, low-level name) pairs
 - Each entry in a directory refers to either files or other directories
- Example :
 - A directory has an entry (“foo”, “10”)
 - A file “foo” with the low level name “10”

Directory Tree (Directory Hierarchy)



An Example Directory Tree

Valid files (absolute pathname) :

/foo/bar.txt
/bar/foo/bar.txt

Valid directory :

/
/foo
/bar
/bar/bar
/bar/foo/

} Sub-directories

File Operations

- **Create**
- **Write** : at write pointer location
- **Read** : at read pointer location
- **Reposition within file (seek)**
- **Delete**
- **Truncate**
- **Open(Fi)** : search directory structure on disk for entry Fi and move content of entry to memory
- **Close(Fi)** : move content of entry Fi in memory to directory structure on disk

Creating Files

- Use `open()` system call with `O_CREAT` flag

```
int fd = open("foo", O_CREAT | O_WRONLY | O_TRUNC);
```

- `O_CREAT` : create file
- `O_WRONLY` : only write to that file while opened
- `O_TRUNC` : make the file size zero (remove any existing content)
- `open()` system call returns **file descriptor**
 - An integer, used to access files

Reading and Writing Files

- An example of reading and writing 'foo' file

```
prompt> echo hello > foo  
prompt> cat foo  
hello  
prompt>
```

- Echo : redirect the output of echo to the file foo
- Cat : dump the contents of a file to the screen

How does the cat program access the file foo ?

We can use `strace` to trace the system calls made by a program.

Reading and Writing Files (Cont.)

```
prompt> strace cat foo
...
open("foo", O_RDONLY|O_LARGEFILE) = 3
read(3, "hello\n", 4096)          = 6
write(1, "hello\n", 6)             = 6 // file descriptor 1: standard out
hello
read(3, "", 4096)                  = 0 // 0: no bytes left in the file
close(3)                           = 0
...
prompt>
```

- Open (file descriptor, flags)
 - Return file descriptor (3 in example)
 - File descriptor 0, 1, 2 is for standard input/output/error
- read(file descriptor, buffer pointer, the size of the buffer)
 - Return the number of bytes it read
- write(file descriptor, buffer pointer, the size of the buffer)
 - Return the number of bytes it wrote

Reading and Writing Files (Cont.)

- Writing a file (A similar set of read steps)
 - A file is opened for writing (`open()`)
 - The `write()` system call is called
 - Repeatedly called for larger files
- `close()`

Reading and Writing, But Not Sequentially

- An open file has a current offset
 - Determine where the next read or write will begin reading from or writing to within the file
- Update the current offset
 - Implicitly : a read or write of N bytes takes place, N is added to the current offset
 - Explicitly : `lseek()`

Reading and Writing, But Not Sequentially (Cont.)

```
off_t lseek(int fildes, off_t offset, int whence);
```

- Fildes : file descriptor
- Offset : position the file offset to a particular location within the file
- Whence : Determine how the seek is performed

If whence is SEEK_SET, the offset is set to offset bytes.

If whence is SEEK_CUR, the offset is set to its current location plus offset bytes.

If whence is SEEK_END, the offset is set to the size of the file plus offset bytes.

Writing Immediately with `fsync()`

- The file system will buffer writes in memory for some time
 - Example : 5 seconds, or 30
 - Performance reasons
- At that later point in time, the write(s) will actually be issued to the storage device
 - Writes seem to complete quickly
 - Data can be lost (e.g. the machine crashes)

Writing Immediately with `fsync()` (Cont.)

- However, some applications require more than eventual guarantee
 - Example : database management system requires force writes to disk from time to time
- `off_t fsync(int fd)`
 - Filesystem forces all dirty (i.e. not yet written) data to disk for the file referred to by the file descriptor
 - `fsync()` returns once all of these writes are complete

Write Immediately with fsync() (Cont.)

- An example of fsync()

```
int fd = open("foo", O_CREAT | O_WRONLY | O_TRUNC);  
assert (fd > -1)  
int rc = write(fd, buffer, size);  
assert (rc == size);  
rc = fsync(fd);  
assert (rc == 0);
```

- In some cases, this code needs to fsync() the directory that contains the file foo

Renaming Files

- `rename(char* old, char* new)`
 - Rename a file to a different name
 - It is implemented as an **atomic call**
 - Ex : change from foo to bar

```
prompt> mv foo bar    // mv uses the system call rename()
```

- How to update a file atomically:

```
int fint fd = open("foo.txt.tmp", O_WRONLY|O_CREAT|O_TRUNC);  
write(fd, buffer, size); // write out new version of file  
fsync(fd);  
close(fd);  
rename("foo.txt.tmp", "foo.txt");
```


Getting Information About Files

- `stat()`, `fstat()` : Show the file metadata
 - Metadata is information about each file
 - Ex : size, low-level name, permission, ...

- Stat structure:

```
struct stat {  
    dev_t st_dev;      /* ID of device containing file */  
    ino_t st_ino;      /* inode number */  
    mode_t st_mode;    /* protection */  
    nlink_t st_nlink;  /* number of hard links */  
    uid_t st_uid;      /* user ID of owner */  
    gid_t st_gid;      /* group ID of owner */  
    dev_t st_rdev;     /* device ID (if special file) */  
    off_t st_size;     /* total size, in bytes */  
    blksize_t st_blksize; /* blocksize for filesystem I/O */  
    blkcnt_t st_blocks; /* number of blocks allocated */  
    time_t st_atime;    /* time of last access */  
    time_t st_mtime;    /* time of last modification */  
    time_t st_ctime;    /* time of last status change */  
};
```

Getting Information About Files (Cont.)

- To see stat information, you can use the command line tool `stat`

```
prompt> echo hello > file
prompt> stat file

File: `file'
Size: 6 Blocks: 8 IO Block: 4096 regular file
Device: 811h/2065d Inode: 67158084 Links: 1
Access: (0640/-rw-r-----) Uid: (30686/ root) Gid: (30686/ remzi)
Access: 2011-05-03 15:50:20.157594748 -0500
Modify: 2011-05-03 15:50:20.157594748 -0500
Change: 2011-05-03 15:50:20.157594748 -0500
```

- File system keeps this type of information in a `inode` structure.

Removing Files

- Rm is a Linux command to remove a file
- Calls `unlink()` system call to remove the file

```
prompt> strace rm foo
...
unlink("foo")          = 0  // return 0 upon success
...
prompt>
```

Why it calls `unlink()`? not "remove or delete"
We can get the answer later.

File Systems Reading

- Chapter 14 : Pg 563 - 589