

# File System Implementation

04/26/2021


Professor Amanda Bienz

# The Way to Think

- Two different aspects to implement file systems
  - Data structures:
    - What types of on-disk structures are utilized by the file system to organize its data and metadata?
  - Access Methods:
    - How does it map the calls made by a process as `open()`, `read()`, `write()`, etc
    - Which structures are read during the execution of a particular system call?

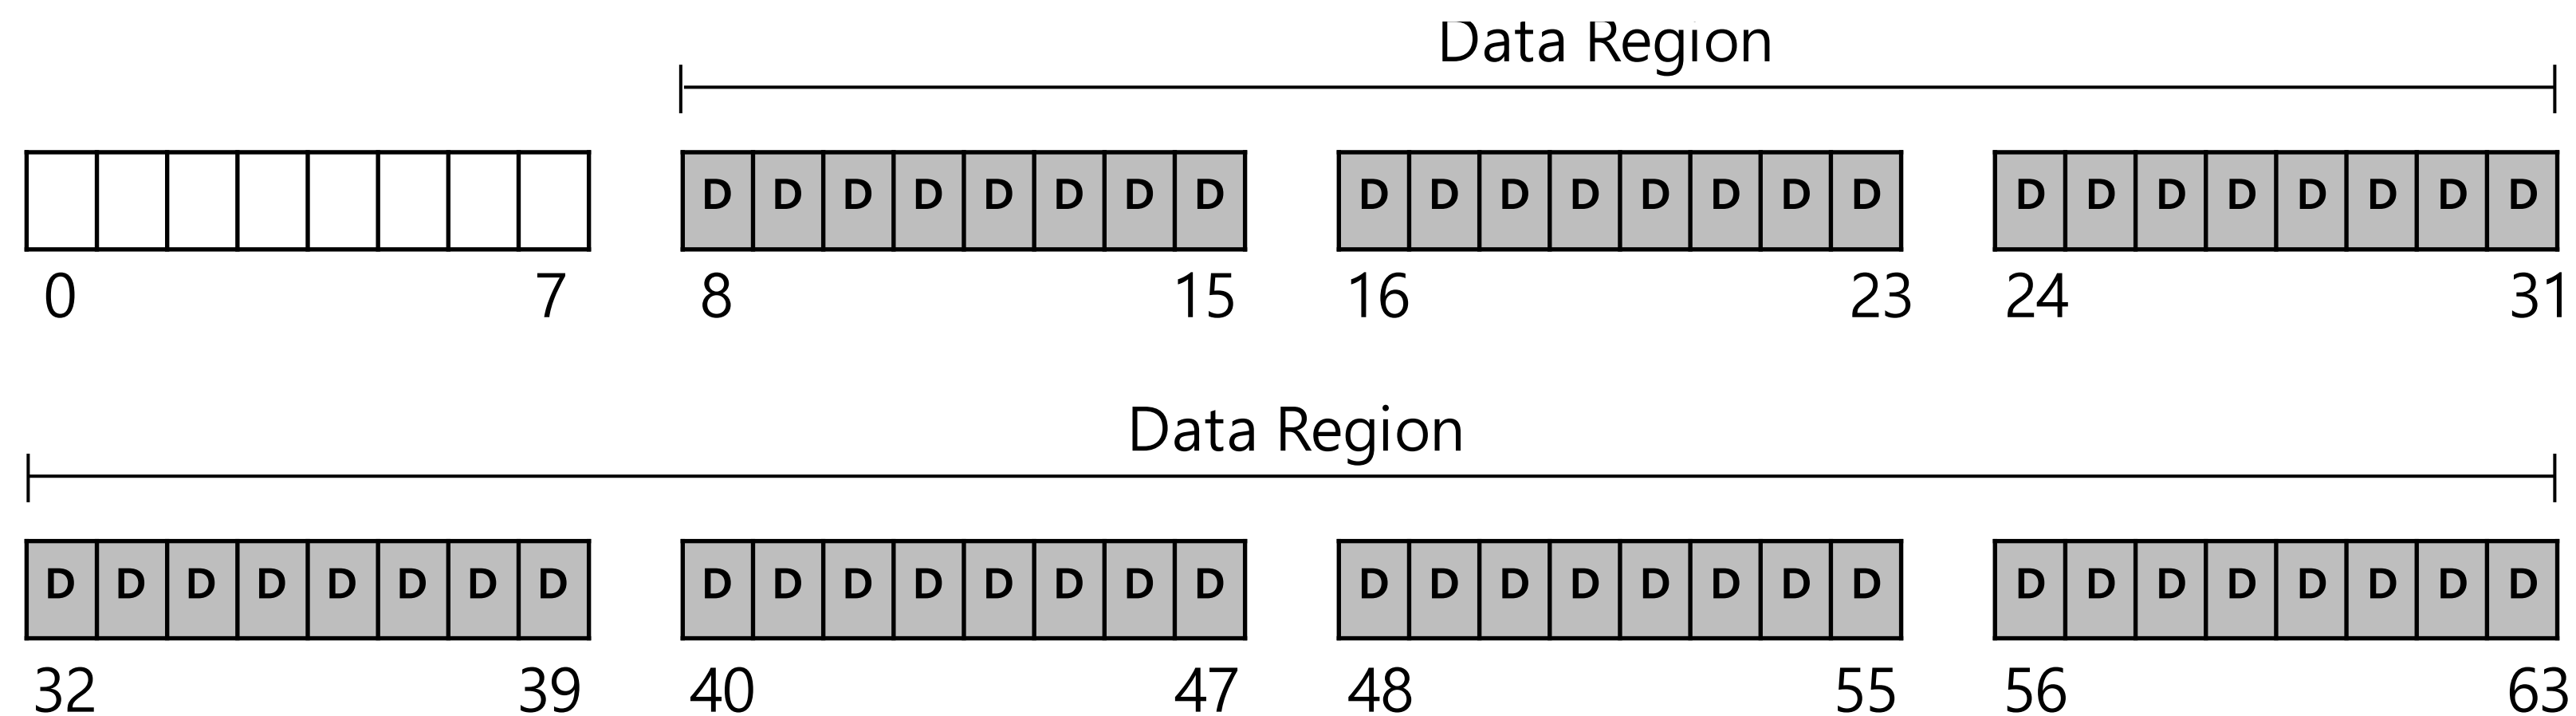
# Overall Organization

- Let's develop the overall organization of the file system data structure
- Divide the disk into **blocks**
  - Block size is 4KB, blocks addressed from 0 to N-1



# Data region in file system

- Reserve **data region** to store user

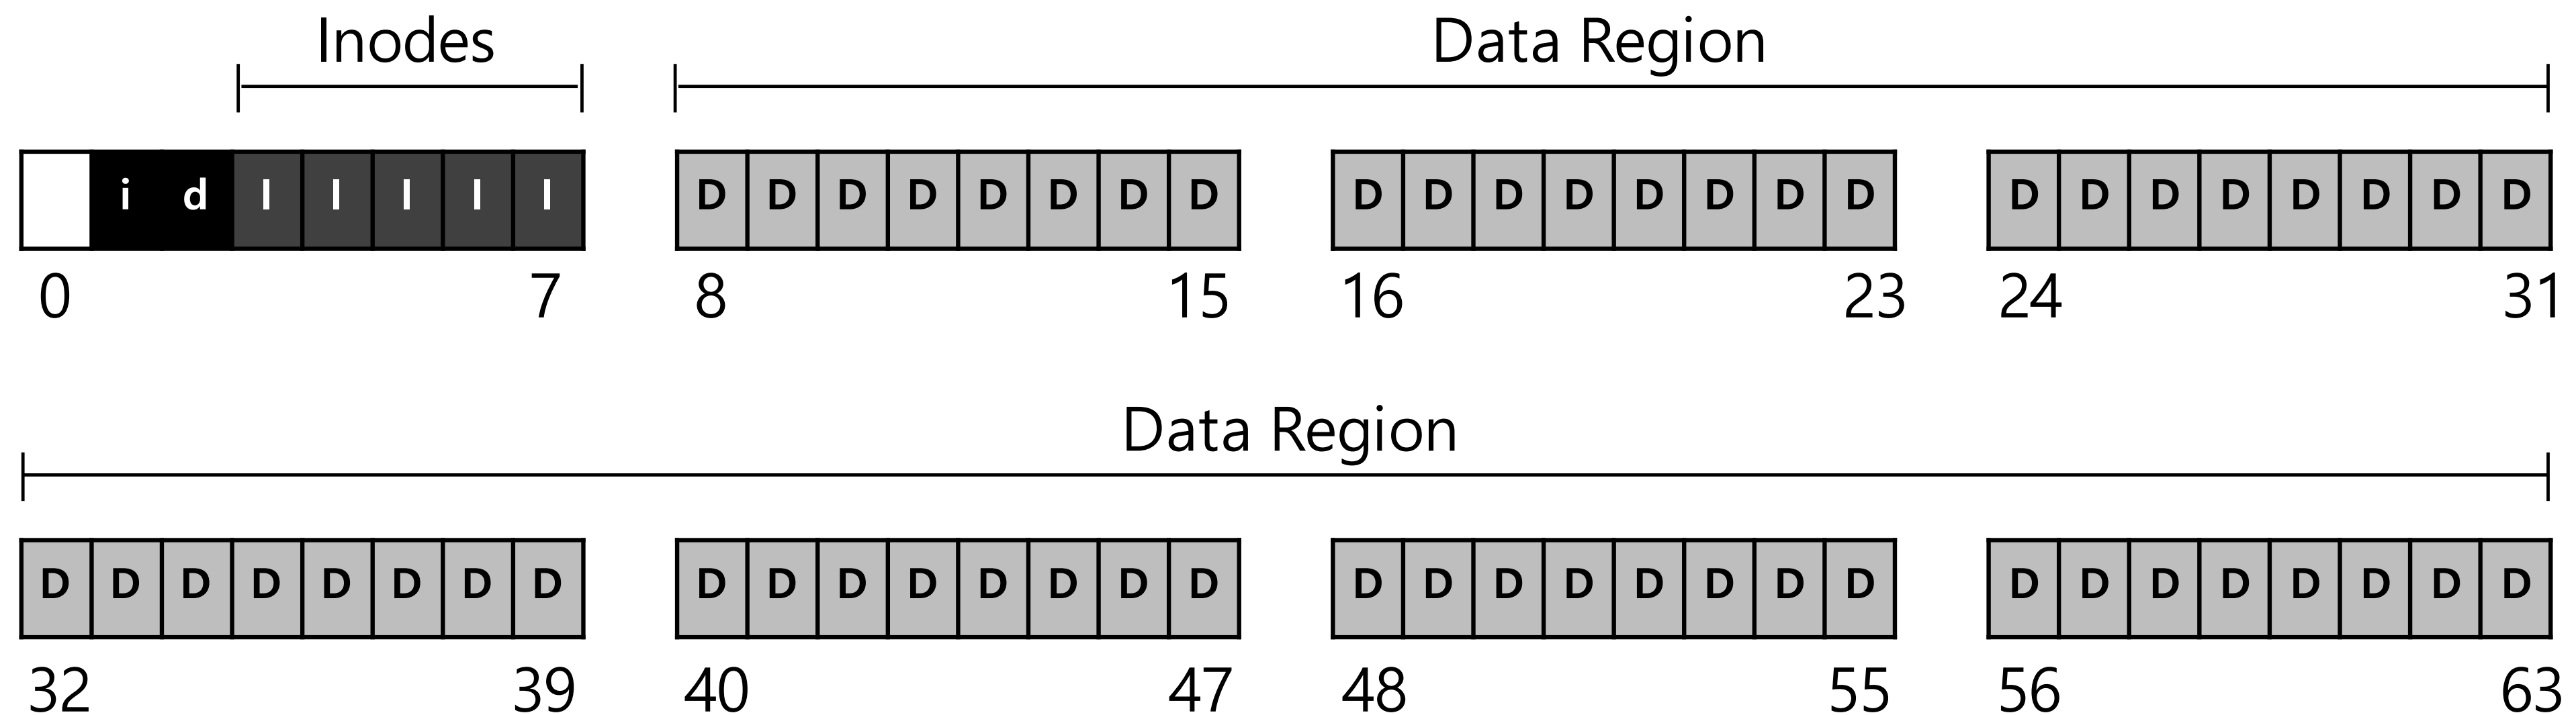


- File system has to track which data block comprise a file, the size of the file, its owner, etc.

How we store these **inodes** in file system?

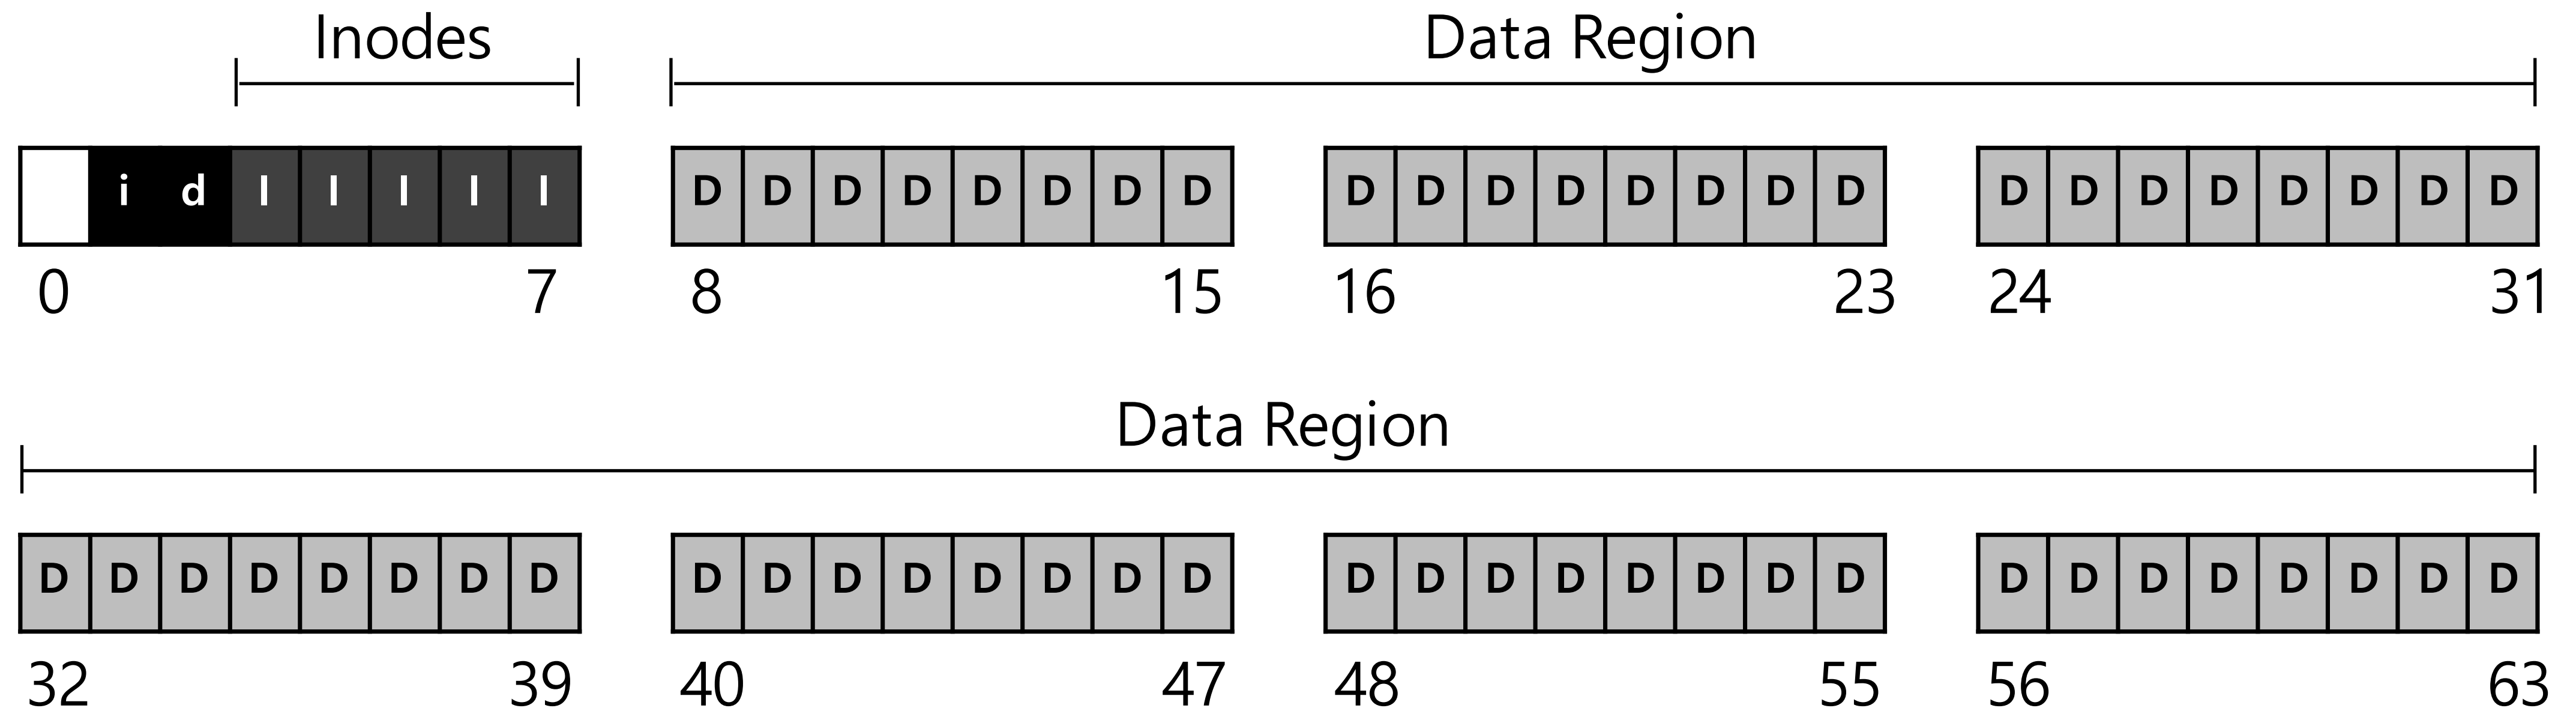
# Inode table in file system

- Reserve some space for **inode table**
  - This holds an array of on-disk inodes
  - Ex) inode tables: 3 ~ 7, inode size: 256 bytes
    - 4KB block can hold 16 inodes
  - The filesystem contains 80 inodes (max number of files)



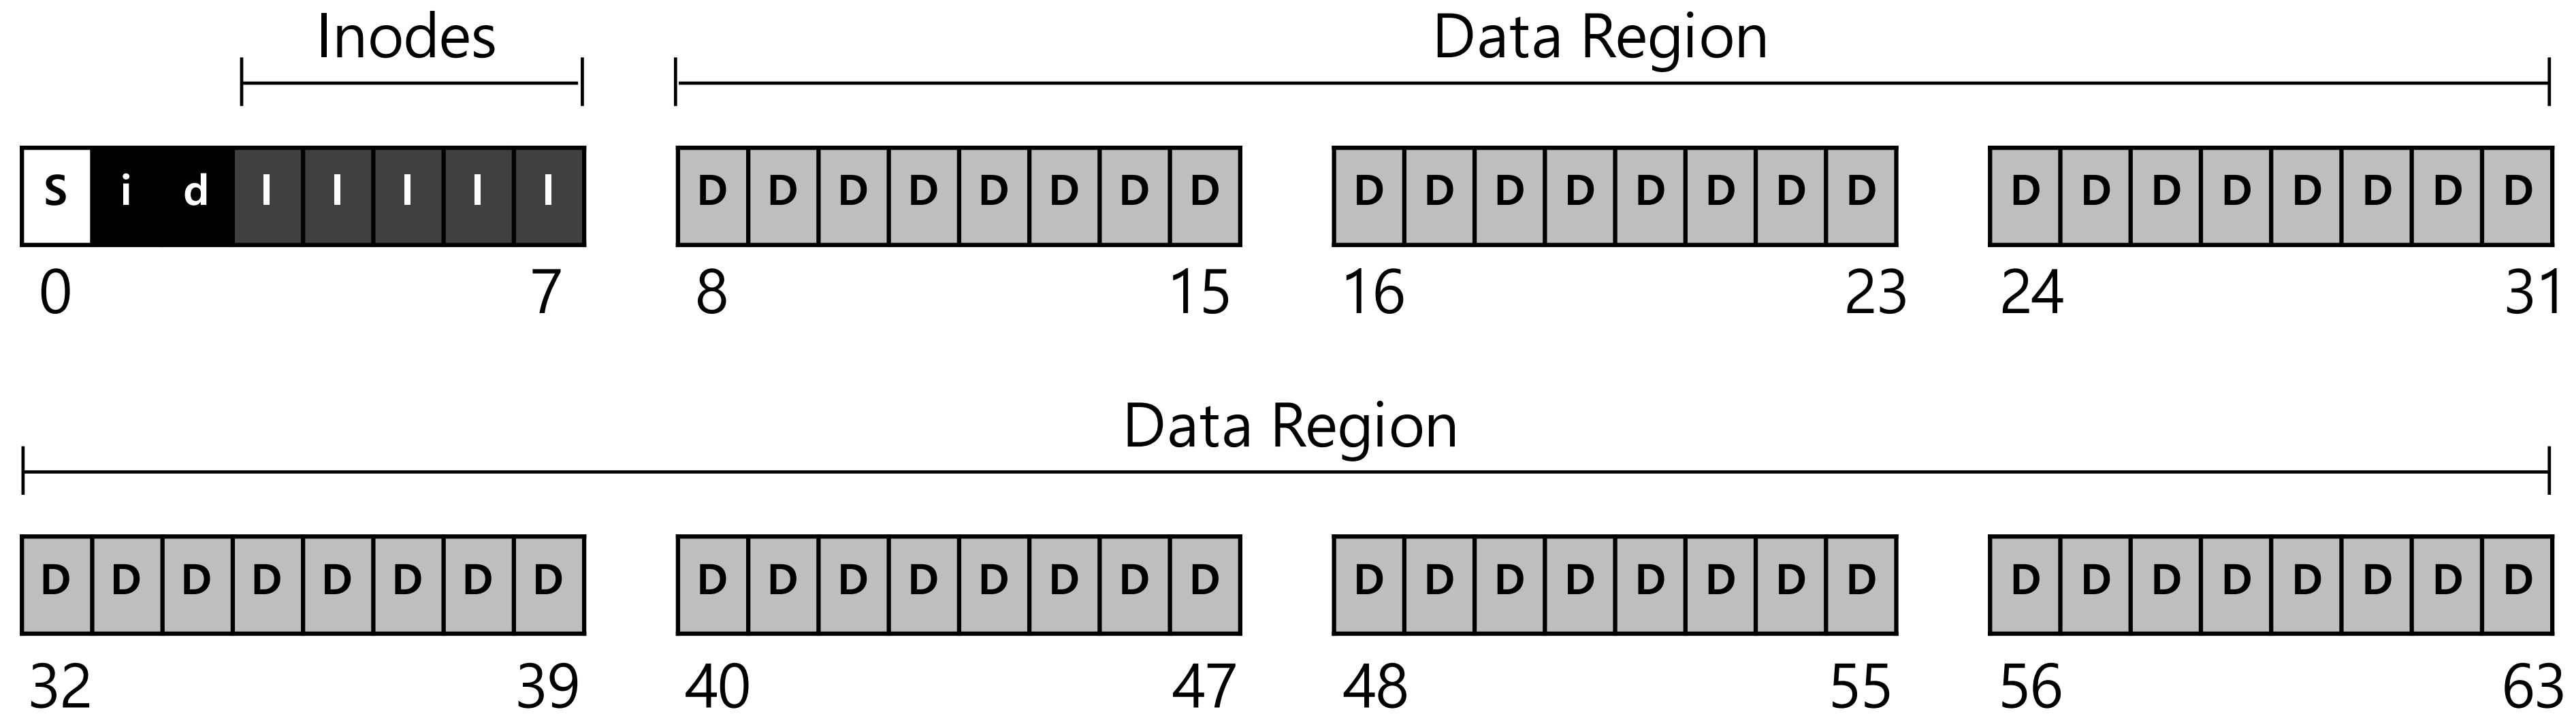
# Allocation Structures

- This is to track whether inodes or data blocks are free or allocated
- Use **bitmap**, each bit indicates free(0) or in-use(1)
  - **Data bitmap:** for data region
  - **Inode bitmap:** for inode table



# Superblock


- Super block contains this **information** for **particular file system**
- Ex ) The number of inodes, begin location of node table



- Thus, when mounting a filesystem, OS will read the superblock first, to initialize various information

# File Organization: The node

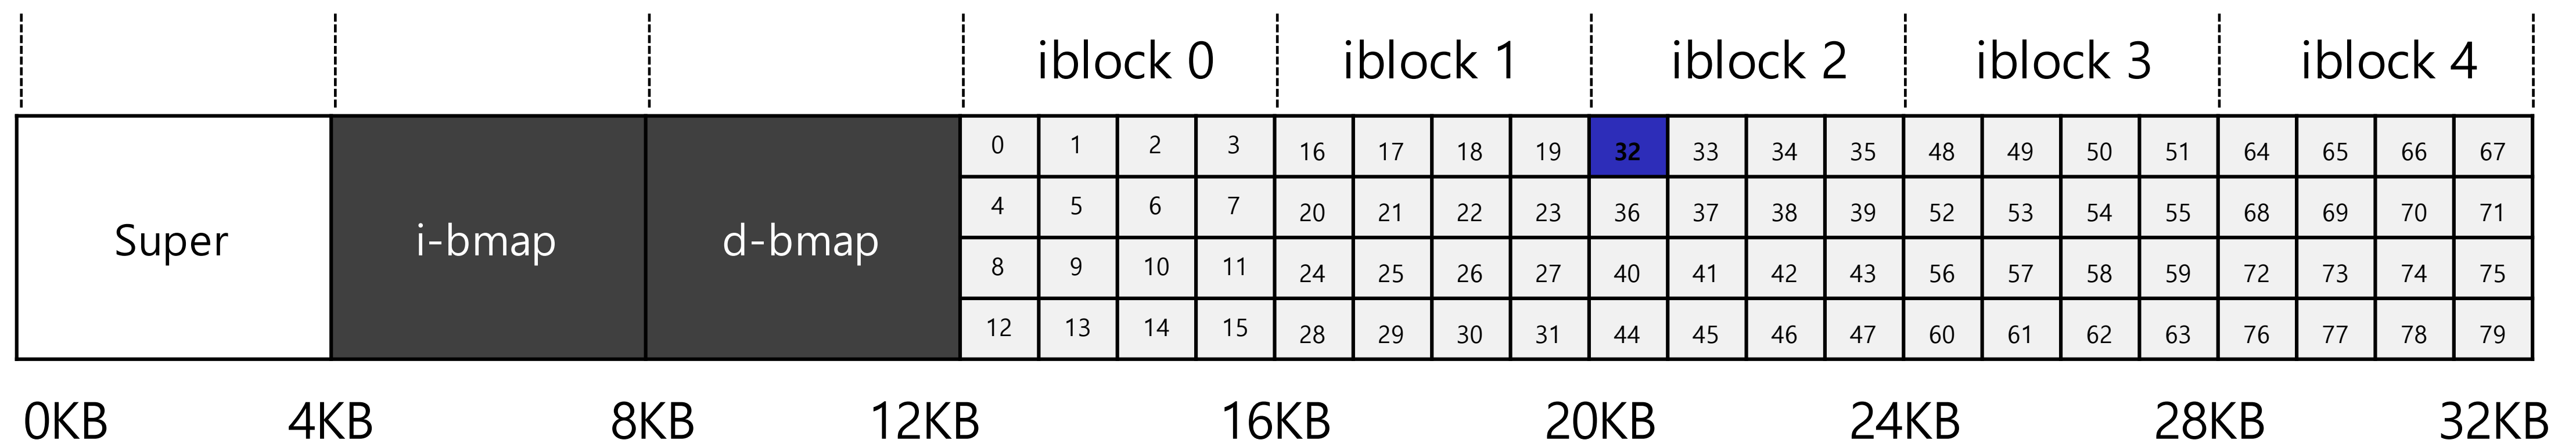
- Each inode is referred to by node number
  - By node number, filesystem calculates where the node is on the disk
  - Ex) inode number : 32
    - Calculate the offset into the node region ( $32 * \text{sizeof}(\text{inode})$ )  
 $(32 * 256) = 8192$
    - Add start address of node table (12KB) + node region (8KB) = 20KB





# File Organization : The inode (Cont.)

- Disks are not byte addressable, sector addressable
- Disks consist of a large number of addressable sectors (512 bytes)
  - Ex) Fetch the block of inode (inode number: 32)
    - Sector address **iaddr** of the inode block:
      - **blk : (inumber \* sizeof(inode)) / blocksize**
      - **Sector : ((blk \* block size) + inodeStartAddr) / sector size**



# File Organization : The inode (Cont.)

- inode have all of the information about a file
  - File type (regular file, directory, etc.)
  - Size, the number of blocks allocated to it
  - Protection information (who owns the file, who can access it, etc)
  - Time information
  - Etc.

# File Organization : The inode (Cont.)

| Size | Name        | What is this inode field for?                     |
|------|-------------|---|
| 2    | mode        | can this file be read/written/executed?           |
| 2    | uid         | who owns this file?                               |
| 4    | size        | how many bytes are in this file?                  |
| 4    | time        | what time was this file last accessed?            |
| 4    | ctime       | what time was this file created?                  |
| 4    | mtime       | what time was this file last modified?            |
| 4    | dtime       | what time was this inode deleted?                 |
| 4    | gid         | which group does this file belong to?             |
| 2    | links_count | how many hard links are there to this file?       |
| 2    | blocks      | how many blocks have been allocated to this file? |
| 4    | flags       | how should ext2 use this inode?                   |
| 4    | osd1        | an OS-dependent field                             |
| 60   | block       | a set of disk pointers (15 total)                 |
| 4    | generation  | file version (used by NFS)                        |
| 4    | file_acl    | a new permissions model beyond mode bits          |
| 4    | dir_acl     | called access control lists                       |
| 4    | faddr       | an unsupported field                              |
| 12   | i_osd2      | another OS-dependent field                        |

The EXT2 Inode

# The Multi-Level Index

- To support bigger files, we use multi-level index
- **Indirect pointer** points to a block that contains more pointers
  - Inodes have fixed number of direct pointers (12) and a single indirect pointer
  - If a file grows large enough, an indirect block is allocated, inode's slot for an indirect pointer is set to point to it
    - $(12 + 1024) * 4K = 4144KB$

# The Multi-Level Index (Cont.)

- **Direct indirect pointer** points to a block that contains indirect blocks
  - Allow file to grow with an additional 1024x1024 or 1million 4KB blocks
- **Triple indirect pointer** points to a block that contains double indirect blocks
- Multi-level index approach to pointing to file blocks
  - Ex) 12 direct pointers, a single and double indirect block
    - Over 4GB in size  $(12+1024+1024^2)*4KB$
- Many file systems use a multi-level index
  - Linux EXT2, EXT3, NetApp's WAFL, Unix file system
  - Linux EXT4 use **extents** instead of simple pointers

# The Multi-Level Index (Cont.)

**Most files are small**

**Average file size is growing**

**Most bytes are stored in large files**

**File systems contains lots of files**

**File systems are roughly half full**

**Directories are typically small**

**Roughly 2K is the most common size**

**Almost 200K is the average**

**A few big files use most of the space**

**Almost 100K on average**

**Even as disks grow, file system remain -50% full**

**Many have few entries; most have 20 or fewer**

**File System Measurement Summary**

# Directory Organization

- Directory contains a list of (entry name, inode number) pairs
- Each directory has two extra files “.” for current directory and “..” for parent directory
- For example, dir has three files (foo, bar, foobar)

| <b>inum</b> | <b> </b> | <b>reclen</b> | <b> </b> | <b>strlen</b> | <b> </b> | <b>name</b> |
|-------------|----------|---------------|----------|---------------|----------|-------------|
| 5           |          | 4             |          | 2             |          | .           |
| 2           |          | 4             |          | 3             |          | ..          |
| 12          |          | 4             |          | 4             |          | foo         |
| 13          |          | 4             |          | 4             |          | bar         |
| 24          |          | 8             |          | 7             |          | foobar      |

**on-disk for dir**

# Free Space Management

- File system track which inode and data block are free or not
- In order to manage free space, we have two simple bitmaps
  - When file is newly created, its allocated an inode by searching the inode bitmap and update on-disk bitmap
  - Pre-allocation policy is commonly used for allocating contiguous blocks



# Access Paths: Reading a File From Disk

- Issue an open (“/foo/bar”, O\_RDONLY)
  - Traverse pathname and thus locate desired inode
  - Begin at the root (/)... in most Unix file systems, root inode number is 2
  - Filesystem reads in the block that contains inode number 2
  - Look inside of it to find pointer to data blocks (contents of root)
  - By reading in one or more directory data blocks, it will find “foo”
  - Traverse recursively the path name until the desired inode (“bar”)
  - Check final permissions, allocate a file descriptor for this process and return file descriptor to user

# Access Paths: Reading a File From Disk (Cont.)

- **Issue read() to read from the file**
  - Read in first block of the file, consulting inode to find the location of such a block
    - Update the inode with a new last accessed time
    - Update in-memory open file table for file descriptor, file offset
- **When file is closed:**
  - File descriptor should be deallocated, but for now, that is all the file system really needs to do. No disk I/Os take place.

# Access Paths: Reading a File From Disk (Cont.)

|           | data<br>bitmap | inode<br>bitmap | root<br>inode | foo<br>inode | bar<br>inode | root<br>data | foo<br>data | bar<br>data[0] | bar<br>data[1] | bar<br>data[2] |
|-----------|----------------|-----------------|---------------|--------------|--------------|--------------|-------------|----------------|----------------|----------------|
| open(bar) |                |                 | read          | read         | read         | read         | read        |                |                |                |
| read()    |                |                 |               |              | read         |              |             | read           |                |                |
| read()    |                |                 |               |              | read         |              |             |                | read           |                |
| read()    |                |                 |               |              | read         |              |             |                |                | read           |

File Read Timeline (Time Increasing Downward)

# Access Paths: Writing to Disk

- Issue **write()** to update the file with new contents
- File may allocate a block (unless the block is being overwritten)
  - Need to update data block, data bitmap
- It generates five I/Os:
  - One to read the data bitmap
  - One to write the bitmap (to reflect new state to disk)
  - Two more to read and then write the inode
  - One to write the actual block itself
- To create file, it also allocates space for directory, causing high I/O traffic

# Access Paths: Writing to Disk (Cont.)

|                      | data<br>bitmap | inode<br>bitmap | root<br>inode | foo<br>inode | bar<br>inode  | root<br>data | foo<br>data   | bar<br>data[0] | bar<br>data[1] | bar<br>data[2] |
|----------------------|----------------|-----------------|---------------|--------------|---------------|--------------|---------------|----------------|----------------|----------------|
| create<br>(/foo/bar) |                | read<br>write   | read          | read         |               | read         | read<br>write |                |                |                |
| write()              | read<br>write  |                 |               | write        | read<br>write |              |               | write          |                |                |
| write()              | read<br>write  |                 |               |              | read<br>write |              |               | write          |                |                |
| write()              | read<br>write  |                 |               |              | read<br>write |              |               |                |                | write          |

File Creation Timeline (Time Increasing Downward)

# Caching and Buffering

- Reading and writing files are expensive, incurring many I/Os
  - For example, long pathname (/1/2/3/.../100/file.txt)
    - One to read the inode of the directory and at least one to read its data
    - Literally perform hundreds of reads just to open the file
- In order to reduce I/O traffic, file systems aggressively use system memory (DRAM) to cache
  - Early file system use fixed-size cache to hold popular blocks
    - Static partitioning of memory can be wasteful
  - Modern systems use dynamic partitioning approach, unified page cache
- Read I/O can be avoided by large cache

# Caching and Buffering (Cont.)

- Write traffic has to go to disk for persistent. Thus, cache does not reduce write I/Os
- File system use write buffering for write performance benefits
  - Delaying writes (file system batch some updates into a smaller set of I/Os)
  - By buffering a number of writes in memory, the file system can then schedule the subsequent I/Os
  - By avoiding writes
- Some applications force flush data to disk by calling `fsync()` or direct I/O