

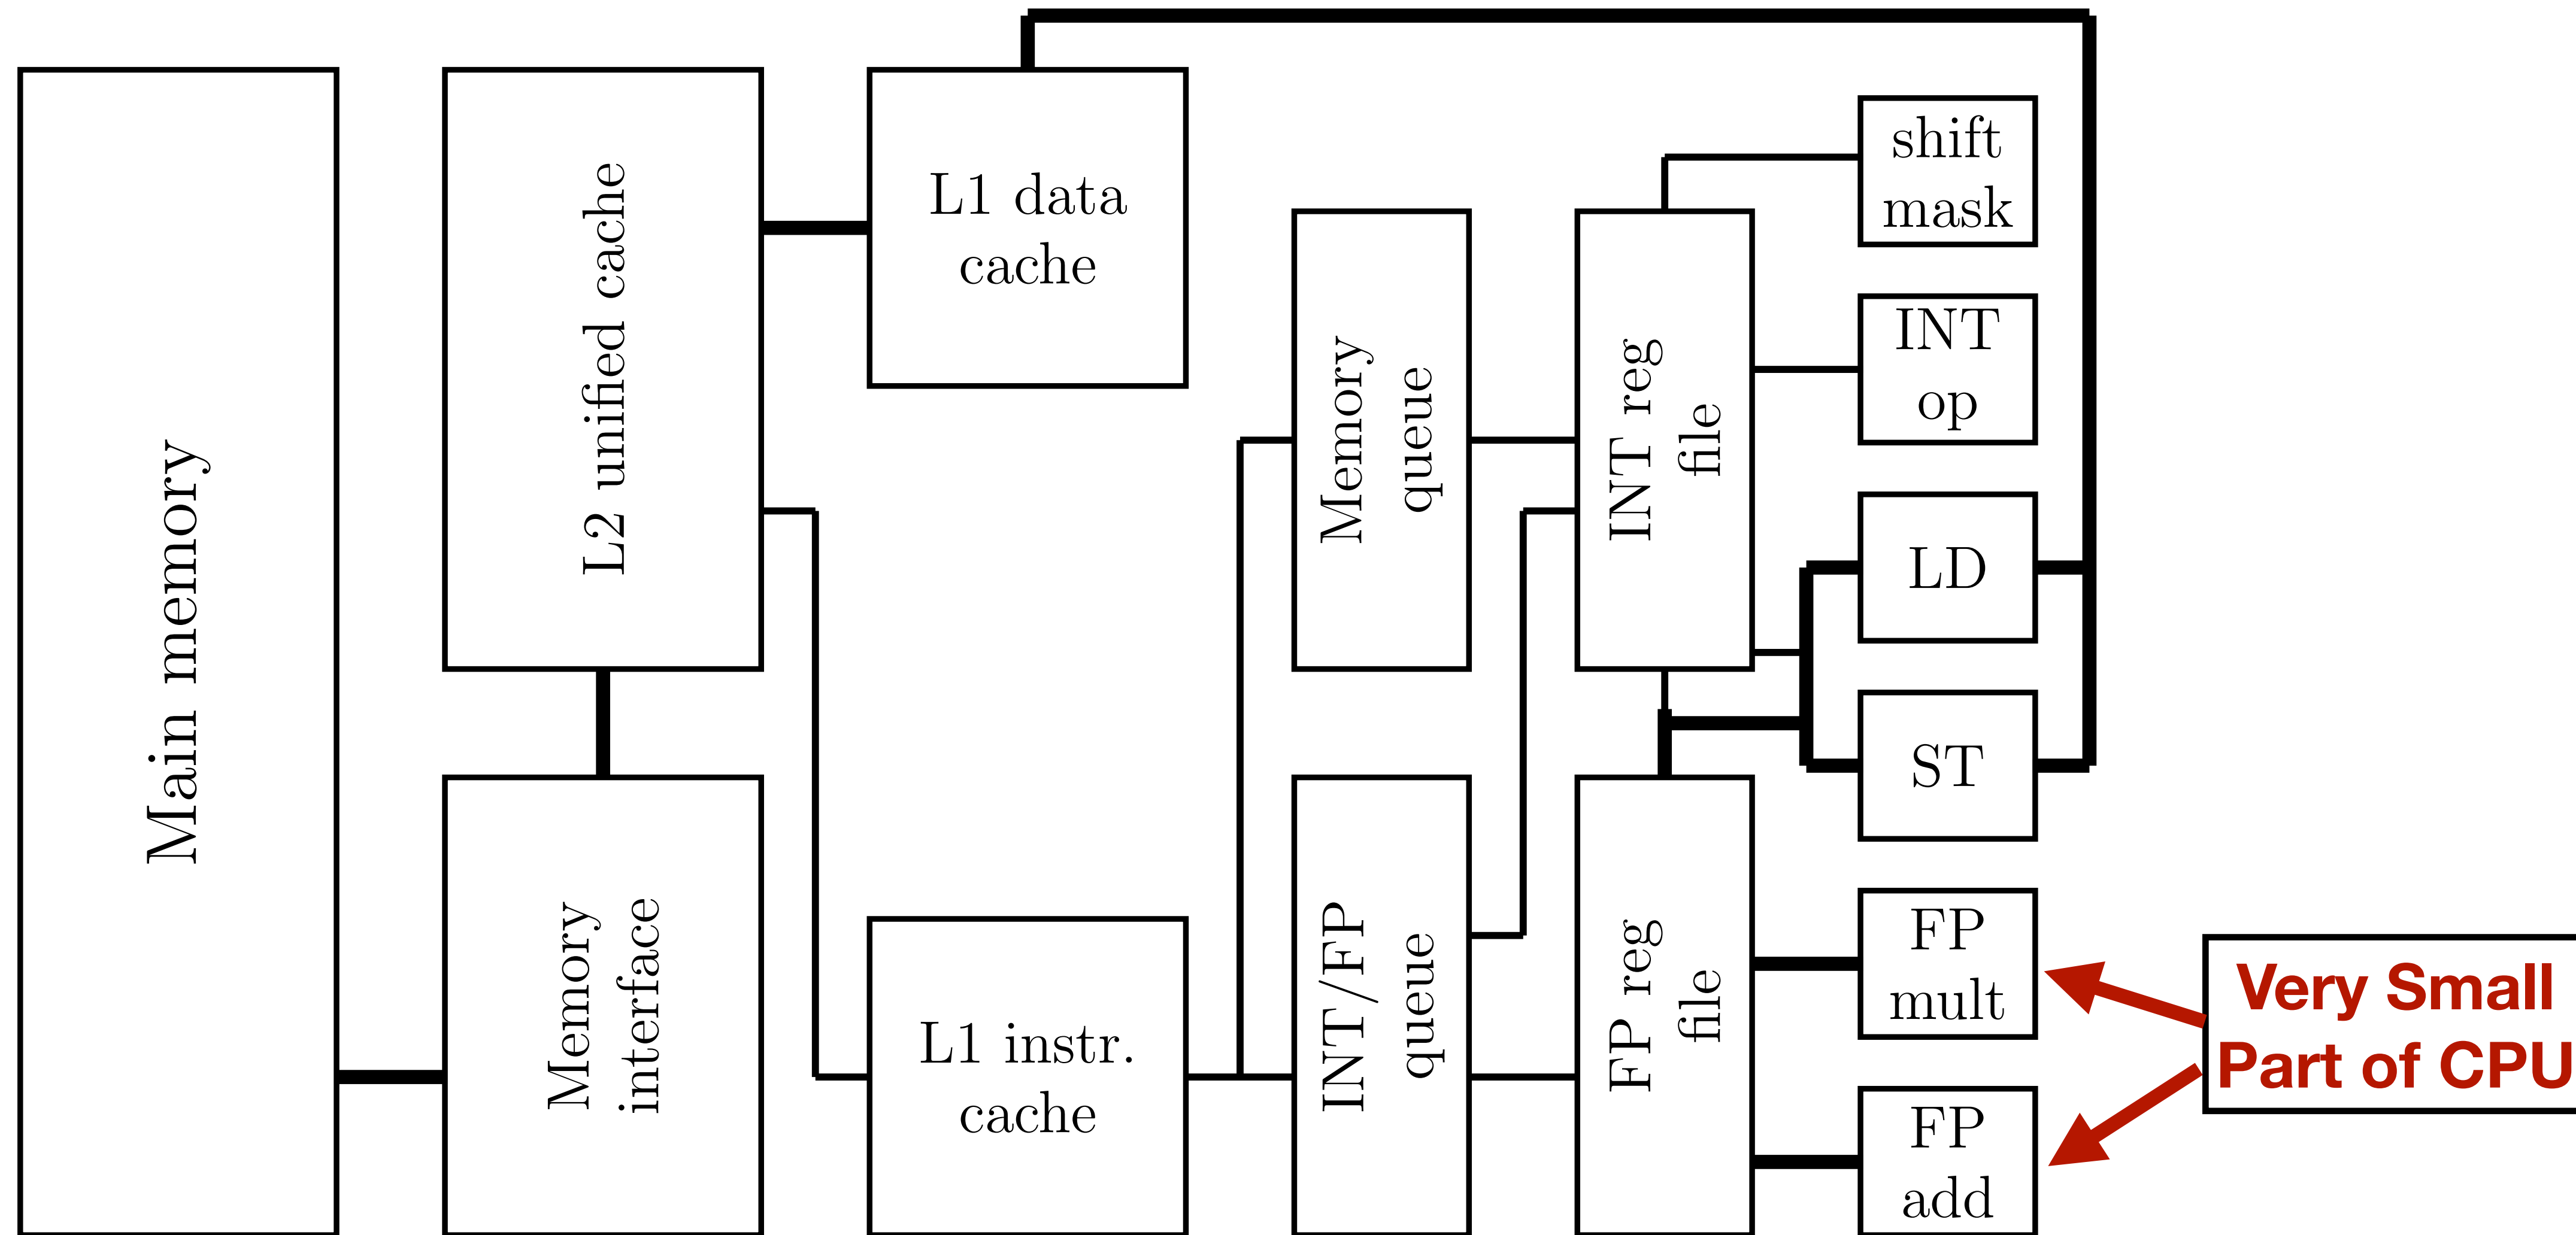
# Introduction to Parallel Processing

Lecture 4 : Vectorization

09/02/2022

Professor Amanda Bienz

# Cache-Based Microprocessor

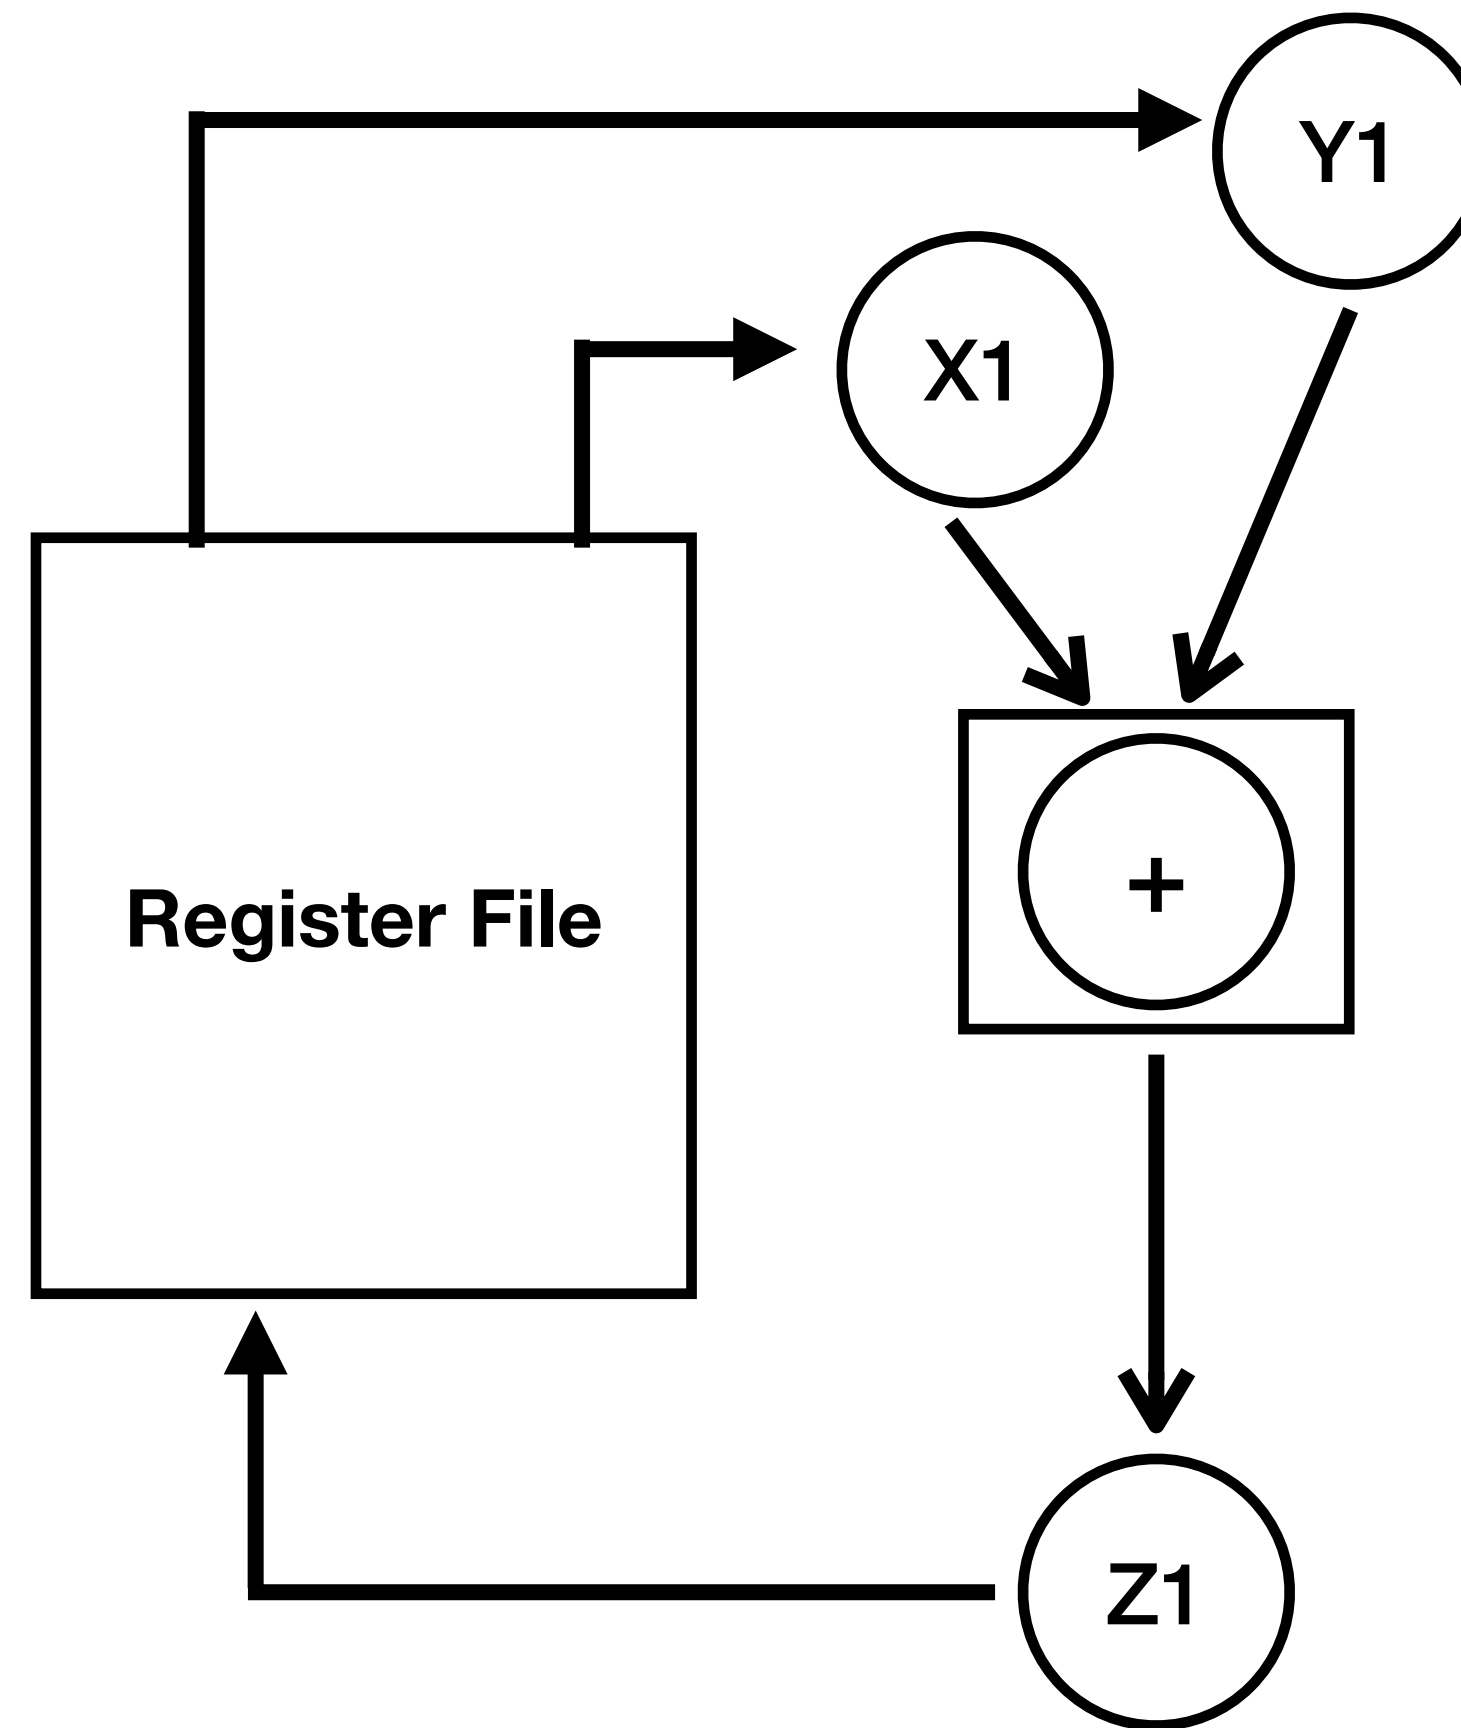


**Want to improve performance, have more operations at once**

# Serial Architecture

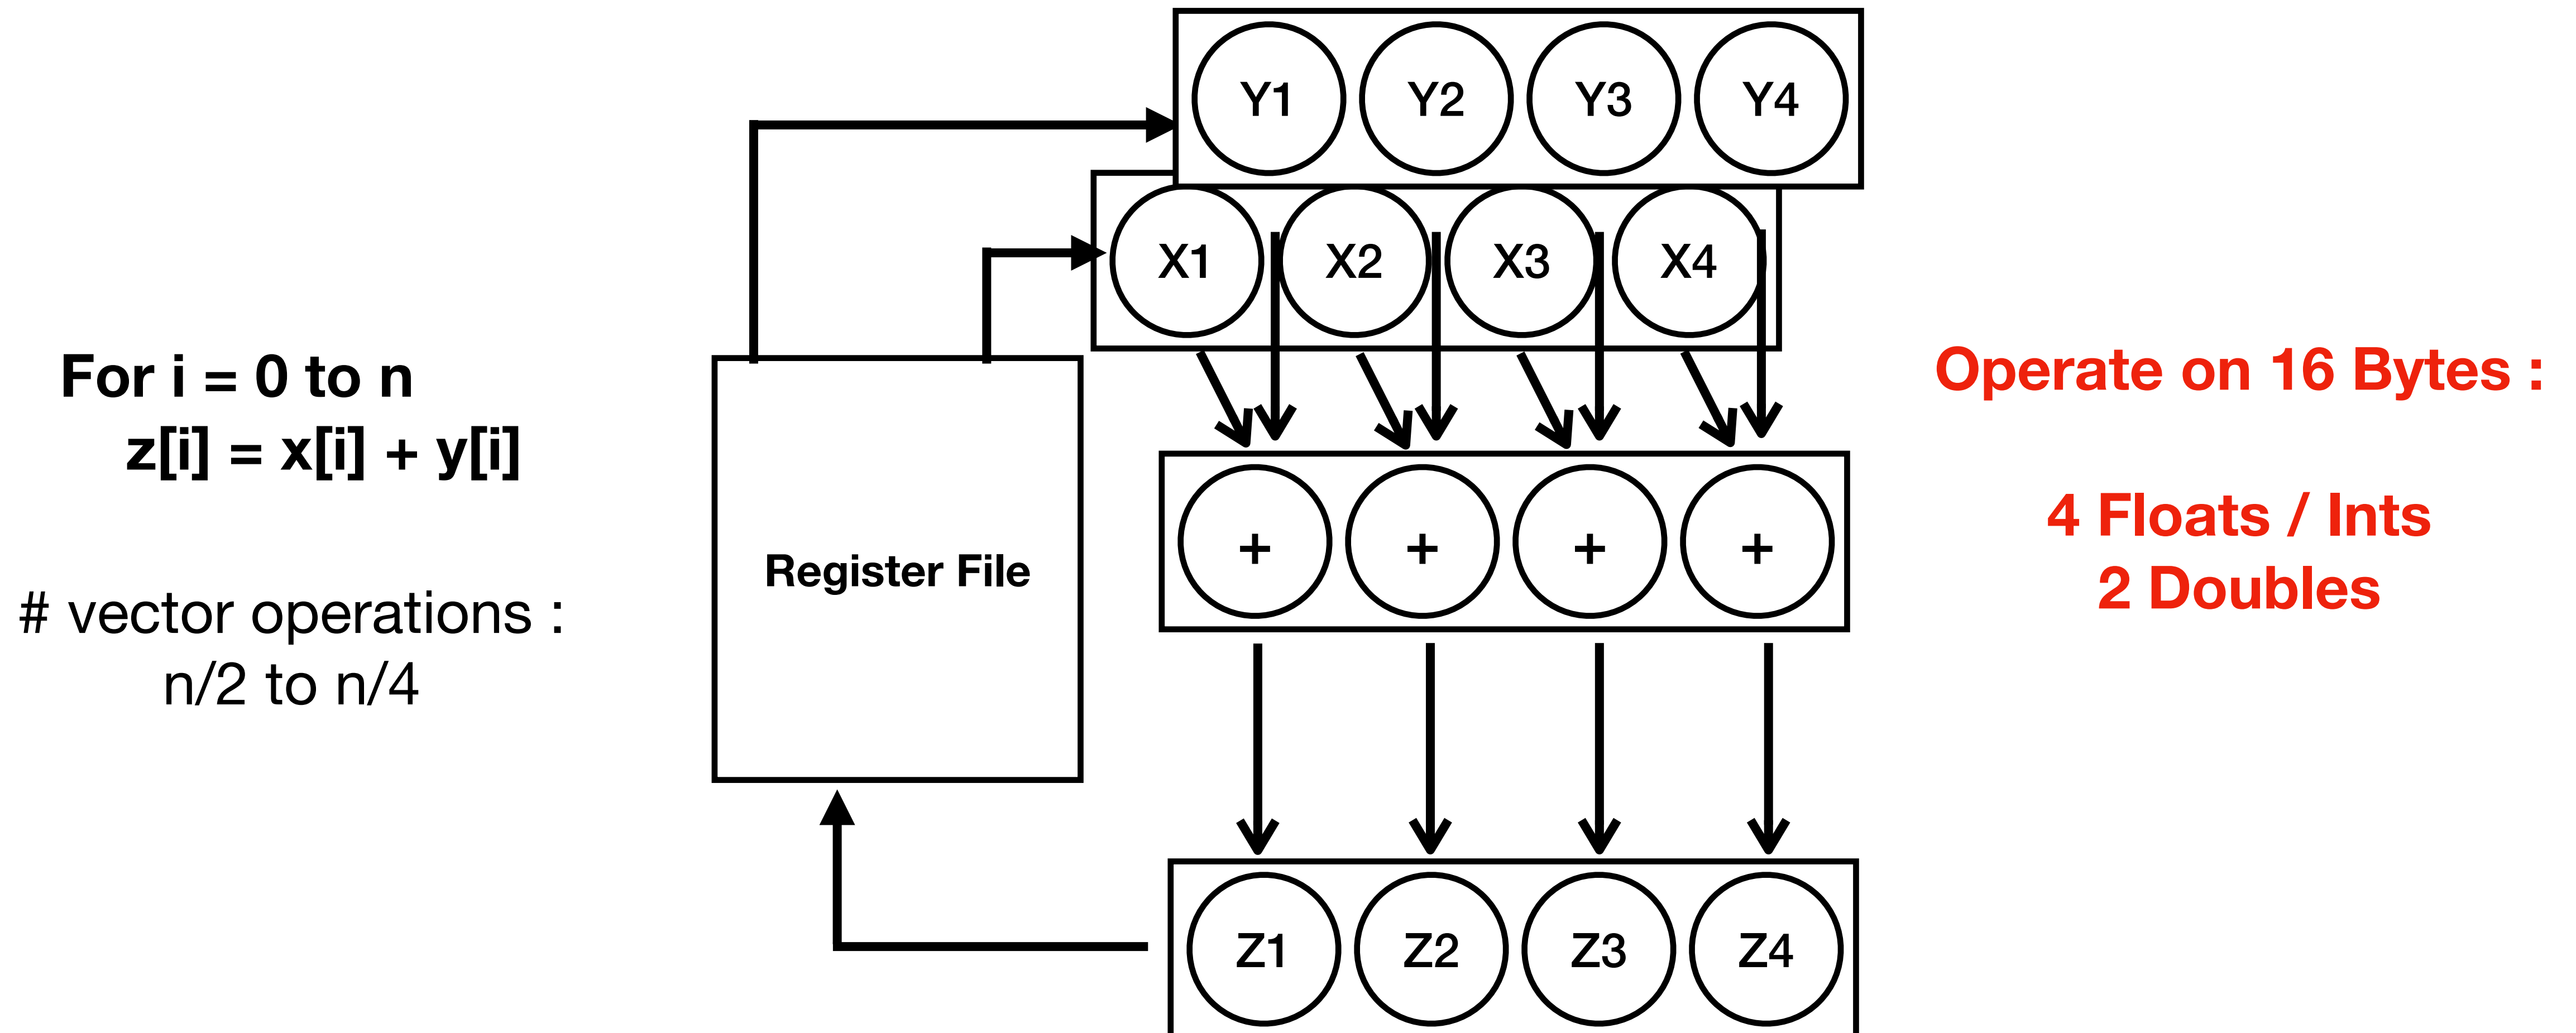
Say we want to add two arrays together :  $Z = X + Y$

For  $i = 0$  to  $n$   
 $z[i] = x[i] + y[i]$   
# scalar operations :  $n$



# Vector Architecture

Say we want to add two arrays together :  $Z = X + Y$



# Vectorized Loops

N	Bytes	Scalar	Vector	Speedup
1,000	4	5.5E-07	2.2E-07	2.6x
1,000	8	5.7E-07	2.8E-07	2x
1,000,000	4	8E-04	5.6E-04	1.4x
1,000,000	8	1.4E-03	1.3E-03	1.1x

# Let's Step Through Matrix Multiplication

$$\begin{bmatrix} A \\ n \times n \end{bmatrix} \times \begin{bmatrix} B \\ n \times n \end{bmatrix} \rightarrow \begin{bmatrix} C \\ n \times n \end{bmatrix}$$

# Data Dependencies

For  $i = 0$  to  $n$

$$z[i] = x[i] + y[i]$$

- Let's look back at this addition operation
- Vector operations : load a block of four floats (or two doubles) into the vector register (e.g load  $x[0]$ ,  $x[1]$ ,  $x[2]$ ,  $x[3]$ )
- Vector operation might look like the following
  - $z[0:3] = x[0:3] + y[0:3]$
  - $z[4:7] = x[4:7] + y[4:7]$

# Data Dependencies

For  $i = 0$  to  $n$

$$z[i] = x[i] + y[i]$$

- What if  $\&z[0]$  points to the same memory address as  $\&x[1]$
- Then:
  - $z[0] = x[0] + y[0]$
  - $z[1] = x[1] + y[1] = \mathbf{z[0]} + y[1]$
  - $z[2] = \mathbf{z[1]} + y[2]$



# Data Dependencies

- A few examples:

- $X = A + B$   
 $C = X + A$

- $A = X + B$   
 $X = C + D$

- $X = A + B$   
 $X = C + D$

# Data Dependencies

- If no data dependencies, instructions can be executed in any order, in parallel, or in a vector operation
- If there are data dependencies, the compiler (and user) cannot perform these optimizations
  - But, sometimes data can be rearranged to avoid dependencies
  - i.e. find two operations that are not dependent and optimize these