Introduction to Parallel Processing

Lecture 4 : Collective Communication Introduction

Professor Amanda Bienz

Collective Operations

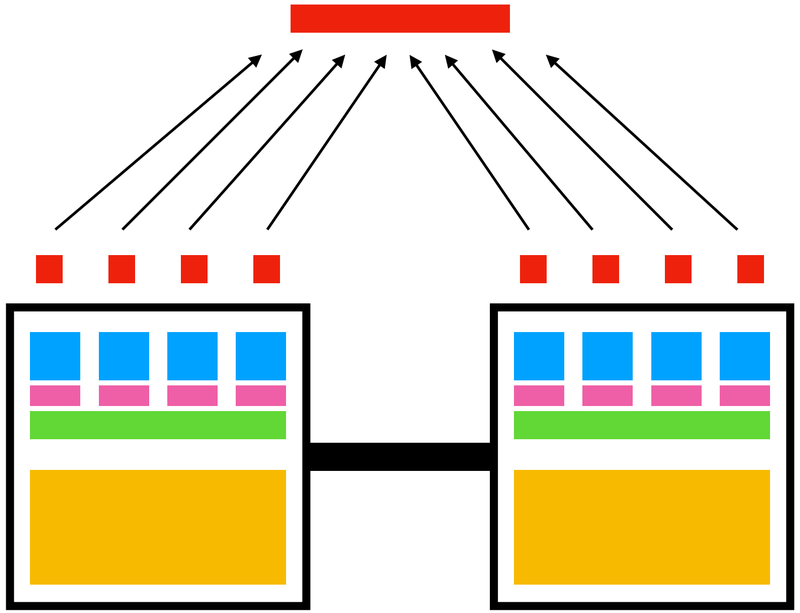

- Consider the case where each process has a portion of data, and we want to combine all of this data in some manner
- E.G. Each process holds a single number, and we want to find the sum of all of these numbers
- How do we handle this?

Reduction

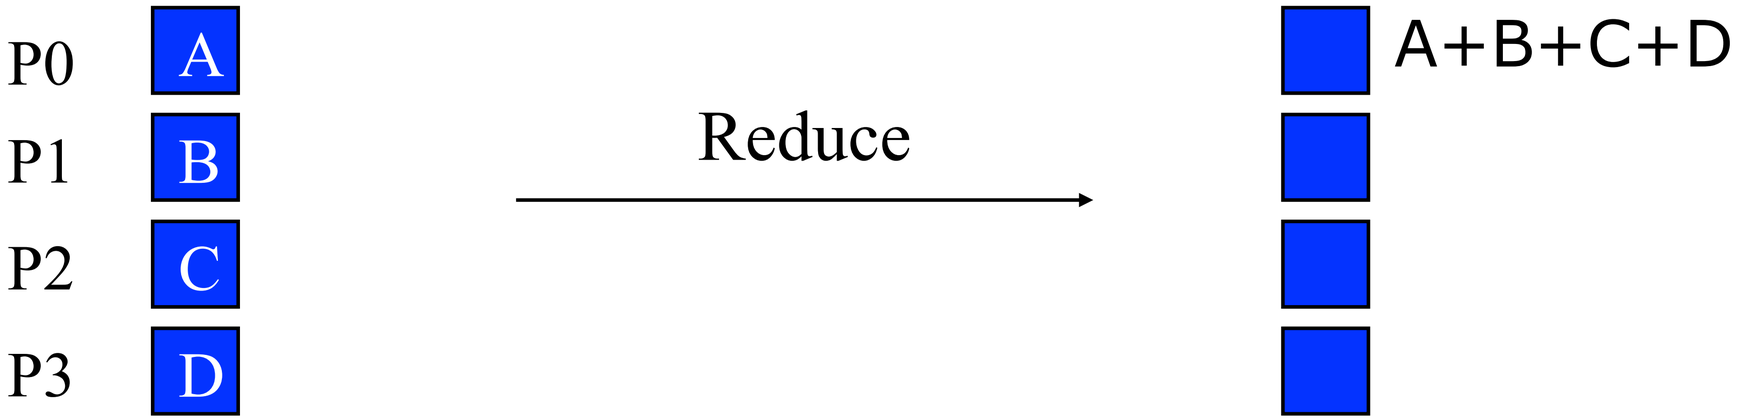

MPI_Reduce

- Reduction to a single processor (only one process will hold final result)

- MPI_Reduce(const void* sendbuf, // Buffer containing original data
  void* recvbuf, // Buffer to hold reduced data
  int count, // Buffer size
  MPI_Datatype datatype, // Type of data (e.g. MPI_INT)
  MPI_Op op, // Operation (MPI_SUM, MPI_MAX,…)
  int root, // Process to hold reduced value
  MPI_Comm comm) // MPI_COMM_WORLD

- **MPI_Reduce(&sum, &global_sum, 1, MPI_DOUBLE, MPI_SUM,**
  0, MPI_COMM_WORLD);

When is MPI_Reduce used?

- When do we want to reduce a variable only to a master ‘root’ process?
- Often, if something should be printed out
- For example, if each process times a method, we only want one process to print out the maximum time
- Similarly, an MPI program that calculates and prints a sum

Reduction

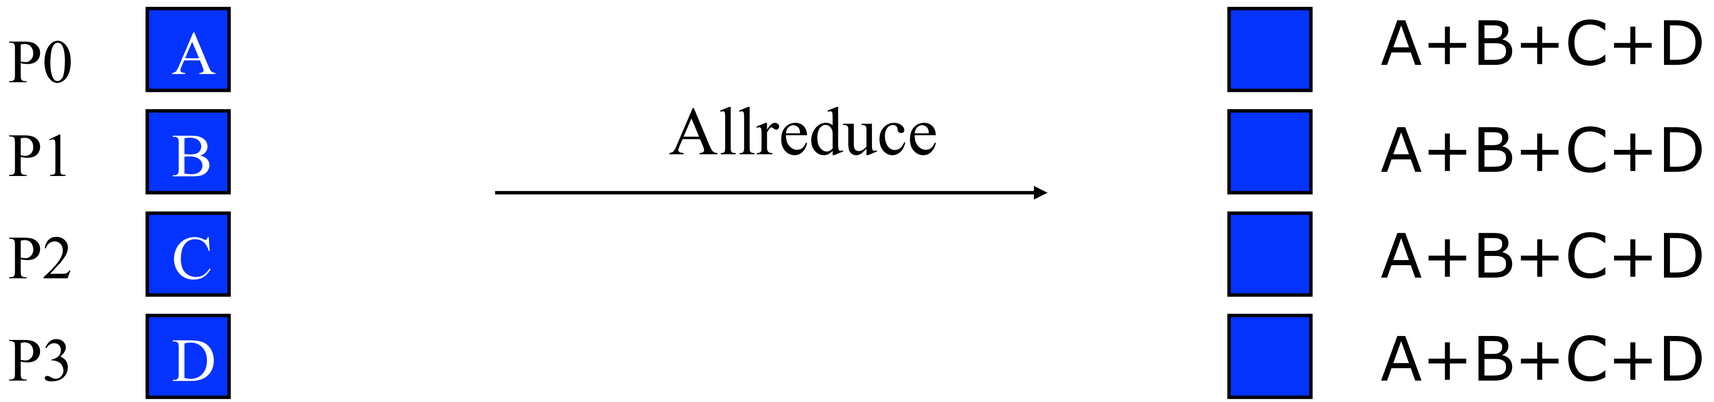

MPI_Allreduce

- Reduction to every process (all will hold final result)

- MPI_Allreduce(const void* sendbuf, // Buffer containing original data
  void* recvbuf, // Buffer to hold reduced data
  int count, // Buffer size
  MPI_Datatype datatype, // Type of data (e.g. MPI_INT)
  MPI_Op op, // Operation (MPI_SUM, MPI_MAX,…)

  MPI_Comm comm) // MPI_COMM_WORLD

- **MPI_Allreduce(&sum, &global_sum, 1, MPI_DOUBLE, MPI_SUM, MPI_COMM_WORLD);**

When is MPI_Allreduce used?

- Iterative Methods : Converging to solution and want to stop when close enough

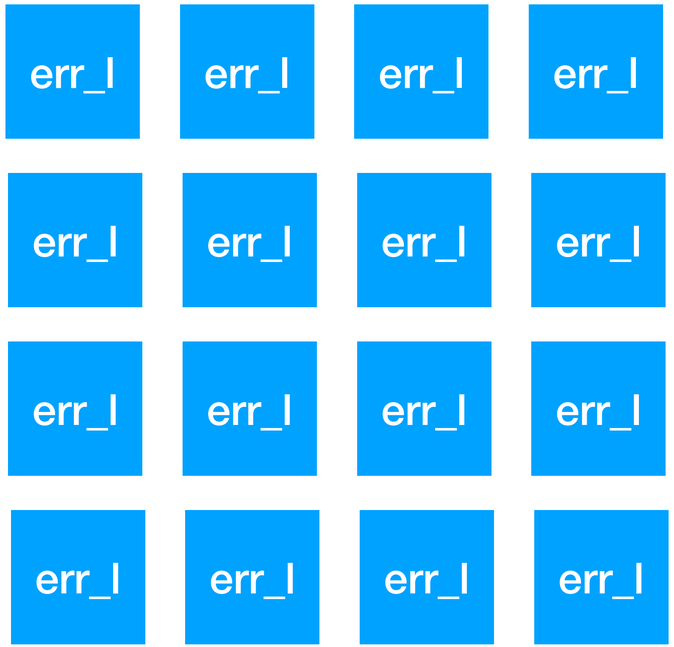

16 Process, each hold a local error

When is MPI_Allreduce used?

- Dot product of two vectors:
  for (i = 0 to n)
      sum += x[i] * y[i]

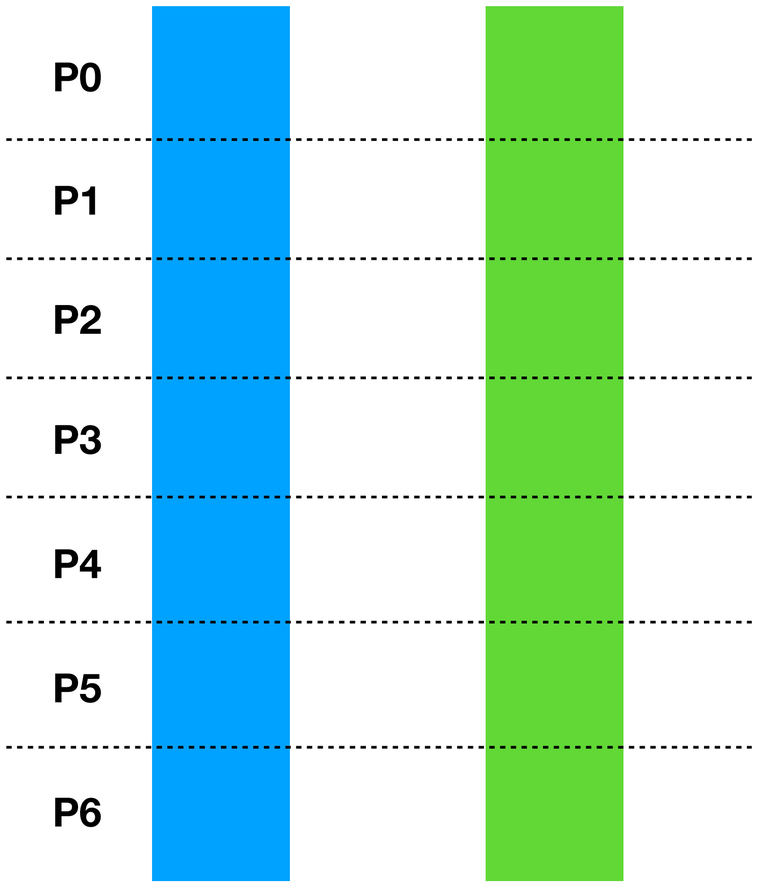

Two vectors, split across 7 processors

When is MPI_Allreduce used?

- Training neural networks
- Stochastic Gradient Descent : finds gradients, representing direction to go towards the minimum
- Tons of data spread across processors, each resulting in gradients
- Often have hundreds of millions of gradients, or directions, that need reduced

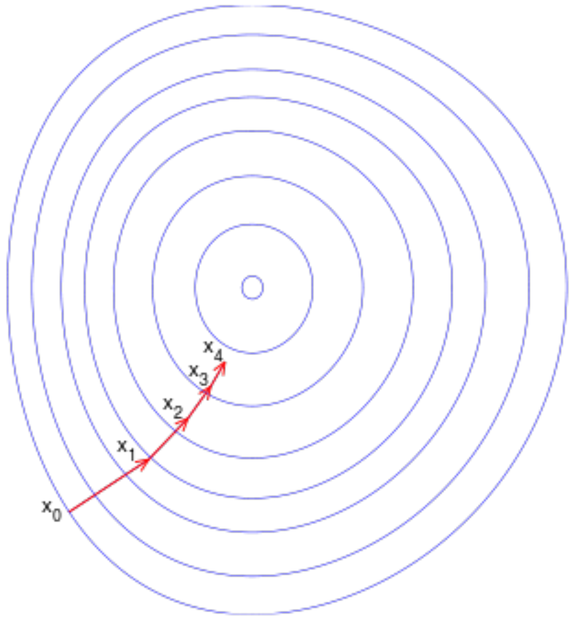

Broadcast

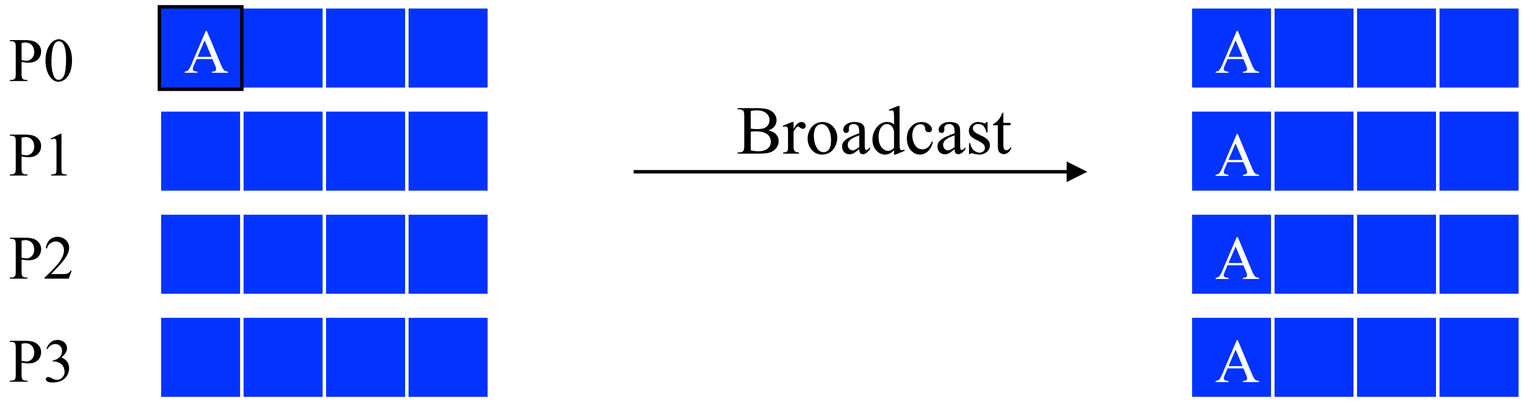

Broadcast

- Send data from one process to all other processes

- MPI_Bcast(void* buffer,
  int count,
  MPI_Datatype datatype,
  int root,
  MPI_Comm comm)

- int size = rand();
  MPI_Bcast(&size, 1, MPI_INT, 0, MPI_COMM_WORLD);

Scatter and Gather

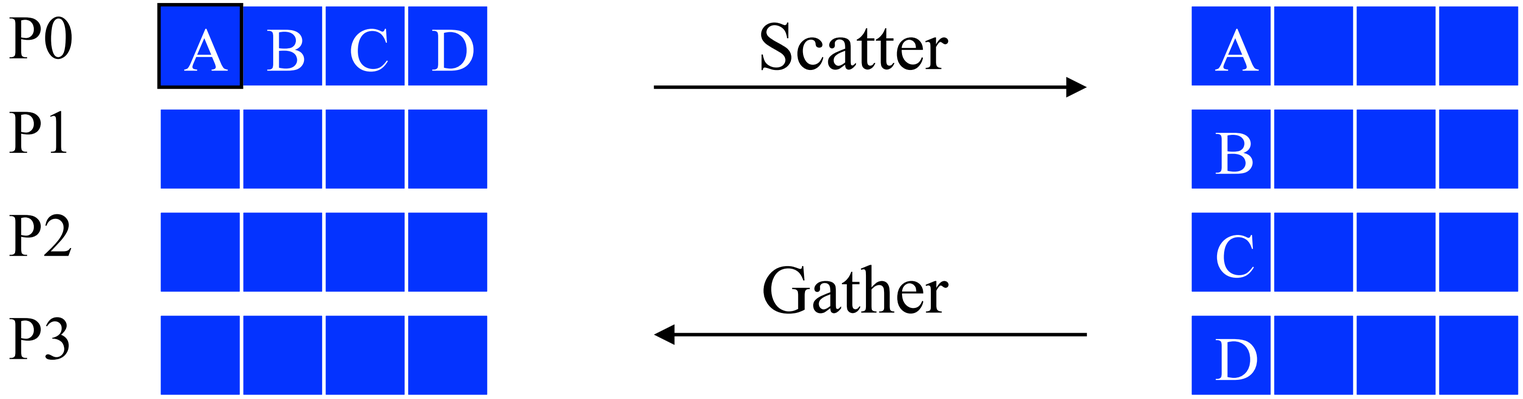

Scatter

- Sends data from one process to all other processes. Each process gets separate portion of data.

- MPI_Scatter(const void* sendbuf,
  int sendcount, // Number vals to send to each proc
  MPI_Datatype sendtype,
  void* recvbuf,
  int recvcount, // Number vals to recv
  MPI_Datatype recvtype,
  int root,
  MPI_Comm comm)

- int* vals = new vals[num_procs];
  int recv_val;
  MPI_Scatter(vals, 1, MPI_INT, &recv_val, 1, MPI_INT, 0, MPI_COMM_WORLD);

Gather

- Opposite of scatter, each process starts with some values and want to gather all of these values onto a single process

- MPI_Gather(const void* sendbuf,
  int sendcount, // Number vals, send all to every proc
  MPI_Datatype sendtype,
  void* recvbuf,
  int recvcount, // Number vals to recv from each proc
  MPI_Datatype recvtype,
  int root,
  MPI_Comm comm)

- int* vals = new vals[num_procs];
  int send_val = …;
  MPI_Gather(&send_val, 1, MPI_INT, vals, 1, MPI_INT, 0, MPI_COMM_WORLD);

Allgather

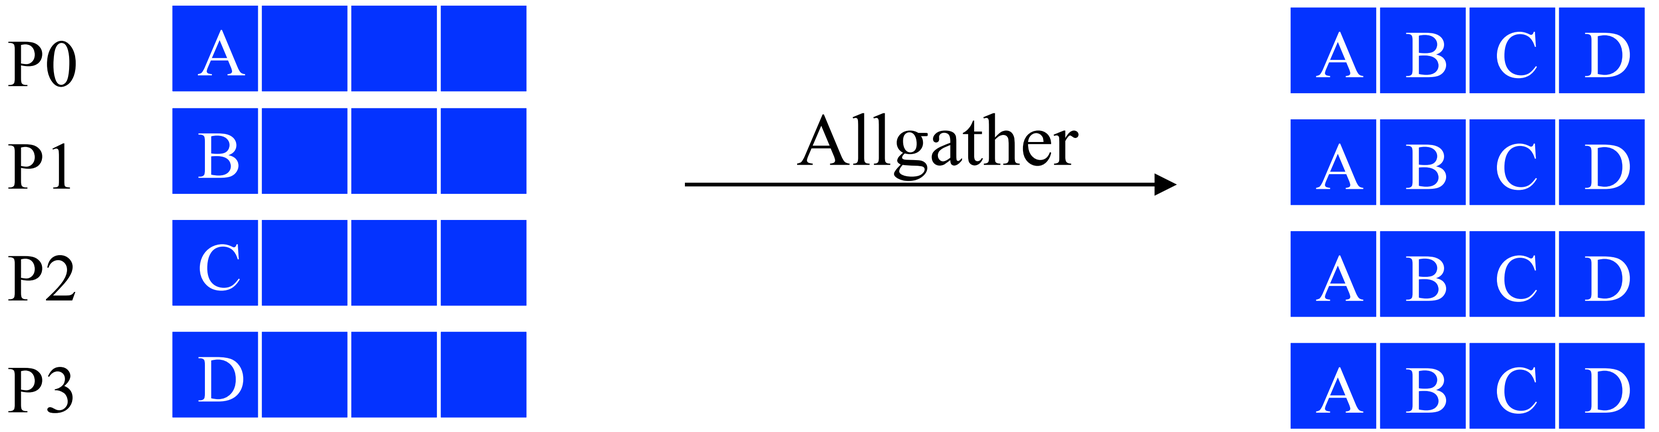

Allgather

- Same as gather, but have all data end up on every process

- MPI_Allgather(const void* sendbuf,
  int sendcount, // Number vals, send all to every proc
  MPI_Datatype sendtype,
  void* recvbuf,
  int recvcount, // Number vals to recv from each proc
  MPI_Datatype recvtype,
  MPI_Comm comm)

- int* vals = new vals[num_procs];
  int send_val = …;
  MPI_Allgather(&send_val, 1, MPI_INT, vals, 1, MPI_INT, MPI_COMM_WORLD);

All-to-all

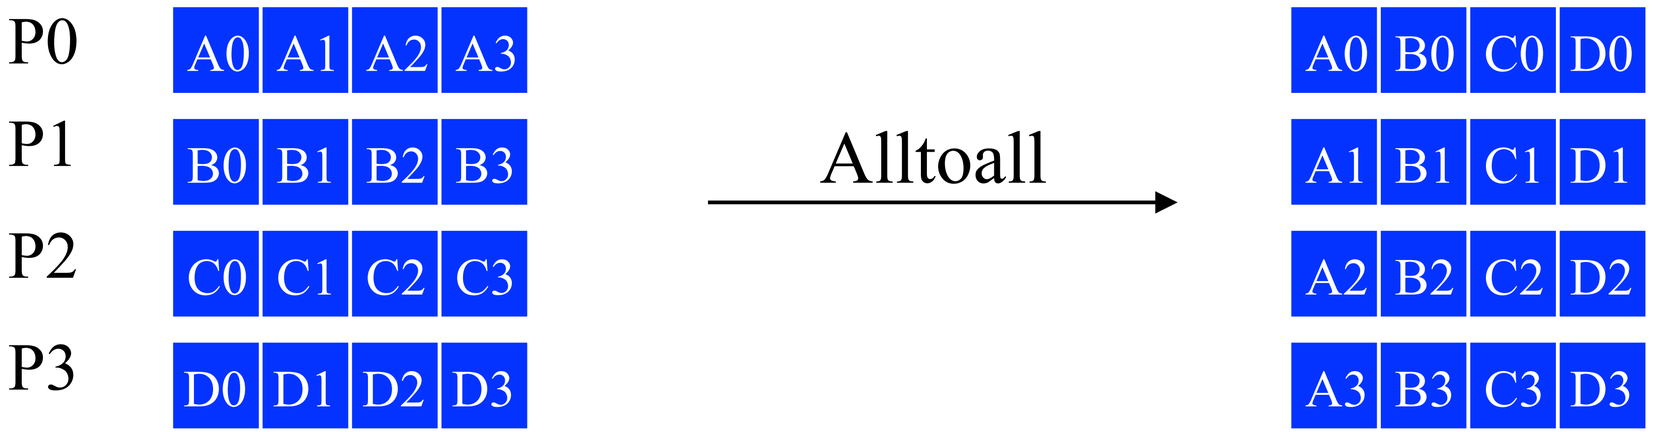
P0
P1
P2
P3
A0 A1 A2 A3
B0 B1 B2 B3
C0 C1 C2 C3
D0 D1 D2 D3
Alltoall
A0 B0 C0 D0
A1 B1 C1 D1
A2 B2 C2 D2
A3 B3 C3 D3

All To All

- Similar to allgather, but instead of getting entire array, each process gets different part of data

- MPI_Alltoall(const void* sendbuf,
  int sendcount, // Number vals to send to each proc
  MPI_Datatype sendtype,
  void* recvbuf,
  int recvcount, // Number vals to recv from each proc
  MPI_Datatype recvtype,
  MPI_Comm comm)

- int* send_vals = new int[num_procs] …
  int* recv_vals = new int[num_procs] …
  MPI_Alltoall(send_vals, 1, MPI_INT, recv_vals, 1, MPI_INT, MPI_COMM_WORLD)

Collective Operations With Computation

- Reduce, Allreduce (previously discussed) have computation such as summing values together
- Scan, Exscan : Combination of data from all prior ranks
- Reduce_scatter : All to all, but combines results

Collective Computation

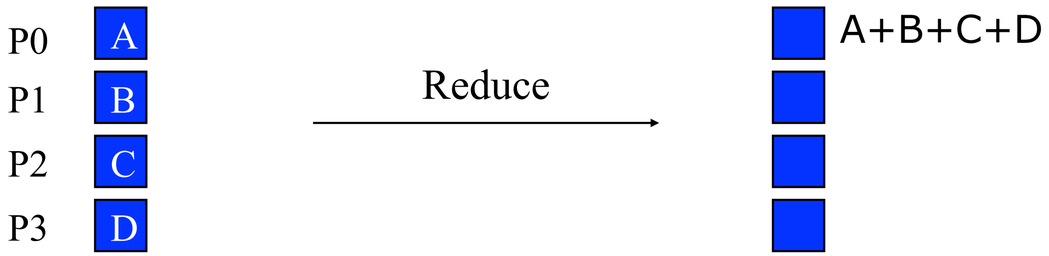

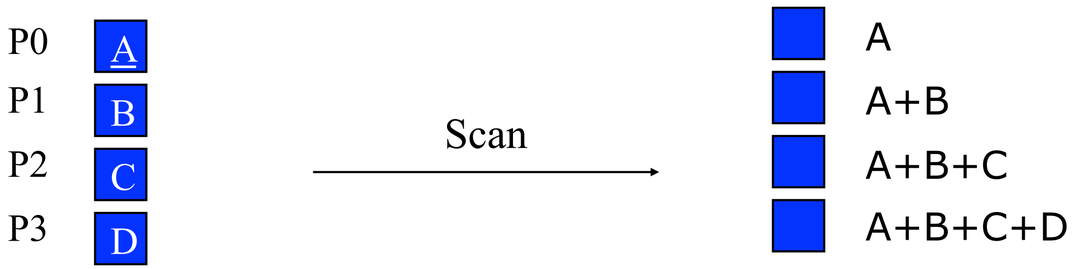

PARALLEL@ILLINOIS

https://wgropp.cs.illinois.edu/courses/cs598-s15/lectures/lecture29.pdf

Collective Computation

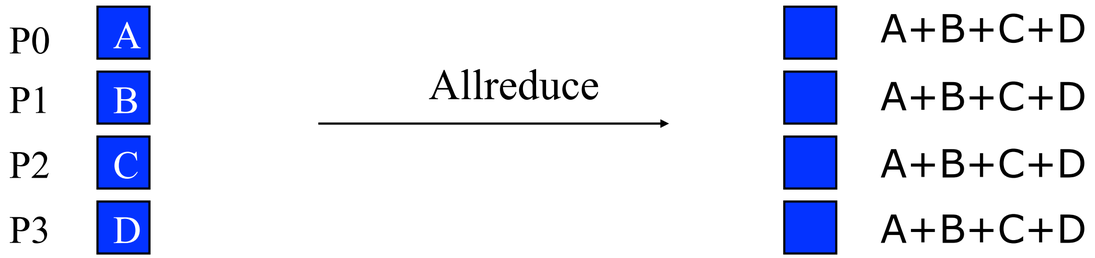

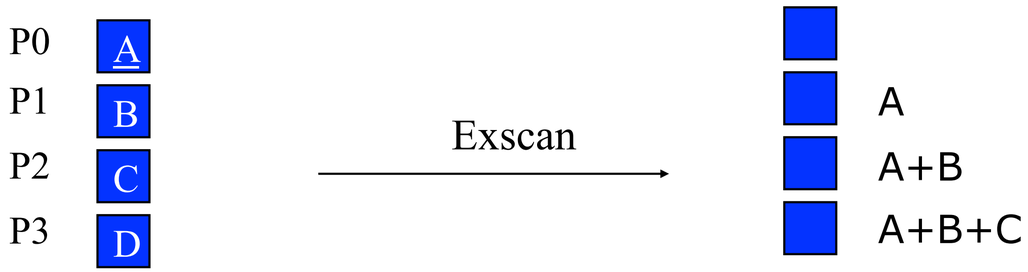

PARALLEL@ILLINOIS

https://wgropp.cs.illinois.edu/courses/cs598-s15/lectures/lecture29.pdf

Built In MPI_Op Values

MPI Built-in Collective Computation Operations

• **MPI_MAX**	Maximum
• **MPI_MIN**	Minimum
• **MPI_PROD**	Product
• **MPI_SUM**	Sum
• **MPI_LAND**	Logical and
• **MPI_LOR**	Logical or
• **MPI_LXOR**	Logical exclusive or
• **MPI_BAND**	Bitwise and
• **MPI_BOR**	Bitwise or
• **MPI_BXOR**	Bitwise exclusive or
• **MPI_MAXLOC**	Maximum and location
• **MPI_MINLOC**	Minimum and location