

Introduction to Parallel Processing

Lecture 7 : Advanced Datatypes

Professor Amanda Bienz

First, MPI Refresher

- Setup
 - MPI_Init called first by each process
 - MPI_Finalize called last by each process
- Datatypes:
 - Basic datatypes MPI_INT, MPI_CHAR, MPI_DOUBLE, etc
 - Can derive complex datatypes such as arrays, structs, striped types (**today**)
- Communicators: (i.e. MPI_COMM_WORLD)
 - All communication happens inside a communicator
 - Each process in communicator has rank in that communicator
 - Will talk about more complicated communicators later (i.e. subset of processes in a group)

First, MPI Refresher

- Basic point-to-point MPI communication is between pairs of processes in a communicator
- Communication is two-sided — sender must call send, receiver must call receive
- Sender specifies destination and a tag (*integer value*) associated with message
- Receiver specifies source and tag
 - MPI_Status can be used to receive message of any size or from any source **(today)**
- Must be careful to avoid deadlocks

MPI Datatypes


- The data in a message to be sent or received is described by:
 - Address of data
 - Count (size of data)
 - Datatype
- Predefined datatypes : MPI_INT, MPI_DOUBLE, MPI_CHAR, etc
- Can also have:
 - Contiguous array of multiple MPI_Datatypes
 - Strided block of datatypes
 - Arbitrary structure of datatypes
- Can also construct custom datatypes such as array of pairs, columns of a matrix stored row-wise

Why Datatypes?

- MPI implementation can support heterogeneous communication between machines
 - Different memory representations
 - Different lengths of elementary datatypes
- Specifies application-oriented layout of data in memory
 - Can reduce memory-to-memory copies
 - Allows use of special hardware (scatter/gather) when available

Non-Contiguous Datatypes

- Provided to allow MPI implementations to avoid copy



- Not widely implemented (yet)
- Handling of important special cases (constant stride, contiguous structures)

Sending Columns of a Matrix

- Assume you want to send a column of the matrix $A[m,n]$
- $A[i, \text{column}]$ is not adjacent in memory to $A[i+1, \text{column}]$
- **One solution: Send each element separately:**
for $i = 0$ to m :
 `MPI_Send(&(A[i, column], 1, MPI_DOUBLE, ...)`
- Why don't you want to do this?

Sending Columns of a Matrix

- Assume you want to send a column of the matrix $A[m,n]$
- $A[i, \text{column}]$ is not adjacent in memory to $A[i+1, \text{column}]$
- **Another solution: Pack into contiguous buffer:**
for $i = 0$ to m :
 $\text{buffer}[i] = (A[i, \text{column}], 1, \text{MPI_DOUBLE}, \dots)$
 $\text{MPI_Send}(\text{buffer}, m, \text{MPI_DOUBLE}, \dots)$
- Why don't you want to do this?

MPI Type Vector

- Create a single datatype representing elements separated by a constant distance (stride) in memory
 - m items separated by a stride of n
 - MPI_Datatype newtype
 - MPI_Type_vector(m, 1, n, MPI_DOUBLE, &newtype)
 - MPI_Type_commit(&newtype)
 - Type_commit required before using the new type in an MPI communication operation
 - MPI_Type_free(&newtype)
- Now, can send the entire column in one instance :
MPI_Send(&A[0, column], 1, newtype,...)

MPI Derived Datatype Calls

- `MPI_Type_contiguous` : allows you to create a datatype for a contiguous array (i.e. a row in C, column in Fortran)
- `MPI_Type_struct` : allows for heterogeneous datatypes
 - More complicated, structs usually have padding determined by compiler
 - Have to pass an array of types and an array of offsets
 - C primitive `offsetof(struct, member)` is useful for computing these offsets
- `MPI_Type_indexed` : allows for non-contiguous, unevenly spaced data (2D blocks of a matrix)

Packing Data

- MPI_Pack : pack data of any type into a contiguous char* buffer
- Predefined datatype : MPI_PACKED
- Useful for sending sparse matrices:
 - int* rowptr
 - int* col_indices
 - double* data
- Want to send full matrix as a single message

Packing Data

- `MPI_Pack(const void* inbuf,
int incount,
MPI_Datatype datatype,
void* outbuf,
int outsize,
int* position,
MPI_Comm comm)`
- Packs datatype into contiguous memory
- Can pack single value at a time or group of values at one time

Sending Packed Data

- `MPI_Send(char *packed_buffer, packed_size, MPI_PACKED, int destination, int tag, MPI_Comm comm)`
- Sends contiguous packed buffer

Unpacking Data

- `MPI_Unpack(const void* inbuf,
int insize,
int* position,
void* outbuf,
int outcount,
MPI_Datatype datatype,
MPI_Comm comm)`
- Unpacks received `char*` to original datatype
- Can unpack single value at a time, or in group of values at once

Helpful Packing Method

- `MPI_Pack_size(int count,
MPI_Datatype,
MPI_Comm comm,
int* pack_size)`
- Know the original size of data, this method will return the size of the data once it is packed

How do we actually send / recv this?

- Now we can send all data at once time!
- This is great, because there is a large overhead associated with sending each message, regardless of size
- Fewer messages = cheaper
- However, how do we receive a sparse matrix without knowing the size we will be receiving?

Receiving a Message of any size

- `MPI_Probe(int source,
int tag,
MPI_Comm comm,
MPI_Status* status)`
 - Waits until the process finds a message from 'source' of any size
- `MPI_Get_count(MPI_Status* status,
MPI_Datatype datatype,
int* count)`
 - Returns the size of the message found by probe
- Then, call `MPI_Recv` with this count!

Receiving a message from any process

- `MPI_ANY_SOURCE` receives a message from any source
- `MPI_STATUS.MPI_SOURCE` : the sending process attached to a message
- Can use this in a few ways
 - `MPI_Recv(MPI_ANY_SOURCE, count, ...)` recv a message of size 'count' from any process
 - Only works if all messages will have size count
 - `MPI_Probe(MPI_ANY_SOURCE, int tag, MPI_Comm comm, MPI_Status* status)`
 - Finds a message sent to my process, from any source, of any size
 - Can then find the source (`status.MPI_SOURCE`) and size (`MPI_Get_count(...)`)