

Introduction to Parallel Processing

Lecture 17 : MPI I/O


10/21/2022

Professor Amanda Bienz

Why MPI IO?

- MPI allows for parallel performance when performing I/O
- Can write to a single file from all processes, instead of one file per process
- MPI provides a lot of helpful I/O methods

Non-Parallel I/O




- Non-parallel
- Performance worse than sequential
- Legacy from before application was parallelized
- Either MPI or not



Independent Parallel I/O


- Each process writes to a separate file



- Pro: parallelism
- Con: lots of small files to manage
- Legacy from before MPI
- MPI or not



Cooperative Parallel I/O



- Parallelism
- Can only be expressed in MPI
- Natural once you get used to it



Why MPI is a natural fit for I/O

- Writing is similar to sending messages
- Reading is similar to receiving messages
- Parallel I/O needs:
 - Collective operations
 - User-defined datatypes for both memory and file layouts
 - Communicators to separate application level message passing from I/O related message passing
 - Non-blocking operations

Parallel I/O

- What do we mean by parallel I/O?
 - Concurrent reads to multiple processes from a common file
 - Concurrent writes from multiple processes to a common file
 - A parallel file system and hardware that support these concurrent accesses

Simple IO Process

- Standard I/O:
 - Open the file
 - Read or write data
 - Close the file
- Parallel I/O with MPI:
 - Open the file : MPI_File_open
 - Write to the file: MPI_File_write
 - Close the file: MPI_File_close

File Open

- Collective over communicator
- Used to support collective I/O, which is important for performance
- Modes are similar to Unix open
- MPI_Info variable can provide additional hints for performance

File write

- Independent (example shows this... only called write from one rank)
- Use `MPI_File_write` or `MPI_File_write_at`
- Use `MPI_MODE_WRONLY` or `MPI_MODE_RDWR`
- If the file doesn't previously exist, must also use `MPI_MODE_CREATE`
- Can pass multiple flags with '|' in C or '+' in Fortran

Ways to Access a Shared File

- `MPI_File_seek`
 - `MPI_File_read`
 - `MPI_File_write`
 - `MPI_File_read_at`
 - `MPI_File_write_at`
 - `MPI_File_read_shared`
 - `MPI_File_write_shared`
- } like Unix I/O
- } combine seek and I/O for thread safety
- } use shared file pointer



File close

- Like open, close is also a collective operation
- This operation is similar to `MPI_Comm_free`

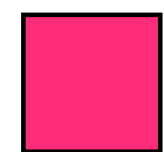
Why independent I/O?


- Collectives add synchronization... if all processes are highly unsynchronized during I/O, collectives could add high cost
- The overhead of collective calls (especially for small messages) can outweigh their benefits

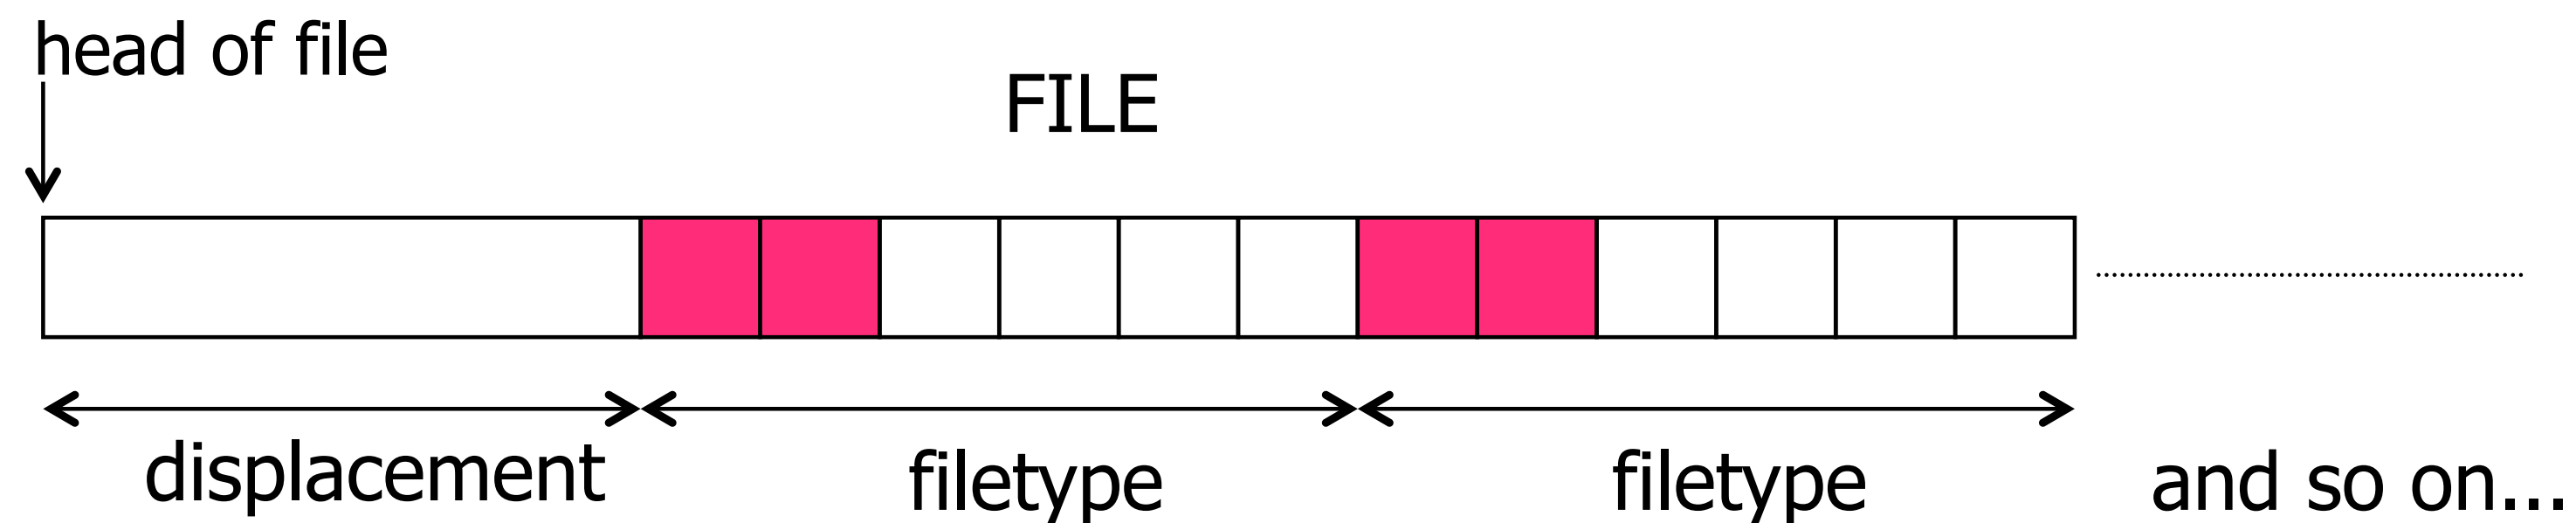
Noncontiguous I/O

- Each process describes the part of the file for which it is responsible
- Only the part of the file described by the file view is visible to the process; read and write will access these locations
- ‘File view’ created with `MPI_File_set_view`
 - Displacement : number of bytes to be skipped from start of file (e.g. to skip a file header)
 - Etype : basic unit of data access (basic or derived datatype)
 - Filetype : which portion of the file is visible to the process?

A Simple Noncontiguous File View Example


 etype = MPI_INT

 filetype = two MPI_INTs followed by a gap of four MPI_INTs



Collective I/O

- Allows communication of the big picture to a file system
- Basic idea : build large blocks so that reads and writes will be large
 - Requests from different processes may be merged together
 - Particularly effective when the accesses of different processes are noncontiguous and interleaved



Collective I/O

- `MPI_File_write_at_all`
 - `_all` says that all processes in the communicator (which was passed to `MPI_File_open`) will call this function
- Each process specifies its own information for the write
 - Calls the same information as if writing without a collective, but tells MPI to aggregate data

The Other Collective I/O Calls

- `MPI_File_seek`
 - `MPI_File_read_all`
 - `MPI_File_write_all`
 - `MPI_File_read_at_all`
 - `MPI_File_write_at_all`
 - `MPI_File_read_ordered`
 - `MPI_File_write_ordered`
- like Unix I/O
- combine seek and I/O for thread safety
- use shared file pointer



File Systems

- NFS (Network File System) : distributed file system
 - Allows a user on a computer to access files from some storage over a computer network
 - Typically only uses a single node

Parallel File Systems

- PNFS (Parallel Network File System):
 - Scalable parallel access to files distributed among multiple servers
 - Separates data and metadata
 - Moves metadata server out of the data path