

# Introduction to Parallel Processing

Lecture 12 : MPI Communicators

Professor Amanda Bienz

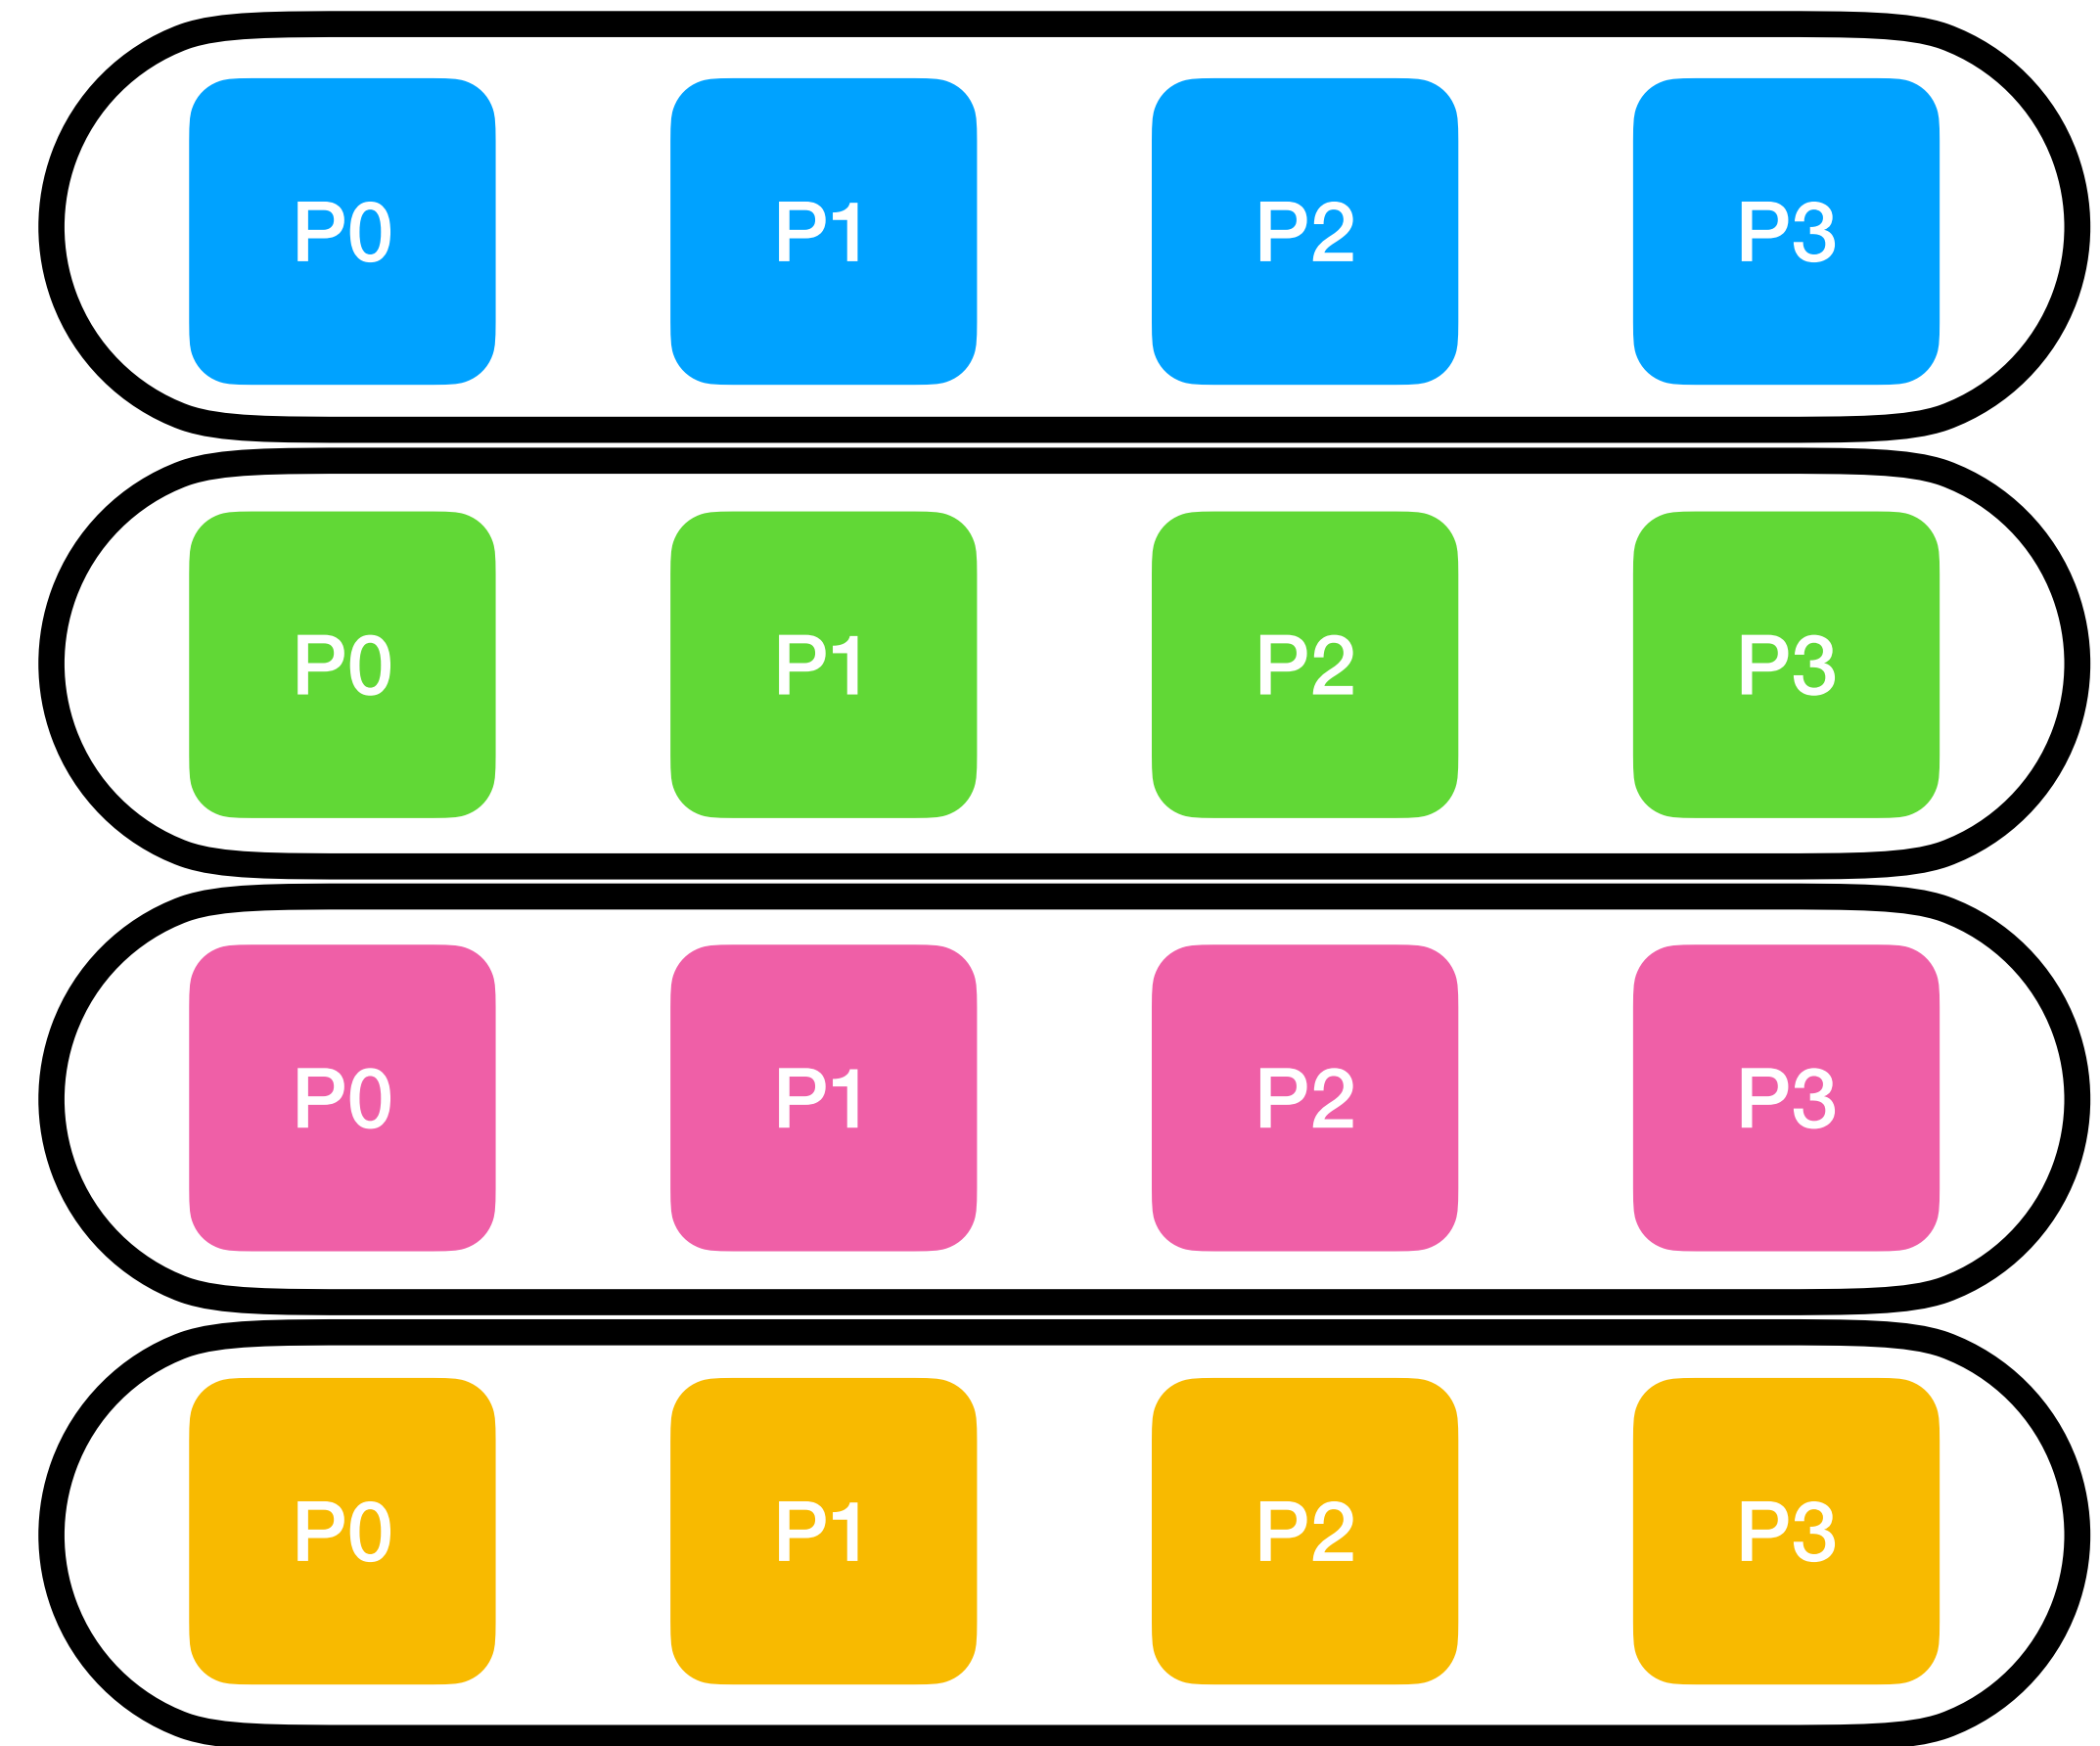

# MPI Communicators

- So far we have only talked about MPI\_COMM\_WORLD
- This is a communicator, or group, that contains all processes that were initialized during the 'mpirun' call
- A communicator provides a space in which processes can communicate with one another.

# Custom Communicators

- Can create custom communicators : sub-group of processes, also known as a sub-communicator
- `MPI_Comm_split(MPI_Comm orig_comm,  
                  int color,     Same color <—> same new communicator  
                  int key,  
                  MPI_Comm* new_comm)`     **Key implies rank in new communicator;  
if tie, based off lowest original rank (in orig\_comm)**
- `MPI_Comm_create(MPI_Comm orig_comm,  
                  MPI_Group group,  
                  MPI_Comm* new_comm)`

# Sub-Communicators



# Topology-Aware Communicators

- As we know from performance models, relative locations of processes has a large impact on cost of inter-process communication
- Often, we would like to create custom communicators based on the topology
- You can do this through external parameters, such as setting mapping of ranks to nodes with variable such as `MPICH_RANK_REORDER_METHOD`
- `MPI_Comm_split_type(MPI_Comm orig_comm,`  
    **`int split_type,`**                   **Most common `split_type` : `MPI_COMM_TYPE_SHARED`**  
    `int color,`  
    `int key,`  
    `MPI_Comm* new_comm)`
- New splitting types available in MPI 4

# OpenMPI Splitting Types

- OMPI\_COMM\_TYPE\_NODE : same as MPI\_COMM\_TYPE\_SHARED
- OMPI\_COMM\_TYPE\_HWTHREAD : splits by hardware thread
- OMPI\_COMM\_TYPE\_CORE : splits by core / processing unit
- OMPI\_COMM\_TYPE\_L1CACHE : splits by L1 cache
- OMPI\_COMM\_TYPE\_L2CACHE : splits by L2 cache
- OMPI\_COMM\_TYPE\_L3CACHE : splits by L3 cache
- OMPI\_COMM\_TYPE\_SOCKET : splits by socket
- OMPI\_COMM\_TYPE\_NUMA : splits by NUMA-region
- OMPI\_COMM\_TYPE\_BOARD : splits by board
- OMPI\_COMM\_TYPE\_HOST : splits by host
- OMPI\_COMM\_TYPE\_CU : splits by computational unit
- OMPI\_COMM\_TYPE\_CLUSTER : splits by cluster

# MPI\_Info


- Stores an unordered set of (key, value) pairs
- Many routines take an info argument (can be MPI\_INFO\_NULL)
- Ignores keys that are not recognized
- Allows for hints to be passed to MPI implementation
- However, you need to know what keys to set...

# Types of Communicators

- There are two types of communicators :
  - **Intra-communicator:** the type that we have been talking about.
    - All processes in the communicator can send messages to each other and participate in collective communication
  - **Inter-communicator:**
    - Can send messages between processes belonging to disjoint intra-communicators




# Inter-Communicator Example



- Communication Requirements:
  - Group 0 to communicate with Group 1
  - Group 1 to communicate with Groups 0 and 2
  - Group 2 to communicate with Group 1

# Inter-Communicator Example




- `MPI_Intercomm_create(MPI_Comm local_comm,  
int local_leader,  
MPI_Comm peer_comm,  
int remote_leader,  
int tag,  
MPI_Comm& inter_comm)`

# Virtual Topology


- Communicator attributes include **topology**
- In MPI, we use virtual topologies (such as cartesian coordinates / grids and graphs) rather than mapping directly to the underlying hardware
- This is very useful! Many programs hold either a structured grid or some type of unstructured graph

# Cartesian Topology




- What if all processes holding neighbors of a structured graph (or stencil code) need to communicate with one another
- We can tell MPI this is our virtual topology, and MPI will assume only neighboring nodes in this topology communicate
- ```
MPI_Cart_create(MPI_COMM_WORLD,  
               int n_dims,  
               int* dims,  
               int* periods,  
               int reorder,  
               MPI_Comm* comm_cart)
```

# Unstructured Graphs




- What if we have an unstructured graph (or sparse matrix) partitioned across the processes, and a random subset of processes need to communicate with one another
- ```
MPI_Graph_create(MPI_Comm old_comm,  
                 int nodes,  
                 const int indx[],  
                 const int edges[],  
                 int reorder,  
                 MPI_Comm * comm_graph
```

# Distributed Graph



- Typically, only know what is on my rank and which processes my rank is connected to
- `MPI_Dist_graph_create(MPI_Comm comm_old, int n, const int sources[], const int degrees[], const int destinations[], const int weights[], MPI_Info info, int reorder, MPI_Comm* comm_dist_graph)`

# Distributed Graph



- Typically, only know what is on my rank and which processes my rank is connected to
- `MPI_Dist_graph_create_adjacent(`  
    `MPI_Comm comm_old,`  
    `int indegree,`  
    `const int sources[],`  
    `const int sourceweights[],`  
    `const int outdegree,`  
    `const int destinations[],`  
    `const int destweights[],`  
    `MPI_Info info,`  
    `int reorder,`  
    `MPI_Comm* comm_dist_graph)`

# Neighborhood Collectives

- One of the reasons to use `MPI_Cart_create` and `MPI_Dist_graph_create_adjacent` is that you now know which processes should communicate with one another
- Instead of implementing that communication by hand, you can call a method such as `MPI_Neighbor_alltoallv`



# MPI\_Neighbor\_allgather

- Each process sends the same message to all neighbors
- Each process receives a message from each of their neighbors ('in degree')
- `MPI_Neighbor_allgather(const void* sendbuf,  
int sendcount,  
MPI_Datatype sendtype,  
void* recvbuf,  
int recvcount,  
MPI_Datatype recvtype,  
MPI_Comm comm)`

# MPI\_Neighbor\_alltoallv

- Each process sends the a message to each of the destination neighbors ('out degree' messages)
- Each process receives a message from each of their source neighbors ('in degree' messages)
- `MPI_Neighbor_alltoallv(const void* sendbuf,  
    const int sendcounts[],  
    const int sdispls[],  
    MPI_Datatype sendtype,  
    void* recvbuf,  
    const int recvcounts[],  
    Const int rdispls[],  
    MPI_Datatype recvtype,  
    MPI_Comm comm)`

# Neighborhood Collectives

- Unfortunately, the implementation behind neighborhood collectives is not currently optimized (this would be a great, longer term research project)
- This API allows for some really great optimizations under the hood (node-aware communication!)
- Ideally, these should be implemented to optimize communication under the hood (you know all information that is being sent / received, and so it is possible to aggregate that data on node and remove off-node communication)